MAD about AI

"Mutual Assured Destruction as a goal is, almost literally, mad. MAD." Donald Brennan coined the acronym to mock the policy, and he lost that argument. Decades later, thinking about MAD is still a trip through the looking-glass: defense is unstable, clarity is threatening, and insanity is rational.

Winter-Levy and Lalwani (2025) defend MAD's core assumptions in *Foreign Affairs* — second-strike survivability, C3 resilience, missile-defense asymmetry — and concede that perceptions of AI progress may drive destabilizing postural shifts. On missile defense I agree; the wall is kinematic and it holds. On second-strike and C3 I push back. On the perception concession I want to extend the argument.

C3 doesn't have to break — the commander only has to believe it might. Stuxnet's lesson from the target's side was: assume compromise, and if you detect an intrusion during a crisis, launch while you still can. AI accelerates the offensive side of this spiral. It doesn't have to pwn Chinese NC3; Beijing only has to worry that it might have, and cannot prove otherwise, because absence-of-intrusion is exactly what you cannot verify. Entanglement (Acton 2018) makes it worse: NC3 shares sensors and networks with conventional C2, so espionage is indistinguishable from disarming-strike prep.

AI is also the prize. If either side comes to believe the other is about to achieve a durable, militarily decisive AI lead, the incentive is to act before the window shuts — preventive-war logic (Copeland 2000) on a visible clock. Taiwan is the mechanism by which AI-as-prize reaches the nuclear layer: once China has domestic leading-edge fabs, the "broken nest" (McKinney and Harris 2021) reverses, and destroying TSMC becomes an access-denial lever rather than a cost to avoid.

The deepest problem is putting AI in the loop itself. Schelling argued that giving up control of your response makes threats more credible; Kavka noted rational agents cannot sincerely intend retaliation that saves nothing and kills millions. The Cold War leaned on training and duty. But Arkhipov reconsidered; from the posture designer's perspective, he was a reliability failure. Moscow's answer was Perimeter. AI is the natural extension — it resolves Kavka's paradox because the machine doesn't need to *want* retaliation, it only needs to lack the ability to change its mind. The alignment literature calls that ability corrigibility.

Every alignment desideratum — interruptibility, human-in-the-loop, a reliable off-switch — is anti-commitment. Posture designers chasing credibility will spec systems that safety engineers are trying to make impossible. To at least one important customer, a corrigible AI is a defective product.

NC3-AI integration is where arms control has to move next, and nobody is negotiating it.