Gladstone AI

Defense in Depth: An Action Plan to Increase the Safety and Security of Advanced AI

This publication was produced for review by the United States Department of State. It was prepared by Gladstone February 26, 2024 and the contents are the responsibility of the authors. The authors' views expressed in this publication do not reflect the views of the United States Department of State or the United States Government.

Organization

Gladstone AI Inc. (hello@gladstone.ai)

Authors

Edouard Harris*

Jeremie Harris

Mark Beall

* Lead and corresponding author.
Contact: edouard@gladstone.ai

Table of Contents

Executive summary 11

LOE1: Establish interim safeguards to stabilize advanced AI development 15

LOE2: Strengthen capability and capacity for advanced AI preparedness and response 16

LOE3: Increase national investment in technical AI safety research and standards development 17

LOE4: Formalize safeguards for responsible AI development and adoption by establishing an AI regulatory agency and legal liability framework 18

LOE5: Enshrine AI safeguards in international law and secure the AI supply chain 19

0. Introduction 21

0.1 Background 22

0.2 Categories of AI risk 24

0.2.1 Weaponization 24

0.2.2 Loss of control 24

0.2.3 Other risk categories 26

0.2.4 Risks addressed by this action plan 27

0.3 Potential sources of catastrophic AI risk 27

0.4 Arguments against regulation for catastrophic AI risk 29

0.4.1 Self-regulation will be sufficient 29

0.4.2 Regulation could damage U.S. innovation and competitiveness 30

0.4.3 Catastrophic AI risk could divert attention from other issues 32

0.4.4 Catastrophic AI risk mitigation is unnecessary 32

0.5 Challenges 33

0.5.1 Technical challenges 33

0.5.1.1 The worst-case outcomes for AI risk are considered both plausible and extremely severe 33

0.5.1.2 Timescales for catastrophic risk are uncertain 35

0.5.1.3 The degree of risk from loss of control is uncertain 36

0.5.1.4 Frontier labs lack safety and security measures to detect and prevent loss of control 37

0.5.1.5 Frontier labs lack sufficient security to prevent critical IP theft 38

0.5.1.6 Open-access development can elicit dangerous capabilities from previously safe AI models 39

0.5.1.7 Closed-access AI models are vulnerable to black-box exfiltration and other attacks 40

0.5.2 Political challenges 40

0.5.2.1 AI advances faster than the ordinary policy process 40

0.5.2.2 The information environment around advanced AI makes grounded conversations challenging 41

0.5.3 Economic and strategic challenges 42

0.5.3.1 Frontier labs face strong incentives to develop and deploy increasingly advanced AI systems with limited safeguards 42

0.5.3.2 AI supply chain proliferation cannot be mitigated after the fact 43

0.5.4 Legal challenges 45

0.5.4.1 The current legal environment is inadequate to address the most extreme risks from advanced AI 45

Action plan to increase the safety and security of advanced AI 47

LOE1: Establish interim safeguards to stabilize advanced AI development 48

1.1 National security threats addressed by this LOE 49

1.2 Establish an AI Observatory for advanced AI 50

1.2.1 Functions 51

1.2.2 Staffing 52

1.2.3 Interagency coordination 52

1.3 Establish responsible AI development and adoption safeguards for U.S. private industry 53

1.3.1 Implementation mechanism and statutory authorities 54

1.3.2 RADA safeguards principles 54

1.4 Establish an AI Safety Task Force for RADA oversight 56

1.4.1 Mission 57

1.4.1.1 Priority Objective: Finalize and secure agreement on RADA safeguards 57

1.4.1.2 Sustainment Component: Oversee compliance with RADA safeguards 61

1.4.1.3 Sustainment Component: Develop recommendations for future regulations 61

1.4.2 Resources and staffing 62
1.4.3 Financial requirements and budget 63
1.4.4 Authorities 63
1.4.5 Location 65
1.5 Securing the advanced AI supply chain 65
1.5.1 AI model weights 67
1.5.2 Cloud computing 67
1.5.3 AI hardware 69
1.5.4 Research collaborations 71
1.5.5 Education of foreign nationals 72
LOE2: Strengthen capability and capacity for advanced AI preparedness and response 73
2.1 Coordinate interagency working groups 74
2.2 Advanced AI education and training 75
2.2.1 Key stakeholders and learning outcomes 76
2.2.2 Suggested training topics 78
2.3 Indications and warnings 82
2.3.1 Key categories and sources of catastrophic risk 83
2.3.2 Improving I&W through bug bounties 86
2.3.3 Improving I&W through direct government investment 86
2.4 Contingency planning 87
2.4.1 Updating the National Preparedness System 88
LOE3: Increase national investment in technical AI safety research and standards development 89
3.1 Federally funded research in AI safety and security and AGI-scalable alignment 90
3.1.1 Levels of AI safety and security research 91
3.1.1.1 Open-source and public research 91
3.1.1.2 Research that requires access to proprietary AI models 92
3.1.1.3 National security research 93
3.1.2 Organizational framework 94
3.1.2.1 Considerations for AGI alignment research 95

3.1.2.2 Research collaborations 96

3.1.2.3 Sensitive research areas 96

3.2 Standards for AI evaluations and RADA safeguards 97

3.2.1 Limitations of AI evaluations 98

3.2.1.1 AI evaluations are not comprehensive 98

3.2.1.2 AI evaluations cannot confirm that a dangerous capability is absent 98

3.2.1.3 AI evaluations are highly vulnerable to manipulation 99

3.2.1.4 AI evaluations could fail systematically in high-capability regimes 100

3.2.2 Addressing the limitations of AI evaluations 100

3.2.3 Standards for RADA safeguards under catastrophic risk 101

LOE4: Formalize safeguards for responsible AI development and adoption by establishing an AI regulatory agency and legal liability framework 102

4.1 Establish the Frontier AI Systems Administration (FAISA) 103

4.1.1 FAISA appropriations and staffing 104

4.1.2 FAISA activities 105

4.1.2.1 Oversee licensing regime for safe AI scaling 105

4.1.2.2 Monitor AI hardware concentrations and AI programs 106

4.1.2.3 Enforce licensing and reporting requirements 107

4.1.2.4 Update technical thresholds for licensing 107

4.1.2.5 Maintain information repositories 108

4.1.3 Comprehensive framework for RADA safeguards 109

4.1.3.1 AI hardware designers (AIHDs) 112

4.1.3.1.1 Tier 1: AIHDs at or below the hardware threshold 114

4.1.3.1.2 Tier 2: AIHDs above the hardware threshold 114

4.1.3.2 Data center infrastructure providers (DCIPs) 114

4.1.3.2.1 Tier 1: DCIPs operating data center facilities at or below the power consumption threshold 116

4.1.3.2.2 Tier 2: DCIPs operating data center facilities above the power consumption threshold 116

4.1.3.3 AI hardware owners (AIHOs) 117

4.1.3.3.1 Tier 1: AIHOs who are at or below the aggregate compute capacity threshold 119

4.1.3.3.2 Tier 2: AIHOs who fall above the aggregate compute capacity threshold 119

4.1.3.4 AI model developers (AIMDs) 120

4.1.3.4.1 Tier 1: AI models below the registration threshold 122

4.1.3.4.2 Tier 2: AI models above the registration threshold, but below the approval threshold 123

4.1.3.4.3 Tier 3: AI models above the approval threshold, but below the controlled threshold 125

4.1.3.4.4 Tier 4: AI models above the controlled threshold 128

4.1.3.5 General provisions for all licensed entities 129

4.1.4 Publication controls 131

4.2 Legislative environment 133

4.2.1 Civil liability 133

4.2.2 Criminal liability 135

4.2.3 Emergency powers 137

4.3 Advanced AI in national security systems 139

LOE5: Enshrine AI safeguards in international law and secure the AI supply chain 141

5.1 Desired end state of this LOE 142

5.2 Build international consensus and partner capacity on catastrophic AI risk reduction 143

5.2.1 Coordinate domestic and international messaging and capacity-building 144

5.2.1.1 Policymaker education and outreach 145

5.2.1.2 Technical education and outreach 146

5.2.1.3 Other forms of capacity-building 148

5.2.2 Articulate and reinforce an official U.S. government position on catastrophic AI risk 149

5.3 Enshrine AI safeguards in international law 149

5.3.1 Treaty structure considerations 150

5.3.2 Preparation for a treaty 151

5.4 Establish an international regulatory agency 152

5.4.1 Monitoring and verification 153

5.4.2 Standard-setting 154
5.4.3 Convening researchers 154
5.5 Allied multilateral initiatives to manage the AI supply chain 155
5.5.1 Strategic considerations 156
5.5.2 Multilateral partnerships 156
5.6 Open challenges 157
Conclusion 159
Funding disclosure 162
Acknowledgements 163
Bibliography 170
Annex A: Glossary of terms 204
Annex B: The full challenge of AGI alignment 213
Annex C: Example AI alignment failure scenarios 214
C.1 Negligible impact: AI cheating at a video game 214
C.2 Low impact: AI sweeping bot 214
C.3 Medium impact: Dangerously creative drone 214
C.4 High impact: Electrical grid failure 215
Annex D: Advanced AI landscape 217
D.1 Frontier AI labs 217
D.2 China-based entities 218
D.3 Open-access developers 218
D.4 Elite quantitative hedge funds 219
Annex E: Funding in AI safety 221
Annex F: Persuasion and manipulation 222
Annex G: Primer on AI and compute 224
G.1 Compute as a pathway to frontier AI development 224
G.2 Implications for risk mitigation strategy 225

G.3 Key thresholds and existing compute concentrations 225

Annex H: AIO activities 229

Annex I: Voluntary Charter for responsible AI 231

I.1 Information sharing 231

I.2 Compute reporting threshold 233

I.3 Model evaluation protocols 234

I.4 Capability prediction protocols 242

I.5 Security measures 243

I.6 Model containment measures 245

I.7 AI safety and AGI alignment research 247

I.8 Dangerous capability ban 248

I.9 Capability research controls 248

I.10 Risk governance 249

I.11 Caps on cloud services for scaled training runs 250

I.12 Incremental adoption 250

I.13 Final considerations for negotiating Charter terms 251

Annex J: Effective compute 252

Annex K: ASTF activities and task-organization 254

K.1 Oversee and support compliance with RADA safeguards 254

K.2 Develop recommendations for a legal and regulatory regime 256

K.3 ASTF task-organization 257

Annex L: AI safety and security research topics 260

L.1 AGI-scalable alignment between AI behaviors and human values 260

L.2 Monitoring of AI systems to detect misbehavior and preempt failures 261

L.3 Robustness in the face of adversaries and unforeseen circumstances 262

L.4 Hardware-based verification 262

Annex M: Secure temporary storage of model weights 263

Annex N: Training approvals process for high-risk AI models 265

N.1 Dangerous capability predictions 266
N.2 Capability evaluations 267
N.3 Capability predictions at training checkpoints 267
N.4 Data controls 267
N.5 Risk management 268

Annex O: Training stage monitoring for high-risk AI models 269

O.1 Training journal 269
O.2 Ongoing capability predictions 270
O.3 Ongoing capability evaluations 271
O.4 Model output predictions 272
O.5 Training kill switch 273
O.6 Failed evaluation protocol 274

Annex P: Deployment stage approvals for high-risk AI models 276

P.1 Lifecycle funds 277
P.2 Pre-deployment KYC 278
P.3 Risk and contingency planning 278
P.4 Weight storage for model evaluations 278
P.5 Model augmentations 278
P.6 Failed evaluation protocol 279
P.7 Fine-tuning approval 279

Annex Q: Deployment stage monitoring of high-risk AI models 281

Q.1 Usage pattern reporting 282
Q.2 Red teaming 283
Q.3 Failed evaluation protocol 283

Table of Figures

Figure 1. Overview of the action plan and its component LOEs. 13

Figure 2. Visualization of an approximate spectrum of AI risks, ranging from individual and societal risks to catastrophic national security risks. 27

Figure 3. Visualization of the AI supply chain, potential sources of catastrophic risk, and categories of catastrophic risk. 29

Figure 4. Simplified representation of the components of the AI supply chain. 43

Figure 5. Safeguards to introduce to frontier AI developers, AI hardware owners, and data center infrastructure providers in the event that RADA safeguards can (Plan A) and cannot (Plan B) be mandated. 60

Figure 6. Summary of an example licensing regime for RADA safeguards. 111

Figure 7. Summary of an example AI hardware designer licensing regime for RADA safeguards. 113

Figure 8. Summary of an example data center infrastructure provider licensing regime for RADA safeguards. 115

Figure 9. Summary of an example AI hardware owner licensing regime for RADA safeguards. 118

Figure 10. Summary of an example AI model developer licensing regime for RADA safeguards. 121

Figure 11. Summary of key licensing requirements at each stage of the AI product lifecycle under RADA safeguards. 127

Figure 12. Visualization of our definitions of advanced AI, frontier AI, and AGI. 209

Figure 13. Organizational chart for the AI Safety Task Force. 257

Table of Tables

Table 1. U.S. government stakeholders and potential target learning outcomes of AI education. 77

Table 2. Training topics associated with target learning outcomes. 79

Table 3. Potential sources of catastrophic AI risk, corresponding risk categories, and information requirements associated with each. 83

Table 4. Potential inaugural members of an ASCCR and the components of the AI supply chain over which they exert the most influence. 157