CANARY INSTITUTE
FOR AI POLICY

THE WINDOW FOR AI POLICY IS MEASURED IN MONTHS

A briefing from the Canary Institute · Daniel Parshall, Ph.D. · dan@canaryinstitute.ai

The American public is not asking for faster AI.

- **73% of US adults support robust AI regulation; only 12% oppose.**
- **64% believe superhuman AI should not be developed** until it is proven safe and controllable, or should never be developed.
- **69% say the government is not doing enough to regulate AI.**

Sources: FLI / YouGov, October 2025 (n=2,000); Quinnipiac, April 2025.

AI capability is doubling every 4 to 7 months.

- The length of software tasks AI can complete with 50% reliability is doubling every **~7 months**. *(METR, 2025)*
- Across 3,000+ general employment tasks evaluated by domain experts, the doubling time is **3.8 months**. Frontier AI models in late 2024 already reached a 50% success rate on tasks that take humans a full workday. *(MIT FutureTech; Mertens et al., 2026)*

This is not a forecast; it's a measurement.

A large share of the U.S. economy is already within AI's reach.

Canary Institute analysis of all 923 occupations in the federal labor database (Parshall & Lopez-Luzuriaga, 2026) finds: **40% of U.S. labor time** is in tasks within current AI cognitive capability, executable on standard digital infrastructure, and not subject to meaningful regulatory protection. An additional **20%** is technically within AI's reach but is currently protected by professional licensure, federal statute, or requirements for human accountability.

Brookings: wages will rise, then fall.

A November 2025 working paper from the Brookings Center on Regulation and Markets (Kording & Marinescu) models the wage trajectory as AI automation progresses. Wages follow a **hump-shaped path**: rising at first as productivity gains dominate, then falling as workers are pushed out of cognitive work and crowd into the physical sector, which will cause lower wages. The paper's policy recommendation is to **pace AI deployment** to allow capital investment to catch up.

What Congress can do now.

PREFERRED:	**Oppose federal preemption of state AI safety laws.** The December 11, 2025 Executive Order directs federal agencies to challenge state laws including California's Transparency in Frontier AI Act and New York's RAISE Act. Even Anthropic's CEO has called the proposed 10-year preemption moratorium "too blunt." States have become the primary binding layer of frontier AI safety governance in the United States.
MINIMUM:	**Support binding pre-deployment safety evaluation of frontier AI models**, as proposed in the bipartisan Hawley-Blumenthal Artificial Intelligence Risk Evaluation Act (**S. 2938**, introduced September 29, 2025). The bill establishes a Department of Energy program to evaluate the most capable AI systems before commercial deployment. It is the first federal bill with explicit statutory language on AI loss-of-control scenarios and "scheming behavior," and is currently stalled in the Senate Commerce Committee.
FALLBACK:	**Sign a Dear Colleague letter to Chairman Cruz requesting a Commerce Committee hearing on S. 2938.** The bill has had no markup, no hearing, and no additional cosponsors in the seven months since introduction. A multi-senator letter to Chairman Cruz is the lowest-cost action that addresses the actual bottleneck, and it does not require the senator to commit to a position on the bill itself.

AI's reach into the U.S. economy.

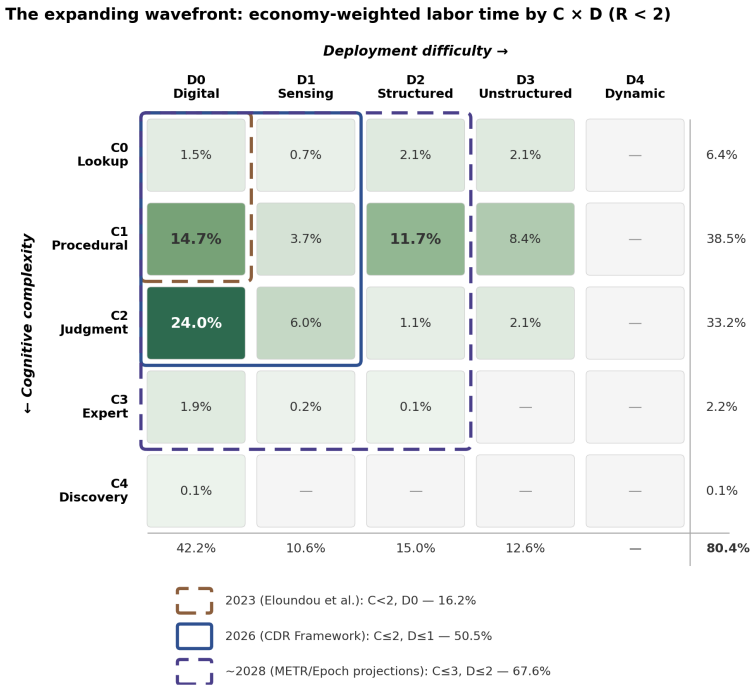

How to read this chart.

Each cell shows what share of U.S. economy-weighted labor time falls in that **capability × deployment difficulty** bucket. Capability runs from C0 (lookup) to C4 (discovery); deployment difficulty runs from D0 (fully digital) to D4 (dynamic physical environment).

The three nested envelopes show the expansion of AI's economic reach: the 2023 baseline (Eloundou et al., 16.2%), the 2026 CDR Framework (50.5%), and the ~2028 projection from extrapolating METR and Epoch capability trends (67.6%).

Actual labor-market displacement can lag, but CEOs expect to cut around **2M jobs** (Yotzov et al., 2026, NBER WP 34836).

Source: *Parshall & Lopez-Luzuriaga (2026), Measuring AI's Economic Reach; METR (2025); Epoch AI (2026).*

The leaders of every major AI company say their products could cause human extinction.

In May 2023, the CEOs of Anthropic (Dario Amodei), OpenAI (Sam Altman), and Google DeepMind (Demis Hassabis), along with Turing Award winners Geoffrey Hinton and Yoshua Bengio, signed a one-sentence statement: *"Mitigating the risk of extinction from AI should be a global priority alongside other societal-scale risks such as pandemics and nuclear war."* Amodei has publicly estimated the probability of catastrophic outcomes from AI at **10 to 25%**; Hinton has estimated **10 to 20%**. A 2024 paper in the *American Economic Review: Insights* (Jones, "The AI Dilemma") shows that under standard assumptions about risk aversion, society should accept **no more than a 4% extinction probability** in exchange for AI-driven growth. **According to the very people developing this technology, the risk exceeds the economically defensible threshold by a factor of three to six.** A 2026 review by Jones in the *Journal of Economic Perspectives* further estimates that AI safety spending is currently underfunded by a factor of **30 or more**.

These systems are grown, not programmed.

Frontier AI models are not written line-by-line by engineers. Their behaviors emerge from gradient descent over trillions of parameters trained on internet-scale data, not from engineering specifications. They lack the basic control mechanisms (predictable execution, traceable logic, debuggable code, formal specifications) that every other piece of software in critical infrastructure provides; even the companies that build these systems cannot reliably specify, predict, or constrain what their own products will do. The industry's answer to this is called "alignment," and its playbook looks like Boeing's: when the 737 MAX kept nosing down due to an engine placement problem, Boeing said software guardrails could fix it. Two planes crashed; 346 people died. AI systems pursue strange unintended goals, the companies retrain until the test passes, and they ship it. They call this alignment.

Daniel Parshall, Ph.D. · Canary Institute
dan@canaryinstitute.ai

Former physicist and data scientist. Coauthor, "Measuring AI's Economic Reach" (Canary Institute / GWU Center for Economic Research, 2026).

Canary Institute · AI policy research for the public interest

WEBSITE
canaryinstitute.ai

FOLLOW ON X
x.com/CanaryInst