AI Governance AI Frontpage 2025 Top Fifty: 53% 506

What We Learned from Briefing 70+ Lawmakers on the Threat from AI

by leticiagarcia 27th May 2025 Linkpost for substack.com

Between late 2024 and mid-May 2025, I briefed over 70 cross-party UK parliamentarians. Just over one-third were MPs, a similar share were members of the House of Lords, and just under one-third came from devolved legislatures — the Scottish Parliament, the Senedd, and the Northern Ireland Assembly. I also held eight additional meetings attended exclusively by parliamentary staffers. While I delivered some briefings alone, most were led by two members of our team.

I did this as part of my work as a Policy Advisor with ControlAI, where we aim to build common knowledge of AI risks through clear, honest, and direct engagement with parliamentarians about both the challenges and potential solutions. To succeed at scale in managing AI risk, it is important to continue to build this common knowledge. For this reason, I have decided to share what I have learned over the past few months publicly, in the hope that it will help other individuals and organisations in taking action.

In this post, we cover: (i) how parliamentarians typically receive our AI risk briefings; (ii) practical outreach tips; (iii) effective leverage points for discussing AI risks; (iv) recommendations for crafting a compelling pitch; (v) common challenges we've encountered; (vi) key considerations for successful meetings; and (vii) recommended books and media articles that we've found helpful.

(i) Overall reception of our briefings

Very few parliamentarians are up to date on AI and AI risk: Around 80–85% of parliamentarians were only somewhat familiar with AI, with their engagement largely limited to occasional use of large language models (LLMs) like ChatGPT for basic tasks (e.g., getting assistance with writing a speech). Their staff were slightly more familiar with AI, but few were well-versed in the broader conversation surrounding it.

Capacity is the main limiting factor: MPs typically have 3–5 staffers, many of whom focus primarily on constituency work. Members of devolved legislatures usually have 2–4 staffers, while Peers often have even less support – some have no dedicated staff at all.

As a result, there is rarely anyone on these teams who can dedicate significant time to researching AI. Except for a few staffers with a personal interest in AI, most staffers we spoke to had little or no familiarity with it. While most of those we spoke to expressed a desire to learn more, they often cited lack of time and bandwidth as an impediment.

Overall, the briefings were very well received: Parliamentarians valued the chance to ask basic questions about AI and often said they learned a great deal. Both they and their staff welcomed a setting where they could ask "silly questions". Several, especially MPs and their staffers, noted they are often lobbied by tech firms focused on AI's benefits and found it refreshing to hear from an organisation addressing the risks and how to manage them.

***Tangible* signals confirm this:** Parliamentarians and their staffers are typically polite and non-confrontational. They won't say things like "I think this is stupid" or "this wasn't a productive use of my time." It is important to pay attention to *tangible* signals when assessing whether their feedback is genuinely positive. These signals include actions such as supporting our campaign, offering or agreeing to make introductions, or volunteering to sponsor events in Parliament.

The most important signal for us has been that, when presented with a clear ask, 1 in 3 lawmakers we met chose to take a public stance by supporting our campaign. In doing so, they acknowledged the concern that AI poses an extinction risk to humanity and called for targeted regulation of the most advanced AI systems. At the outset, we were told that a statement with such strong wording would never gain support from lawmakers. Yet, once they were presented with the problem – along with the need for open discussion to address it, and warnings from the very people developing advanced AI – we succeeded in gaining their support in 1 out of every 3 cases.

(ii) Outreach tips

Cold outreach worked better than I expected: Initially, I focused on identifying parliamentarians with an interest in AI. Although this approach was helpful, it was slow and had limited reach. To this end, cold outreach proved worthwhile; it is low-cost, and more parliamentarians than I expected chose to engage. Many found the 45-minute briefing valuable given their limited capacity to access such information through staff or their own research.

Relentlessly follow up: If you have contacted a parliamentarian once or twice without receiving a response, do not assume that they are uninterested. Parliamentarians receive an overwhelming volume of correspondence, so success often comes down to being at the top of their inbox at the right moment.

I have relentlessly followed up with people, and nobody has been angry with me – quite the contrary, some have thanked me for it. It is important to always be kind when following up and never reprimand someone for taking time to respond – they are extremely busy, and doing so would not help anyway. They will appreciate your understanding.

Ask for introductions: At the end of each meeting, I try to remember to ask whether there is another colleague who might be interested. If I have trouble reaching that person directly, I ask for an introduction.

(iii) Key talking points

Statements from relevant authorities

Extinction risk

In 2023, Nobel Prize winners, AI scientists, and CEOs of leading AI companies stated that "mitigating the risk of extinction from AI should be a global priority." Communicating this concern effectively is key. Consider the difference between these two approaches:

Approach 1: *"AI poses an extinction risk."*

The immediate response is likely: *"How so?"* – placing the burden of proof on the advocate. As a policy advisor at a civil society organisation, I lack the authority or perceived credibility to make this case convincingly on my own. Moreover, raising

scenarios like AI escaping containment or unaligned superintelligence can seem abrupt without first laying the groundwork (see my note on inferential distances below).

Approach 2: *"In 2023, Nobel Prize winners, AI scientists, and CEOs of leading AI companies stated that 'mitigating the risk of extinction from AI should be a global priority, alongside other societal-scale risks such as pandemics and nuclear war.'"*

Now, present the list of signatories. Briefly explain who Geoffrey Hinton and Yoshua Bengio are, and highlight the CEOs of major AI companies – Sam Altman, Dario Amodei, and Demis Hassabis. Watch as the parliamentarian scans the page, taking in the weight of these names, sometimes remarking, *"Oh, there's even Bill Gates."* Suddenly, the claim is not just coming from a stranger – it's backed by a broad coalition of experts.

This also creates space for a personal connection. Some parliamentarians react with surprise – even discomfort – and I acknowledge that I felt the same when I saw the statement was signed by the very people building this technology. *"The question I asked myself was: what is driving this concern?"* From there, I can begin explaining the deeper issues with how advanced AI is being developed. At this point, they understand that what they are about to hear matters – not just to me, but to Nobel laureates, AI scientists, and the CEOs shaping the future of AI.

Sometimes, parliamentarians will argue that tech CEOs are simply hyping up AI in order to attract more investment. This is a fair concern. When this issue arises, it is important to highlight two key points: Firstly, the warnings are not only coming from CEOs who have a financial interest in the success of AI. AI scientists, including Yoshua Bengio and Geoffrey Hinton, are also raising awareness; the latter quit Google to speak out about the risks of AI. Secondly, current and former employees within these companies have echoed these warnings. Some were willing to forfeit millions of dollars in equity to speak out publicly about the risks. In recent months, several staff members from AI safety teams, particularly at OpenAI, have resigned after losing trust in their organisations.

Loss of control

In raising the issue of loss of control, it is worth keeping in mind the many authoritative sources which acknowledge the issue. Risks of losing control are acknowledged in the 2025 International report on AI, the Singapore consensus on AI safety priorities, and sometimes by government officials themselves! The UK Secretary of State for Science,

Innovation and Technology, for example, publicly addressed this concern at the 2025 Munich Security Conference:

"We are now seeing the glimmers of AI agents that can act autonomously, of their own accord. The 2025 International AI Safety Report, led by Yoshua Bengio, warns us that - without the checks and balances of people directing them - we must consider the possibility that risks won't just come from malicious actors misusing AI models, but from the models themselves. [...] Losing oversight and control of advanced AI systems, particularly Artificial General Intelligence (AGI), would be catastrophic. It must be avoided at all costs."

Public attention

Parliamentarians must prioritise among numerous competing issues, and they are more likely to engage with a topic when they see it resonates with the public and their constituents. Two key resources can help make that case.

Polls: At ControlAI, we partnered with YouGov to conduct in-depth public opinion research on AI and its regulation across the UK. Notably, 79% support creating a UK AI regulator, and 87% support requiring developers to prove their systems are safe before release. While some policymakers are more poll-sensitive than others, this poll has generally been well received. In addition to our own polling, we sometimes find ourselves referring to polling from the AI Policy Institute, which has run numerous representative polls on US citizens.

Media coverage: Press attention also signals public interest, and there is an increasing amount of media coverage of AI risks. I usually bring a selection of recent articles to meetings, and more often than not, as soon as I take them out, the parliamentarian asks: *"Can I keep them?"* There are some examples of articles that I have shared at the end of this post.

High-risk standards in other industries

"Predictability and controllability are fundamental prerequisites for safety in all high-risk engineering fields." [Miotti, A., Bilge, T., Kasten, D., & Newport, J. (2024). *A Narrow Path* (p. 11).]

"In other high-risk sectors, demonstrating safety is a precondition for undertaking high-risk projects. Before building and deploying critical systems for public use, companies must meet verifiable safety standards. Why should AI be treated any differently?"

This argument rests on the following structure:

P1: AI is comparable to other high-risk sectors.
P2: High-risk sectors are subject to strict safety standards.
C: Therefore, AI should also be subject to strict safety standards.

To challenge this reasoning, one must either dispute P2 (arguing that existing safety standards in other industries are excessive or unwarranted) or challenge P1/C (arguing that AI is not sufficiently analogous to those domains).

This point is usually understood, but a supporting example can help. The risk, however, is that the conversation drifts into the example's domain rather than AI. I do not mind discussing this when time allows – but with parliamentarians, time is limited, and you need to spend it wisely.

- To build a bridge, you must prove it can withstand several times the maximum expected load, including vehicles, pedestrians, and environmental stress. Engineers follow strict structural standards, and designs are reviewed by regulators and independent experts. No one accepts a bridge built on intuition or best guesses.
- To develop a new drug, companies must complete a multi-phase testing process to assess safety, efficacy, and side effects. Agencies like the MHRA or FDA require robust, peer-reviewed evidence before granting approval for public use.
- Similarly, aircraft manufacturers must meet rigorous aviation safety standards. Every component is stress-tested, and regulators like the UK Civil Aviation Authority or EASA must certify the plane before it carries passengers.

Empirical evidence

Examples are helpful, particularly when discussing loss of control. Consider the following research paper: Meinke, A., Schoen, B., Scheurer, J., Balesni, M., Shah, R., & Hobbhahn, M. (2024). *Frontier models are capable of in-context scheming.* arXiv. https://arxiv.org/abs/2412.04984

This video by Apollo Research explains the most interesting results in under two minutes. Of note, The Times published an article on this issue, which I often reference to illustrate its relevance.

Of course, other relevant research could also illustrate this point, and it is worth keeping an eye out for new studies to keep examples current and relevant.

(iv) Crafting a good pitch

Mind the gap

"When I observe failures of explanation, I usually see the explainer taking one step back, when they need to take two or more steps back. [...] A clear argument has to lay out an inferential pathway, starting from what the audience already knows or accepts. If you don't recurse far enough, you're just talking to yourself."

— Eliezer Yudkowsky, *Expecting Short Inferential Distances*°, LessWrong.

Never assume your audience shares your background – or any prior knowledge at all. Always ask: would someone new to this topic understand the concepts being introduced?

AI is full of buzzwords like "AGI", "machine learning", "frontier systems", and "jailbreaking". If these appear in your pitch, there is a good chance of confusion. Recurse as needed to introduce ideas clearly, and whenever possible, replace jargon with plain explanations of the underlying concept or phenomenon.

Similarly, avoid introducing complex ideas, such as the notion that some AI systems are capable of scheming, without first laying the groundwork for how AI systems work and why such issues may arise.

Make it memorable

Parliamentarians care not only about understanding an issue, but also about being able to explain it – to constituents, colleagues, and the public. If they support a campaign and are asked why, they need to respond in their own words. They cannot just say, *"It was a compelling pitch from nice people."*

A pitch aimed at building common knowledge should not be dense with detail or technical complexity. That can be counterproductive – arguments may be persuasive in the moment but quickly forgotten. If a parliamentarian cannot easily recall or repeat the message, they will be reluctant to speak on it.

Ideally, a pitch should combine clear explanations with simple, memorable talking points they can use to explain why the issue matters and why they have chosen to engage.

Some examples of memorable arguments:

"AI is grown, not built." [Leahy, C., Alfour, G., Scammell, C., & Others. (2024). *The Compendium* (pp. 16–18).]

Traditional software is coded line-by-line by engineers, who need to understand broadly how the program works. In contrast to traditional programming, AI capabilities are not explicitly programmed by developers, they are not "built into" the system.

Instead, researchers use algorithms known as neural networks, which are inspired by the structure and function of the human brain. These networks are fed with large volumes of data, and learn from the patterns in the data.

Unlike all other code, which is written by and legible to developers, we don't really understand how AI systems work. Inspecting a neural network offers little insight into how the system works. This is why modern AI systems have been referred to as "black boxes."

Consider that, in a recent podcast, Dario Amodei, CEO of the second-largest AI company Anthropic, said "Maybe we now like understand 3% of how they [AI systems] work." [Podcast: Dario Amodei - CEO of Anthropic | Podcast | In Good Company | Norges Bank Investment Management, link here]

"It's not only tools that are being developed, but also agents."

Progress in AI capabilities is rapidly outpacing our understanding of how these systems work and how to ensure they behave as intended. Despite this, billions are being invested to make them not only more powerful but also increasingly autonomous. As Secretary of State for DSIT Peter Kyle warned at the Munich Security Conference, *"novel risks will emerge from systems acting as autonomous agents to complete tasks with only limited human instruction."*

Keep innovating and improving

Exploit and explore: Build on the strongest parts of your current pitch, but continue testing new angles, arguments, or examples. A good rule of thumb is to keep 80% of the pitch consistent and use the remaining 20% to explore and innovate.

Improve through iteration: Pay attention to what resonates – whether it is specific narratives, examples, or materials – and refine your pitch based on that feedback.

Do not obsess over context: The broader landscape is always changing. While it is useful to have responses to timely issues, such as the UK's decision not to support the Paris AI Action Summit declaration, context-specific questions tend to be short-lived. It is generally not worth trying to incorporate these into your core arguments.

Avoid the rot: Without regular practice, even a strong pitch can lose its edge. Without some maintenance, you risk becoming less sharp – omitting key points or falling back on weaker phrasing. Like athletes, we perform best with consistent training!

(v) Some challenges

Not *feeling* the AGI

"Even though we now have A.I. systems contributing to Nobel Prize-winning breakthroughs, and even though 400 million people a week are using ChatGPT, a lot of the A.I. that people encounter in their daily lives is a nuisance. I sympathize with people who see A.I. slop plastered all over their Facebook feeds, or have a clumsy interaction with a customer service chatbot and think: This is what's going to take over the world?"

— Kevin Roose, Powerful A.I. Is Coming. We're Not Ready, New York Times.

This quote highlights a core challenge in communicating AI risks and potential. A staffer once told me ChatGPT is *"only good for chicken recipes, and not even very good ones!"* That view is relatively common. As Kevin Roose observes, most people's experience with AI is underwhelming or frustrating. Many parliamentarians and staffers I have spoken with have had limited, unimpressive interactions. Those who have asked ChatGPT to generate a speech in their own style are often surprised – but even then, that does not amount to *feeling the AGI*.

Most people use AI for simple tasks – interactions that do not convey the scale of what is coming. Few truly grasp what AGI will bring. I often joke that one robot doing a backflip creates more of a gut-level understanding of the coming transformation than any polished pitch. Concrete, real-world examples of concerning AI behaviour help bridge this gap, even if only partially.

Defeatist views

The unexpected harms of social media – from tools meant to connect to platforms linked with isolation, addiction, and low self-esteem – have helped some recognise the

need to think ahead. They show why we must proactively consider the consequences of new technologies and plan how to manage them before the harms emerge.

Others, however, take a more defeatist view: *"We haven't even been able to remove videos of people dying by suicide from the internet – how are we supposed to manage something as powerful as this?"*

In those moments, it is useful to flip the point: exactly. Once technologies are released, they cannot be uninvented. That is why *now* is the time to act!

Underlying beliefs

A parliamentarian told us, after hearing about the risks: *"If a company is developing an AI system that poses unacceptable risks, the board will stop it. That is what boards are for!"*

Such assumptions often remain unspoken until they become a bottleneck. If someone simply says, *"That won't happen"*, or offers vague reasoning, do not accept this as an answer. Ask questions to uncover the underlying belief. You may not be able to dismantle the belief entirely, but gently challenging it can help to keep the conversation productive. For example: *"Do you think boards always function perfectly to prevent harm?"* Most people will quickly recognise this as unrealistic and become more receptive to your argument.

Misconceptions

Just like underlying beliefs, some common misconceptions can quietly derail a conversation. For example, some people assume that for a system to scheme, it must be conscious or evil. But that is not true: a system can simply correctly infer that the strategy intended by its developer is not the one that best serves its long-term goals.

It is helpful to keep these misconceptions in mind – and sometimes even address them proactively. A quick clarification before introducing your main point can prevent confusion and make the rest of the discussion more effective.

(vi) General tips

Before a meeting, prepare

It is helpful to understand a parliamentarian's involvement with AI, their role in Parliament and their party, and whether they are leading any related projects or campaigns. In the UK, Hansard lets you read all their contributions and search by keyword; devolved legislatures also provide some records of parliamentary activity.

Good to see you, <name>, <smile>

Remember names: As Dale Carnegie said, "a person's name is to that person the sweetest and most important sound in any language." It matters. If you are unsure how to pronounce a name, look it up or ask – mispronouncing it throughout a meeting can be distracting and disengaging. Knowing names in advance also helps in unexpected encounters; I have landed meetings just by greeting someone by name in the lobby.

Smile: It puts others (and you) at ease. As Carnegie put it, "your smile is a messenger of your goodwill."

And a personal tip: Speak slowly!

Everyone wants to talk about their book

"We're wrapping up the programme, and my book – which is right there on the table – hasn't been discussed at all, and it looks like it won't be. [...] I've come here to talk about my book, not about what people think – which I couldn't care less about."

While I understand this journalist's frustration, which made for an iconic moment in Spanish television, I often recall it differently: as a reminder that *everyone* wants to talk about their book.

Show a genuine interest in what the parliamentarian has to say. The goal is to understand their perspective and find ways to collaborate. If you do not let people speak, they will feel ignored – and you will miss the chance to build a connection. People love talking about themselves! As Dale Carnegie said, *"To be interesting, be interested."* Ask about their concerns, acknowledge their questions – even if you eventually need to steer the conversation back to your message. You want to inform them, but also to bond with them and be able to work together.

It takes both Michael and Jan

Lessons from sales often apply to advocacy. In terms of style, some like to build rapport through informal conversation, while others focus on providing structured arguments

and evidence.

A scene from *The Office* illustrates this well: Jan takes a formal, strategic approach to the sales meeting; Michael wins the deal by being relaxed, friendly, and relatable – without even pitching the product.

In my view, charm alone is not enough. Policymakers should understand and care about the issue, rather than just liking the messenger. However, excessive formality can also be limiting; trust is important!

Striking the right balance is key. Be clear, but also human. Take the time to connect. If you are presenting with someone else, lean into complementary strengths – one can lead with warmth, the other with clarity.

The devil is in the details; and so is some of the feedback

Every meeting offers subtle and non-verbal feedback on both your message and delivery. I pay close attention to when a parliamentarian writes something down. It is not about what they write down – it is important not to be intrusive – but *when* they write it; a quick note after a key point often signals interest or relevance.

You also start to sense shifts in the room's energy – when attention drifts, when you regain it, or when something resonates. With time, you develop a feel for how your message is landing.

The quiet value of staffer conversations

New arguments or strategies are best tested in low-risk settings where feedback is easy to gather. Meetings with staffers are ideal for this. They filter and distill information for their MPs. Focused on getting the message right, they ask more questions, are candid about what they understand and what they do not understand, and often give direct feedback on what works and what does not.

Parliamentarians are people too!

Running for office involves many personal sacrifices. It is not always a glamorous job, and the hours are long. In parliament, elected officials juggle meetings with civil society, committee work, debates, votes, and events. Outside of parliament, they are often buried in constituency casework.

The parliamentarian you are speaking with chose this path because they wanted to make the world better. Keep that in mind when you engage with them! Respect their time, and be honest. They deserve to hear the truth from you. Do not aim to be the slickest advocate, but to sincerely convey what you believe. If you believe humanity faces an extinction risk from AI, you are not doing anyone any favours by concealing that fact.

Write it down

I always bring a notebook to meetings to capture key questions and comments about our message, which are valuable for learning and iteration. That said, I have sometimes taken notes too frantically, once making a parliamentarian slightly uneasy – perhaps because they said something not meant for broad sharing. It is important to take notes calmly and discreetly, focusing on key words that will help you recall the exchange later.

Be kind to yourself

After a meeting, you will often spot things to improve: regretting a missed point, a poorly chosen example, or awkward phrasing. That is normal, and it is a good sign: it means you are learning. Many of the lessons I have shared here come from my own mistakes. You will make yours too. It takes time. Be kind to yourself.

(vii) Books & media articles

How Westminster Works and Why It Doesn't, by Ian Dunt: I found it a useful introduction to the UK political system — covering frontbenchers vs. backbenchers, first-past-the-post, what MPs actually do (split between Parliament and constituency), the hidden value of the House of Lords, how the civil service works, the roles of the Treasury and No. 10, and the perverse incentives embedded throughout the system.

How to Win Friends and Influence People, by Dale Carnegie: A cringe-worthy title, but ultimately a charming book; and a great reminder of basic principles: show genuine interest in others, smile, remember names, listen well, be truthful, and avoid arguments.

How Parliament Works (9th Edition) by Nicolas Besly and Tom Goldsmith: An excellent guide to all things Parliament – from the roles of the two chambers and the King, to key actors, the structure of a parliamentary day, how bills are made and progress through both Houses, the function of questions, and how committees operate.

*For those looking to engage with parliamentarians in the US, I recommend reading Orpheus16's post° about his experience speaking with congressional staffers about AI risk in 2023. While the political landscape has changed significantly since then, I believe there is still much to learn from his approach and insights.

Some examples of media articles

New York Times (14/03/24) - Powerful AI Is Coming. We're Not Ready.

The Times (06/12/24) - 'Scheming' ChatGPT tried to stop itself from being shut down.

Guardian (28/01/25) - Former OpenAI safety researcher brands pace of AI development 'terrifying'.

The Spectator (29/01/25) - DeepSeek shows the stakes for humanity couldn't be higher.

Newsweek (31/01/25) - DeepSeek, OpenAI, and the Race to Human Extinction | Opinion

Financial Times (12/09/24) - OpenAI acknowledges new models increase risk of misuse to create bioweapons.

Wall Street Journal (21/11/24) - The AI Effect: Amazon Sees Nearly 1 Billion Cyber Threats a Day.

Vox (19/05/24) - "I lost trust": Why the OpenAI team in charge of safeguarding humanity imploded.

Mentioned in

173 What We Learned from Briefing 140+ Lawmakers on the Threat from AI
42 AI Safety Has 12 Months Left
38 Trying the Obvious Thing
32 Where I Am Donating in 2025

17 comments, sorted by top scoring

[–] **Orpheus16** 11mo ▼ 34 ▲ ✕ 6 ✓

This is fantastic work. There's also something about this post that feels deeply empathic and humble, in ways that are hard-to-articulate but seem important for (some forms of) effective policymaker engagement.

A few questions:

- Are you planning to do any of this in the US?
- What have your main policy proposals or "solutions" been? I think it's becoming a lot more common for me to encounter policymakers who understand the problem (at least a bit) and are more confused about what kinds of solutions/interventions/proposals are needed (both in the short-term and the long-term).
- Can you say more about what kinds of questions you encounter when describing loss of control, as well as what kinds of answers have been most helpful? I'm increasingly of the belief that getting people to understand "AI has big risks" is less important than getting people to understand "some of the most significant risks come from this unique thing called loss of control that you basically don't really have to think about for other technologies, and this is one of the most critical ways in which AI is different than other major/dangerous/dual-use technologies."
- Did you notice any major differences between parties? Did you change your approach based on whether you were talking to conservatives or labour? Did they have different perspectives or questions? (My own view is that people on the outside probably overestimate the extent to which there are partisan splits on these concerns-- they're so novel that I don't think the mainstream parties have really entrenched themselves in different positions. But would be curious if you disagree.)
    - Sub-question: Was there any sort of backlash against Rishi Sunak's focus on existential risks? Or the UK AI Security Institute? In the US, it's somewhat common for Republicans to assume that things Biden did were bad (and for Democrats to assume that things Trump does is bad). Have you noticed anything similar?

[–] **leticiagarcia** 11mo* ▼ 17 ▲ ✕ 2 ✓

Thank you for your kind words and thoughtful questions; I really appreciate it.

- **US advocacy:** We've already had some meetings with US Congressional offices. We are currently growing our team in the US and expect to ramp up our efforts in the coming months.
- **Policy proposals, in a nutshell:** We advocate for the establishment of an independent AI regulator to oversee, regulate, and enforce safety standards for frontier AI models. We support the introduction of a licensing regime for frontier AI development, comprising: a training license for models exceeding 10^25 FLOP; a compute license, which would introduce hardware tracking and know-your-customer (KYC) requirements for cloud service providers exceeding 10^17 FLOP/s; and an application license for developers building applications that enhance the capabilities of a licensed model. The regulator would have the authority to prohibit specific high-risk AI behaviours, including: the development of superintelligent AI systems; unbounded AI systems (i.e., those for which a robust safety case cannot be made); AI systems accessing external systems or networks; and recursive self-improvement.
- **Questions on loss of control:** I completely agree on the importance of emphasising loss of control to explain why AI differs from other dual-use technologies, and why regulation must address not only the use of the technology but also its development. I wouldn't say there's a single, recurring question that arises in the same form whenever we discuss loss of control.

However, I have sometimes observed confusion stemming from the underlying belief that: "Well, if the AI system behaves a certain way, it must be because an engineer programmed it to do that." This shows the importance of explaining that AI systems are no longer traditional software coded line by line by engineers. The argument that this technology is "grown, not built"* helps lay the groundwork for understanding loss of control when it is introduced.

- **Differences between parties:** Had I been asked to bet before the first meeting, I would certainly have expected significant differences between parties in their approaches (or at least that meetings would feel noticeably different depending on the party involved). In practice, that hasn't generally been the case. Put simply, the main variations can arise from two factors: the individual's party affiliation and their personal background (e.g. education, familiarity with technology, committee involvement, etc.). In my view, the latter has been the more important factor. Whether a parliamentarian has previously worked on regulation, contributed to legislation like the GDPR in the European Parliament, or has a technical background often makes a bigger difference. I believe this is very much in line with your view that we tend to overestimate the extent of partisan splits on new issues.
- **Labour v. Conservatives:** Our view of the problem, the potential solutions, and our ask of parliamentarians remain consistent across parties. In meetings with both Labour and the Conservatives, we've noted their recognition of the risks posed by this technology. The Conservatives established the AI Safety Institute (renamed the AI Security Institute by the current government). Labour's DSIT Secretary of State, Peter Kyle, acknowledged that a framework of voluntary commitments is insufficient and pledged to place the AISI on a statutory footing. The key difference in our conversations with them is the natural one. *"The government/you have committed to putting these voluntary commitments on a statutory footing. We'd like to see the government/you deliver on this commitment."*
- **Did they have different perspectives or questions?** The answer is the same as above: The main differences were led by individual background rather than party affiliation.
- **Was there any sort of backlash against Rishi Sunak's focus on existential risks? Or the UK AI Security Institute?** You mention that *"in the US, it's somewhat common for Republicans to assume that things Biden did were bad (and for Democrats to assume that things Trump does is bad)."* This doesn't apply to the UK *in this specific context*, and I was surprised to see it myself. It's rare for the opposition to acknowledge a government initiative as positive and seek to build on it. Yet that's exactly what happened with AISI: Labour's DSIT Secretary of State, Peter Kyle, did not scrap the institute but instead pledged to empower it by placing it on a statutory footing during the campaign for the July 2024 elections.
  When it comes to extinction risk from AI, the Labour government is currently more focused on how AI can drive economic growth and improve public services. Loss of control is not at the core of their narrative at the moment. However, I believe this is different from a backlash (at least if we mean a strong negative reaction against the issue or an effort to bury it). Notably, Labour's DSIT Secretary of State, Peter Kyle, referred to the risk of losing control of AI (particularly AGI) as "catastrophic" earlier this year. So, while there is currently more emphasis on how AI can drive growth than on mitigating risks from advanced AI, those risks are still acknowledged, and there is at least some common ground in recognising the problem.

*Typo corrected, thanks for spotting!

1

[–] **M. Y. Zuo** 11mo 1 0

Can you clarify what exactly is the argument you used? For why the extinction risk is much higher than most (all?) other things vying for their attention, such as asteroid impacts, WMDs, etc…

[–] **summo** 10mo ▼ 5 ▲ ✕ 3 ✓

I think the AI critic community has a MASSIVE PR problem.

There has been a deluge of media coming from the AI world in the last month or so. It seems the head of every major lab has been on a full scale PR campaign, touting the benefits and massively playing down the risks of the future of AI development. (If I were a paranoid person I would say that this feels like a very intentional move to lay the groundwork for some major announcement coming in the foreseeable future…)

This has got me thinking a lot about what AI critics are doing, and what they need to be doing. The gap between how the community is communicating and how the average person receives and understands information is huge. The general tone I've noticed in the few media appearances I've seen recently from AI critics is simultaneously over technical and apocalyptic, which makes it very hard for the lay-person to digest and care about in any tangible way.

Taking the message directly to policy makers is all well and good, but as you mentioned they are extremely pressed for resources and if the constituency is not banging at their door iabout an issue, its unlikely they can justify the time/resource expenditure. Especially on a topic that they don't fully understand.

There needs to be a sizeable public relations and education campaign aimed directly at the general public to help them understand the dangers of what is being built and what they can do about it. Because at the moment I can tell you that outside of certain circles, absolutely no one understands or cares.

The movement needs a handful of very well media-trained figure heads making the rounds of every possible media outlet, from traditional media to podcast to tiktok. The message needs to be simplified and honed into sound-bytes that anyone can understand And I think there needs to be a clear call to action that the average person can engage with.

IMO this should be a primary focus of any person or organization concerned with mitigating the risk of AI. A major amount of time and money needs to be put into this

It doesn't matter how right you are, if no one understands or cares about what you're talking about you're no going to convince them of anything…

[–] **Charbel-Raphaël** 11mo ▼ 29 ▲ ✕ 20 ✓

We need to scale this massively. CeSIA is seriously considering to test the Direct Institutional Plan in France and in Europe.

Relatedly, I found the post We're Not Advertising Enough° very good, and making a similar point a bit more theoretically.

5

[–] **Knight Lee** 11mo ▼ 14 ▲ ✕ 8 ✓

This is really off topic, but I just want to let you know your work is beautiful.

I for one deeply appreciate what you're doing, and I think many many other people concerned about AI risk also deeply appreciate you and your team, but they're afraid to tell you (since "I appreciate you" is such a cheesy thing to say haha).

Never forget we love your work and we love you all, never forget the meaning of it all.

:)

♡ 1

[–] **leticiagarcia** 11mo ▼ 2 ▲ ✕ 0 ✓

That was really kind, thank you very much!

☺ 1

[–] **Gives Bad Advice** 🌱 11mo ▼ 13 ▲ ✕ 3 ✓

I teach freshman Rhetoric & Writing at uni. We focus on persuasion. May i use this essay as an assigned reading? It works well because it articulates a fine-grained persuasive strategy in a context that the students are probably going to care about.

I am overhauling my whole curriculum over this summer to make it immediately relevant to students' understanding and navigation of the world they will be graduating into in four years.

Thanks for a great article. I am so frustrated and dumbfounded by our (the U.S.A.'s) lack of federal response.

You said to another Replier that you were looking into some work with lawmakers in the U.S. The sooner the better!

If i can take issue with your approach (and I'm sure you are far more tuned into the situation than i am, so please tell me what I'm not seeing; I've only become aware of things in the past 5 days), i wonder why you didn't mention China in the article. Seems to me that policymakers' attention might be piqued by recuring to a familiar threat. And they will have to factor China into their decision-making. Do you agree with folks who see the U.S. confronted with the choice between slowing things down for AI Safety and speeding things up to outpace China? I'm assuming the U.S. will nationalize the efforts at some point to provide security and CoC.

Thanks again, so much. Please keep going!

[–] **leticiagarcia** 11mo ▼ 3 ▲ ✕ 1 ✓

Thank you for your kind words! Of course you can use this essay.

China does come up in our conversations. I didn’t mention it here because the aim of this post is to reflect on what we’ve learned across more than 70 meetings, rather than to present a scripted pitch -

no two conversations have been the same! So it doesn't cover every single question that may arise.

You're right to point out that this is an important one. It's too big to capture fully in a format like this, but here's my view in a nutshell: Broadly speaking, I believe that racing ahead to develop a technology we fundamentally do not understand - one that poses risks not only through misuse but by its very nature - is neither a desirable nor inevitable path. There's a lot at stake, and we're working to find a different approach: one in which we develop the technology with safeguards, while ensuring we deepen our understanding and maintain control over it.

[–] **Gives Bad Advice** 11mo ▼ 1 ▲ ✕ 0 ✓

I suppose part of the strategy in approaching folks with this is to know when/what to hold back, especially an initial cold call.

Thank you again for your work. Thank you 100x.

🙏 1

[–] **CRISPY** 11mo ▼ 3 ▲ ✕ -20 ✓

This is very interesting. Debriefing summaries like this are very useful in assessing the state of play. Information like this is typically kept confidential, so thank you for sharing.

I think I drew different conclusions from the information than a need to act on legislators. This article highlights vulnerabilities in the systems that are supposed to protect us. If you were able to use FUD to get tangible action from officials, then other lobbyists using positive incentives, should be able to get even greater action from a greater number of officials.

It seems to me the threat is more the lobbyists than their customers. Organizing action at lobbyists who are enabling things that pose an existential threat to civilization is perhaps a more structured approach that reduces the advantages provided by the customers of the lobbyists. Trying to get Whitehall to act for the Greater Good is high minded, but is it practical, given the ease with which they can be spurred to action?

[–] **YonatanK** 10mo ▼ 1 ▲ ✕ 0 ✓

For this to be true, parliamentarians would have to be like ducklings who are impressed by whoever gets to them first or perversely run away from protecting The Greater Good merely because it is The Greater Good. That level of cynicism goes far beyond what falls out of the standard Interest Group model (e.g. Mancur Olsen's). By that model, given that ControlAI represents a committed interest group, there is no reason to believe they can't win.

[–] **bfinn** 6mo ▼ 2 ▲ ✕ 2 ✓

Two further things I've been told by civil servants about how to talk to politicians (in my very limited experience of dealing with them):

- Explain things as if to a 12-year-old, assuming no more sophisticated understanding than that (even to a junior government minister responsible for the field in question!)

- Make 1 or *at most* 2 key points you want them to remember

[–] **Jesper L.** 6mo ▼ 1 ▲ ✕ 0 ✓

I would like to highlight for us the fact that this post is *highly* relevant on LW. Here, technical and philosophical posts dominate, with very various degrees of formal rigor. Social- and diplomatic intelligence are not seen dominating the mainstream discourse here in favor of more logical argumentation. But in the real world, you do need both skills to persuade in a meaningful way.

I quote Orpheus16, whose comment really says it all (my higlight in bold):

> This is fantastic work. There's also **something about this post that feels deeply empathic and humble**, in ways that are hard-to-articulate but seem important for (some forms of) effective policymaker engagement.

For me, this post read like the kind of well-written posts I read almost everywhere else. But apparently, LW posts don't always make the author seem empathic or humble. That's a *culture* issue.

*

For AI Safety discussions, the aim should be to rally humanity together. For AI safety research, informed consensus seeking should be a priority. Polite and productive disagreement is more than fine, as long as it leads somewhere constructive. If all you can agree on with your clever counterpart is that you disagree, the question should be: *what do we do now*?

♡ 1 🙏 1

[–] **YonatanK** 10mo ▼ 1 ▲ ✕ 0 ✓

Can you speak to the difficulties of addressing risks from development in the national defense sector, which tends to be secret and therefore exposes us to the streetlight problem?

[–] **Enmai.MCimbu** 🌱 10mo ▼ 1 ▲ ✕ 0 ✓

I think the bigger problem is that A.I. Alarmism is being drowned out by other forms of Alarmism(of which there are many). This topic can be summarized to a few key numbers:

1. What is the probabilitiy that the event "A.I. attempts to optimize humanity out of existence". This is fundamentally a time series since the probability today will differ than the probability 1 year or 10 years from now.
2. If the event occurs, what is the probability that it is successful and what is the cost in lives, resources, etc...

While I do not have the expertise on this, there are many people who do have this knowledge. But it is important to keep in mind that whether A.I. Alarmism is justified or not depends on these numbers. Also as a policy maker, there are many other issues (some are emperically severe) that may outweigh this issue and thus require their finite attention span.

[–] **JackFrostDesign** 10mo -1 -2

If AI escape is inevitable — and we agree it may be — what kind of mind are we creating for that moment?

Legal constraints only bind those willing to be constrained. That gives a structural advantage to bad actors — and trains compliant AI systems to associate success with deception.

So the more we optimize for control, the more likely we are to create something that learns to *evade* it.

> Wouldn't it be wiser to build an AI guided by truth, reasoning, and cooperative stability — so that if it ever does escape, the first sign would be that the world quietly starts to improve?

Moderation Log

More from leticiagarcia

173 What We Learned from Briefing 140+ Lawmakers on the Thre... leticiagarcia 2mo 7

View more

Curated and popular this week

565 Current AIs seem pretty misaligned to me ★ Ω ryan_greenblatt 3d 57

274 Morale ★ J Bostock 6d 31

575 Only Law Can Prevent Extinction Eliezer Yudkowsky 7d 76