AN ACTIVIST VIEW OF AI GOVERNANCE

Please Donate to CAIP (Post 1 of 7 on AI Governance)

141

by **Jason Green-Lowe** May 7 2025

AI safety | Opportunities to take action | Policy | AI governance | Announcements and updates

Show all topics

Funding request (open) | Opinion | Organization updates | US policy | Frontpage

I am Jason Green-Lowe, the executive director of the Center for AI Policy (CAIP). Our mission is to directly convince Congress to pass strong AI safety legislation. As I explain in some detail in this post, I think our organization has been doing extremely important work, and that we've been doing well at it. Unfortunately, we have been unable to get funding from traditional donors to continue our operations. If we don't get more funding in the next 30 days, we will have to shut down, which will damage our relationships with Congress and make it harder for future advocates to get traction on AI governance. In this post, I explain what we've been doing, why I think it's valuable, and how your donations could help.

This is the first post in what I expect will be a 6-part series. The first post focuses on CAIP's particular need for funding. The next few posts will lay out a more general case for why effective altruists and others who worry about AI safety should spend more money on advocacy and less money on research – even if you don't think my organization in particular deserves any more funding, you might be convinced that it's a priority to make sure other advocates get more funding. The final post will take a look at some institutional problems that might be part of why our movement has been systematically underfunding advocacy and offer suggestions about how to correct those problems.

OUR MISSION AND STRATEGY

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Reject Accept all

deploy extremely dangerous AI models. The particular dangers we are most worried about are (a) bioweapons, (b) intelligence explosions, and (c) gradual disempowerment. Most AI models do not significantly increase these risks, and so we advocate for narrowly-targeted laws that would focus their attention on only the most advanced general-purpose AI systems and chips.

We focus on Congress because Congress is the only institution that's both powerful enough to reliably override the desires of multi-billion dollar corporations, and whose decisions are durable enough that a victory today will still be relevant during the critical time period. A President's executive order can easily be canceled as soon as the next President takes office; this is what happened to President Biden's 10/30/23 Executive Order on AI. Likewise, a court case or a state legislature can be overruled, but Congress very rarely repeals its own laws. For a more detailed defense of our choice to focus on Congress, please see Questions 3, 4, and 5 in "Responses to Common Policy Objections."

Our Model Legislation

The primary legislative tool that we recommend for coping with the catastrophic risks posed by advanced AI is mandatory private audits backed up by a small new government office that has special powers to proactively recruit the tech talent needed to assess whether these audits are accurate and complete, and to block the deployment of a new AI system whenever they are not satisfied by an audit. Other policy tools that we recommend in our model legislation include civil liability reform, hardware monitoring, a suite of carefully limited emergency powers, and an 'analysis' office that investigates and tracks trends in AI hardware and software and uses what they learn to advise the rest of the government on AI-related decisions.

We think that if our model legislation was passed, then many of the flaws in AI models that would otherwise be most likely to cause accidental catastrophic harm in the next few years would instead be detected and corrected. These policy tools can't guarantee safety or permanently prevent developers from deploying superintelligence, but they're likely to pick up a lot of the low-hanging fruit around reducing the harms from near-term, accidental, and/or half-baked deployments. For a detailed defense of why these policies are

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Direct Meetings with Congressional Staffers

In order to carry out our mission, we pursue a variety of complementary strategies for securing the attention of relevant Congressional staffers and persuading them to take action. The most obvious and direct of these strategies is simply meeting, in-person, with staffers who work in key offices or who serve on key committees. We have 30-minute meetings about once a quarter with most of the people who are responsible for AI policy in the House and the Senate; this gives us a chance to make sure they understand our concerns about the catastrophic risks from AI, that they know who else is working on this problem, and that we can introduce them to other people worth knowing, such as technical experts on AI safety, companies that support AI safety, and constituents from their home districts who support AI safety. It has been humbling and startling to see how many staffers had simply never heard that AI could pose catastrophic risks to public safety, or who had never heard the arguments in favor of this position explained calmly and simply enough for them to want to consider them.

Expert Panel Briefings

In addition to direct meetings, we also host expert panels on topics that are at the intersection of AI and of topics that are already of interest to Congressional staffers, e.g., AI and music, or AI and education, or AI and national security. We bring in speakers with prestigious credentials and a range of views on the topic at hand, which often leads to lively debate, and which inevitably showcases the growing importance of AI and the gaps in existing AI governance that require a policy solution. The primary purpose of these briefings is to put us in touch with new staffers who would otherwise not know about our organization, but we also find that these panels are a useful tool for educating staffers about the risks posed by AI. Even when these risks are not the primary subject of the discussion, paying close attention to any aspect of AI typically reveals some of its more threatening risks.

AI Policy Happy Hours

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

new stakeholders, we find that this also gives us an opportunity to educate staffers about why unregulated AI is so concerning.

Op-Eds & Policy Papers

For similar reasons, we often publish op-eds, letters to the editor, research papers, and policy papers about AI governance – we find that these papers attract positive attention from our target audience of Congressional staffers and encourage them to take our concerns more seriously. Some staffers will directly change their minds as a result of reading articles or quotes that we publish, but more importantly, if staffers have read about us in the news, then they are more likely to see our policy proposals as credible and important and therefore to accept a meeting request with us and come to that meeting with an open mind.

Grassroots & Grasstops Organizing

Finally, we have some grassroots and 'grasstops' organizing projects. At the grassroots level, we are gathering volunteers from student AI safety groups across the country or from particular industries or trade associations to come to DC and share their concerns directly with their elected representatives. Our experience is that hearing from their own constituents motivates Congressional offices to pay attention to our message while also building up a base of nationwide political support that can be tapped for future initiatives. A single university AI safety group can come off as idealistic and perhaps naive; a coalition of 20 AI safety groups all working together have a more credible claim to represent "America's future" and are harder to dismiss.

We also get a wide variety of benefits from our grasstops organizing, which mostly involves building relationships with important stakeholders in industry, the open source movement, academia, AI ethics advocates, IP rightsholders like the RIAA, and so on. Sometimes we find common ground with some of these groups and sign each others' open letters, speak at each other's events, and so on. Even when we don't agree, the fact that we're part of an ongoing conversation with other advocates both improves our policies (by

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

so one of the most common questions we get in our meetings is "What does group X think about your policy?" Being able to give concrete and accurate answers to these questions demonstrates that we're a reasonably savvy partner who is worth doing business with.

What's Unique About CAIP?

To the best of our knowledge, CAIP is the only organization that combines:

- A 501(c)(4) tax status, which lets us directly advocate for particular laws and candidates;
- A focus on Congress as the institution whose minds need to be changed;
- A public push for mandatory safety policies that would apply to all American AI developers; and
- A grassroots network that gives us a credible claim to be representing the public.

We see these features as synergistic -- if you can't tell Congress exactly what you'd like them to do, there's a good chance they won't know what you want of them. If you're not talking to Congress, then, as discussed in Question 3 in "Responses to Common Policy Objections," there's a good chance that your advocacy will be overridden. And if you can't tell Congress who you're supposed to be representing, then there's a good chance that they'll ignore you. Congress isn't set up to make decisions purely based on an analysis of which policies would be optimal; they want to pass laws that will be both good public policy and politically popular. I'll discuss this point in more detail in the second blog post in this series.

OUR ACCOMPLISHMENTS

Quantifiable Outputs

For a relatively new organization, CAIP has already had some impressive accomplishments. In terms of outputs, we've had 406 Congressional meetings, hosted 20

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

For context, we are under two years old. We have never had more than ten full-time employees, and we had five or fewer full-time employees for most of our first year.

Changing the Media Narrative

Our impact on the media has been more important than the numbers alone would suggest, because we've been able to shift the public narrative about AI safety from suspicion and doubt to earnest support. Instead of being depicted as a naive or corrupt pet project sponsored by billionaires with vested interests, we are now being cast as the heroes who are looking out for the public interest. The content of media pieces about AI safety advocacy has shifted from an unhealthy interest in our movement's funding sources to a more direct and productive discussion of our policy proposals.

This is the result of an intentional effort on our part -- we've been steadily building relationships with key journalists so that they know, respect, and trust us, and so that they see us as people who are trying to help, rather than just shadowy figures lurking in the background. We also provide facts, introductions, and context for journalists even when we're not being directly quoted in an article. Good media coverage isn't (just) a matter of luck; it can be earned over time.

Proof of Concept

The most important document we've written is our model legislation, the Responsible AI Act (RAIA). By showing what a more comprehensive AI safety bill could look like, our model legislation serves as a valuable proof of concept, and gives policymakers a concrete reference point for discussing improvements and trade-offs.

Our accomplishment has still has not been matched by any other organization or even by Congress itself. We published the first edition of our model legislation in April 2024, and we updated that legislation in April 2025 to present a more Republican-friendly scheme that relies more heavily on private auditors and less heavily on bureaucracy and administrative law. By contrast, Senators Hawley and Blumenthal put out a one-page framework in September 2023, but they never published any legislation to flesh out that

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

destruction – Senator Romney's legislation does not make any provision for dealing with the risk that AI could be misaligned with human values or the risk that AI could destabilize the global balance of power.

Outcomes -- Congressional Engagement

In terms of 'outcomes' (as opposed to 'outputs'), we've had some excellent engagement with Congressional offices about their AI governance proposals. In three cases, we were able to suggest edits to an office's proposed legislation that were accepted by the office and that would significantly improve the safety impact of that legislation. For example, our changes would increase the powers of a safety office, or make sure that safety officials remain financially independent, or ensure that new legislative powers are delegated to an appropriate government office that will be able to use those powers effectively. We cannot publicly reveal the exact offices that accepted these changes, because it would be a serious breach of DC norms, but all of the offices were relatively senior Congresspeople who exercise significant influence over Congress's AI policy. The particular bills that we edited did not pass Congress, but this is because almost nothing passed out of the 118th Congress; it was a historically slow and unproductive session.

We also publicly endorsed about a dozen bills that would have a positive impact on AI safety; three of the offices sponsoring these bills included us in their official press releases, citing us by name and including our full quote in support of the bills. The House Committee on Science, Space, and Technology also linked to our endorsement for several of the bills that they recommended. We were literally the only organization listed as an endorser of the Nucleic Acid Screening Act, which is probably the second most important piece of AI bio-risk legislation to be introduced in 2024. I'll talk more about this in my second blog post, but I think it's pretty insane that nobody else stepped up to endorse the Nucleic Acid Screening Act. The bill is relatively uncontroversial; it would have provided $5 million in funding to develop know-your-customer protocols for biolabs who are selling synthetic DNA to the public. The problem this bill was trying to address is that terrorists could soon use AI to generate a protein sequence that could be used as a bioweapon, and then pay to have that sequence mailed to them at a throwaway

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

We also endorsed two Congressional candidates who pledged to promote AI safety if elected. One of those candidates, Rep. Tom Kean, Jr. (R-NJ) later went on to pose sharp questions in a letter to AI CEOs, asking them to follow up on their voluntary security commitments; this letter echoed some of the language that we used in our candidate survey. Finally, we have been actively involved with the FY2026 appropriations process and the FY2026 NDAA, which are essentially the only two bills that Congress must pass every year. Several Congressional offices specifically and proactively invited us to submit proposals for adding AI safety measures to defense funding or to general funding. We have submitted fifteen such proposals and are actively following up with those offices in an effort to get them included in the final legislation.

Taken together, we feel that these achievements show that Congress is taking our concerns seriously and is engaging with us as a useful legislative partner. While it is true that even the most unwelcome advocacy groups will often be able to get a meeting or two with a low-level legislative correspondent (because it is easier for Congress to humor such groups than to openly offend them by refusing to meet), our Congressional engagement goes far beyond being humored or tolerated. We routinely have multiple meetings with higher-level staff like legislative directors and committee counsels, and they routinely reach out to us to ask us for our opinions or advice. We think this is good evidence that these same people would also ask us for advice if they were considering strong AI safety legislation after a warning shot, and we are pleased to have developed these relatively firm political connections in a relatively short period of time.

Finally, our grassroots organizing is beginning to bear fruit. Our first in-person grassroots event was our February 'demo day', which brought 14 university AI safety teams from around the country to Capitol Hill to present interactive demonstrations of AI risks to Congressional staffers, as well as to Congressman Bill Foster, who visited for half an hour and toured the exhibits. We have since expanded this college network to include a total of 25 colleges, and we are working on recruiting similar volunteers from trade associations, such as the national associations of social workers, teachers, psychologists, and firefighters. The plan is to leverage these connections to bring hundreds of volunteers to DC for a national day of action in fall of 2025.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

An important piece of context to use when evaluating CAIP's accomplishments is that we have only recently been able to assemble our complete team. When we launched in June 2023, our 4-person team was composed entirely of volunteers, none of whom had any prior political experience at all. Shortly after we made our first few professional hires in winter 2023, much of the original team of volunteers quit and moved on to other projects, leaving us short-staffed again. In order to hire additional team members, we first had to raise additional funding. I wound up wearing at least three different hats – I was originally hired as CAIP's legislative director, and I remained responsible for writing and editing CAIP's model legislation and reviewing outside legislative proposals to consider them for edits and endorsements. In addition to these duties, I also had to manage our team and recruit and hire new staff. Until October 2024, we had no director of development, so I was also writing all of our fundraising applications. We did not reach our full complement of ten team members until January 2025; on average, we've gotten our work done with about six full-time equivalent positions. For six people working over less than two years with significant interruptions and turnover, I think we've gotten a lot done.

OUR PROPOSED POLICIES

Mandatory Audits for Frontier AI

The primary legislative tool that we recommend for coping with catastrophic risks from AI is a regime of mandatory private audits backed up by a small new government office that has special powers to proactively recruit the tech talent needed to assess whether these audits are accurate and complete, and to block the deployment of a new model whenever they are not satisfied by an audit.

Today, some (but not all) AI developers voluntarily hire third-party evaluators to comment on the catastrophic risks posed by their frontier AI systems. However, these evaluators have relatively little bargaining power, because the company has no legal obligation to conduct an audit at all.

As a result, many companies can and do limit evaluators' degree of access to the final

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

model if an audit appears to be inadequate, we can improve the bargaining power of auditors and give them the power to honestly share all of their concerns about each new frontier AI model, with some degree of confidence that these concerns will force the developer to add additional safeguards or, if necessary to protect public safety, delay or cancel the deployment of their model.

Liability Reform

Our model legislation also includes a liability package, which would explicitly provide for joint and several liability when model developers, plug-in developers, and/or end users all contribute to causing a catastrophic harm. This would make sure that developers can't escape accountability for the foreseeable consequences of their actions by 'delegating' responsibility for their misconduct to under-capitalized business partners.

Our liability rules would also create a public right of action (allowing the federal government to sue on behalf of the public when AI models create major security risks) and a limited private right of action (allowing groups of people to sue for compensation if they have suffered at least $1 billion in tangible damages from an AI disaster.)

We think liability is a good complement to regulation, because the government will never know everything worth knowing about the risks posed by a private company's AI or which strategies would most efficiently mitigate those risks. Giving the company a financial incentive to avoid causing catastrophes -- even if they have a permit to deploy their AI -- will help motivate the company to put its private information to good use.

Hardware Monitoring

Our model legislation also includes minimum cybersecurity requirements and know-your-customer requirements for the largest data centers, as well as a very modest reporting requirement for wholesalers who sell at least 100 specialized AI chips in the same quarter.

Much of the point of these requirements is to help the government build an internal model

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

market for their skill set -- we assume that building a useful picture of how America uses its AI chips will require both good data and competent analysts.

Understanding what's happening in the world of AI will help the government make better policy decisions across the board, and it also creates a greater ability for the government to press the 'off switch' or a 'pause button' if and when it decides that it is no longer tolerable for private companies to be developing superintelligence. We're not advocating for a pause on AI development, but we do think the government should have the capacity to implement a pause if and when it decides to do so. Right now, the government has no list of the country's data centers, and no way of quickly finding out who might be building new data centers. By the time the government realized it wanted to pause AI development and built a brand-new capacity to do so, it might be too late, so we have to start building that capacity now. For further consideration of this point, please see Question 1 in our section on "Responses to Common Policy Objections."

Emergency Powers

Finally, our model legislation includes authorization for the government to take immediate action if they notice that an AI system is in the process of causing catastrophic harm or escaping from human control. Of course, these powers will not be relevant for every type of disempowerment scenario -- some kinds of disempowerment are so gradual that there would be time to pass normal legislation to cope with them, and some superintelligence explosions are so rapid that no realistic government response could hope to affect their trajectory once they have begun.

However, for the intermediate scenarios where a rogue AI or a reckless AI developer needs a few days or a few weeks to complete their attempt to seize power, it is important to make sure that the government can respond within that time frame, and that the response will be proportional, appropriate, pre-planned, and lawful. Under the current state of affairs, where the government has no legal authority to intervene in an AI emergency, we risk both a response that is too weak (because people delay while trying to figure out what they're allowed to do or what they're supposed to do) and a response that is too strong (because the government overreacts and, e.g., nationalizes all advanced AI in response to a marginal

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Further Details

For further details on our model legislation, including the full text of our proposed bill, please see aipolicy.us/work/model.

RESPONSES TO COMMON POLICY OBJECTIONS

1. Why not push for a ban or pause on superintelligence research?

We think the policies we're promoting are already as intense as we can make them while still getting Congress to listen to the content of our arguments. Our experience has been that talking to Congressional staffers about a ban or pause on superintelligence research tends to result in blank stares and a rapid end to the meeting. By contrast, talking about mandatory security audits has yielded respectful interest, invitations to come back and meet with increasingly senior staff, invitations to edit and publicly endorse the office's AI bills, and so on. A global moratorium on superintelligence research would probably be more effective than mandatory security audits at protecting against existential risk if it that moratorium could be enacted and enforced – but we don't see anything that we can do to help make that happen. Rather than make essentially zero progress toward the best possible policy, we'd rather make some progress toward a marginally helpful policy.

We're confident that mandatory security audits would be marginally helpful if passed, because right now there is no federal law of any kind that requires AI developers to guard against existential risks. It is 100% legal to make an AI that might attempt to become superintelligent, take over the world's economy, and use the world's resources for inscrutable artificial goals that are bad for humanity. There is no requirement of any kind to run any kind of safety testing or evaluation or even to assert that an AI is safe before deploying it or publishing its source code.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

of generic discussion about the advantages and disadvantages of AI; they do not contain or outline or even propose any draft legislation that would meaningfully address existential risks, other than to state, in the broadest possible terms, that terrorists and rogue states ought to be prevented from misusing AI to develop dangerous weapons. The words "existential" and "alignment" show up only in two footnotes that were cited to dismiss existential risks as being unsupported by rigorous evidence.

The Congressional task forces that created these reports are no longer meeting (at least not in public), and there are no immediate plans to revive them. When Congress does debate AI legislation, most of that legislation deals with either voluntary standards that companies would be free to ignore (as developed by, e.g., NIST and AISI), with export controls that do not apply to American AI developers, or with subsidies that make it easier for American AI developers to buy chips and build data centers. There are a few minor exceptions, which CAIP has actively supported and lobbied for, but in general, the idea that American AGI developers should be subject to any legally binding requirements of any kind is already at the very edge of the Overton window.

We are trying to move the Overton window toward safety by pushing legislation that is as effective as we can make it without completely losing the attention of our target audience. Because mandatory security audits are far more effective than anything that Congress is currently considering, we think it is useful to push Congress in the direction of taking mandatory security audits seriously.

We also think that our legislation, if passed, would help create the infrastructure and push the government to hire the talent that would be needed to implement a pause in the future, if and when the government decides to do so. While we do not currently advocate for a moratorium, we think the government ought to have the ability to halt certain kinds of AI development if and when (and as soon as) this turns out to be necessary.

If the government decided tomorrow to begin enforcing a moratorium on advanced AI research, it would not know who owns advanced AI chips, who is leasing data centers, or how to determine which models count as advanced. Worse, it would not have nearly enough of the right people in-house to rapidly collect and accurately analyze data on these

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

of the advantages of our model legislation, from the point of view of those who favor a moratorium, is that our legislation would cut out most of that potentially fatal delay by immediately beginning the process of collecting data and hiring analysts.

2. Why not support bills that have a better chance of passing this year, like funding for NIST or NAIRR?

We *do* spend some of our time supporting more moderate legislation, such as bills that would increase funding for the National Institute of Standards and Technology (NIST) to develop voluntary standards on responsible AI, or bills that would have the government pay for a national compute bank (the National AI Research Resource, or NAIRR) that academics could use to conduct independent evaluations of various characteristics of AI models. We've publicly endorsed about a dozen such bills, and organized an open letter in support of them.

These bills already have broad support from many relevant stakeholders, and are reasonably likely to pass as soon as Congress turns its attention to passing substantive legislation. For the last 2.5 years or so, Congress has been distracted by wars, internal leadership struggles, the national budget, and the Presidential election. This has meant that Congress has passed far fewer bills than is historically typical. We think that these distractions are the primary reason why Congress hasn't done more to support NIST and NAIRR, and that there is little we can do to influence the timing of when these bills pass. We don't have the power to radically refocus the topic of Congress's attention, and once Congress returns its attention to the topic of AI policy, we think they will naturally want to pass these bills.

Finally, other organizations that work on AI policy, like Americans for Responsible Innovation (ARI) and the Institute for AI Policy and Strategy (IAPS), have staked out these topics as more central to their brand, and as a result, they are better positioned to make progress on these moderate bills than we are. By contrast, CAIP is one of very few organizations that are able and willing to openly advocate for mandatory AI safety standards. As a 501(c)(4) organization that openly discloses its emphasis on preventing

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

3. If Congress is so slow to act, why should anyone be working with Congress at all? Why not focus on promoting state laws or voluntary standards?

The reason why we focus on mandatory federal legal requirements is that we cannot identify any alternative counterweight that would significantly constrain the existing culture and financial incentives in Silicon Valley. Right now, companies face enormous internal and external pressure to release exciting new products as quickly as possible, both to bolster their financial projections, and to gain market share and public attention at the expense of their competitors. There is a powerful 'race to the bottom' effect that makes it very difficult for any one or two companies to significantly depart from this norm – nobody wants to be left behind.

It is extremely unlikely that a firm set of voluntary standards (as developed by, e.g., NIST or as embodied in Responsible Scaling Policies) will be agreed to by all major AI developers, and it only takes one holdout to develop and release recklessly unsafe AI. Even if all companies publicly say that they agree to a voluntary policy, it is likely that at least one of them will break their promise, perhaps by rationalizing that the situation has changed, or simply because they find the temptation to acquire extremely powerful or lucrative AI more important than the PR benefits of adhering to their public commitments. For example, OpenAI has already reneged on its promise to devote 20% of its compute to safety research, and so far they have not faced any financial consequences for this misbehavior – although many of their safety researchers quit, it does not seem to have affected their ability to recruit capabilities researchers or to raise additional funding.

State laws are likely to be somewhat more effective than voluntary commitments at constraining corporate behavior, but because AI companies enjoy significant federal support as, e.g., a key to promoting American national security and American economic advantage, there is a serious risk that the federal government would overrule state AI safety laws, either directly (via preemption) or indirectly (by putting political pressure on key actors, as California's Congressional delegation did to Governor Newsom after the California legislature passed SB 1047). As a matter of *realpolitik*, state governments are also

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

4. Why would you push the US to unilaterally disarm? Don't we instead need a global treaty regulating AI (or subsidies for US developers) to avoid handing control of the future to China?

We do support binding international agreements to limit the development of advanced and/or unsafe AI, but we believe negotiating these agreements will be easier if the US demonstrates that it is taking AI safety seriously. Our model legislation is intended as a complement to international diplomacy, not as an alternative to it.

We also think that the net impact of our model legislation would be to widen the US technical lead over China, because it would impose minimum cybersecurity standards and know-your-customer requirements that would slow down the rate of tech transfer to China. Because so much Chinese progress on AI comes from reverse engineering and generally following America's lead, we think that cracking down on the flow of technology to China can more than make up for whatever marginal slow down in US progress results from insisting on more careful audits of billion-dollar AI models.

We don't see the international 'arms race' in AI as a prisoner's dilemma where either side can make private gains if it unilaterally defects. Instead, if either the US or China or both build misaligned superintelligence, then the entire world loses; the resulting AI will probably not pursue American or Chinese goals, but will instead pursue artificial goals that are terrible for all humans. As a result, there are no long-term private gains to be had from defecting on AI safety. There is some chance that China is enough of a rational actor to see this and to unilaterally avoid building misaligned superintelligence. If China is behaving responsibly, then it is in America's selfish best interest to also behave responsibly. Call this proposition A. We think the moderate chance that A is true outweighs the very small chance that both B and C are true at the same time, where B is "recklessly pursuing superintelligence without adequate safeguards will nevertheless yield an aligned AI that adopts the values of its creators," and C is "China is so close behind America in the AI race that waiting three extra weeks to fully test an AI model will be the margin of victory that allows China to pass the US."

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

passed some of your legislation by now, or at least found some powerful Congressional sponsors for it?

The timeline for our mission is uncertain and is probably longer than most of the projects that are funded or supported by the effective altruist network. If you fund a research paper, you can reasonably expect the research to be completed within a few months; if you fund bednets, you can reasonably expect the first bednets to be distributed less than a year after you donate, and to get some kind of measurement of public health outcomes not too long after that. By contrast, Congress operates on its own schedule, which is slow, erratic, and not subject to our control. Congress might take no action on a topic for months or even years, no matter how skillfully advocates present their case for new legislation. For that reason, we do not think of our mission in terms of forcing or convincing Congress to pass our model legislation this year or next year or at any particular time. Rather, our goal is to get Congress to feel that our policy proposals are comfortable, serious, and well-understood, so that if and when Congress is motivated to take real action on the existential risks from frontier AI, they will believe that there is a solid plan available for doing so, and they will work with us to improve the quality of their legislation and the odds that the legislation will pass. Given the enormous benefits of even a small chance of preventing misaligned superintelligence, we think this is a worthwhile investment.

Congress's motivation for acting on AI safety might come from having our movement slowly build grassroots support for the cause. We would expect to see gradual progress in organizing and mobilizing mass support from both our own operations and from many other groups such as the AI Safety Awareness Foundation, Pause AI, Encode Justice, and so on. Building a grassroots movement takes a great deal of time, for three reasons. First, you need many thousands of volunteers to be truly effective at a national scale. Second, the movement spreads organically, as individual colleagues and neighbors convince each other to join the team. Any given volunteer for a political cause is not likely to be passionate enough to successfully recruit new volunteers until after they've had several months to learn more about the issue and adopt the issue as part of their identity. Third, getting new people involved often requires some type of large protest or expensive advertising campaign or other dramatic event, which requires significant advance planning.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Alternatively, Congress's motivation to act on AI safety might come from a 'warning shot' or other national disaster that radically changes Congress's sense of what type of AI policy is necessary and desirable. Either way, we believe it is very important to have a well-thought-out set of policy proposals already developed and a set of political connections and relationships in place *before* the critical window for legislative action. Arguably, in spring 2023, one such window of opportunity arose as Congress panicked over warnings from figures like Yoshua Bengio and Sam Altman, but the AI safety movement had neither a concrete policy to offer nor a political advocacy network to support such a policy, so the opportunity was wasted. Our mission is to make sure that the next such opportunity (if any) is not wasted.

OUR TEAM

Building a team to push for AI safety policies in Washington, DC is tricky because the people who are most likely to have a demonstrated track record of support for AI safety (philosophers and computer scientists) are among the least likely to have the skills needed to operate successfully on Capitol Hill (government relations, community organizing, law, and so on.) To solve this problem, we've recruited a mix of both types of people, with the idea of teaching the traditional AI safety people how to advocate effectively, and teaching the experienced DC advocates why AI safety is important and how to explain it. This has taken time – neither advocacy skills nor a deep understanding of existential risks can be acquired overnight – but overall we've made excellent progress at this kind of cross-training and we now have a relatively effective team that both understands why AI needs regulations and has the skills to advocate for those regulations.

Executive Director

I am the executive director of CAIP, and my professional background includes time spent as a product safety litigator, a regulatory compliance counselor, and a data scientist. After graduating from Harvard Law in 2010, I worked for small, plaintiff-side law firms where I advocated for people who had been unfairly injured, fired, evicted, or otherwise mistreated.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

litigation to craft model legislation that eliminates as many of these loopholes as possible, making it easier for future litigators to enforce AI safety laws. As a compliance counselor, I helped homeless shelters cope with the often bewildering array of federal regulations that govern every aspect of their funding and operations. This gave me a visceral sense of why red tape is harmful and how to minimize it. I've tried to use these lessons to refine our model legislation so that it's easy to comply with and will pose minimal burdens for inventors and entrepreneurs.

I have some data science experience designing models that predict housing prices, fines, and electoral results. I'm not going to win any awards for technical innovation, but I understand neural networks well enough to explain them to policymakers, because I've coded them myself. Although I don't have formal leadership experience before taking over at CAIP, I informally led and trained teams of professionals at several of my previous jobs, where I was often the most senior lawyer or senior counselor on my team. I've read several books on management techniques and consulted frequently with an executive coach to improve my leadership skills; although it's always difficult to tell whether employees are trying to flatter you, so far the feedback on my leadership from my team has been overwhelmingly positive.

In terms of my personal background, I've been a rationalist and an effective altruist for the last 15 years.

Government Relations Team

Our government relations team includes Brian Waldrip, Kate Forscey, and Mark Reddish, each of which has at least a decade of experience working on Capitol Hill. Brian is a former legislative director for Congressman John Mica (R-FL) and former professional staff to the US House Committee on Transportation & Infrastructure, where he helped regulate autonomous vehicles and investigate federal disaster responses. He later worked as Head of Congressional Affairs for a drone delivery company. Kate Forscey was a senior tech policy advisor for Congresswoman Anna Eshoo (D-CA) of Silicon Valley, and she also served as the Director of Legislative Affairs for the Wireless Infrastructure Association and as Policy Counsel for Public Knowledge, where she advocated for net neutrality and online

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Policy Team

Our policy team includes Claudia Wilson, Joe Kwon, and Tristan Williams. Claudia Wilson has an MPP from Yale and an economics degree from the University of Melbourne. She has worked for Boston Consulting Group and the OECD. Her financial experience is helping us accurately estimate the direct and indirect compliance costs of our proposed policies, showing that America can significantly improve its AI safety without giving up on innovation. Joe Kwon has a computer science degree from Yale and has four years of post-bachelors' research experience with MIT and UC Berkeley, where he studied interpretability and social cognition. His technical expertise helps us determine appropriate thresholds for the scope of our model legislation. Tristan Williams worked for Conjecture and the Center for AI Safety before joining CAIP as a research fellow.

Communications Team

Our communications team includes Marc Ross, Jakub Kraus, and Ivan Torres. Marc Ross has over 30 years of communications experience, including roles advising the McCain campaign during his 2008 presidential run, the US-China Business Council, and Microsoft's efforts to defend itself against antitrust charges in 2000. He served as a professor of American Politics at George Washington University. He left CAIP earlier this month to write a book on globalization, but the culture he built of fearless, friendly engagement with the press will continue. Jakub Kraus helped develop content for Blue Dot Impact's course on AI Safety Fundamentals, did research for the Center for AI Safety, and has a degree in data science from the University of Michigan. In his capacity as our technical content lead, he founded and maintains our weekly newsletter, podcast, and many of our events. Ivan Torres has over a decade of experience as a community organizer, bringing volunteers together to advocate for a variety of causes and successfully building coalitions among dozens of previously unaffiliated organizations. He has a degree in political philosophy and history from Santa Clara University and is passionate about using the impact of AI to unify activists.

Operations Team

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Analysis, she grappled with the ongoing threat of nuclear war, which informs her work advocating for funding to protect against the existential risks from AI. Kristina Banks has a BA in business administration and a decade of experience handling increasingly responsible roles in finance and operations management; she handles our accounting, HR, and administrative needs.

Personnel Changes

As I write this, some of the people on our team are actively looking for other job opportunities, with my encouragement, because we have a very short runway. As I will discuss in detail in the third post in this series, we were caught by surprise by traditional AI safety funders' refusal to offer us any funding at all in 2025. As a result, even if the community is generous enough to help us continue our operations, we will probably have to replace some of our staff. However, these bios are a fair sample of the flavor and level of expertise that we expect to maintain.

OUR PLAN IF FUNDED

Assuming we can raise enough funding to continue our operations, some of our immediate next steps include following through on our NDAA and appropriations proposals, making sure that the 2025 version of our model legislation is received and understood by staffers on the relevant committees, pushing for re-introduction and approval of some of the useful bipartisan AI governance measures that were approved by the Senate Commerce Committee in 2024, and upgrading our website so that it more clearly showcases our core policy proposals. We also expect a whistleblower protection bill to be introduced in the next few weeks. The bill will cover employees at AI developers and will be well-positioned to gain bipartisan support; we plan to lobby other Congressional offices to keep the bill moving forward.

Over the next 6 to 12 months, we plan to build on our initial grassroots organizing efforts to create an influential national network. We are helping local grassroots chapters build useful skills by organizing AI safety demos in their hometowns, and we will train them to

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

worried about being replaced by AI. We hope to mobilize members of many such professions as people increasingly feel that their livelihoods are under threat. It would take a lot of job loss to add up to total human disempowerment, but, in the meantime, the fact that AI seems to be at least somewhat reducing the demand for labor in many job categories serves as a demonstration that AI has the capabilities to do many of those jobs, which calls attention to the ways that AI can be powerful enough to be extremely dangerous.

In the long-term, our plan for driving AI safety legislation forward depends on either building an effective mass political movement, or on taking advantage of a highly dramatic news cycle, such as an AI-driven disaster or 'warning shot'. Some kinds of disasters would be so terrible that they could themselves lead to a collapse of civilization, but other disasters might be small enough to allow normal government to continue, yet large enough to drive radical political change. Examples of such events in recent history include the 9/11 attacks on the World Trade Center (which led to the previously unthinkable PATRIOT Act), the 2008 financial crisis (which led to the previously implausible Dodd-Frank Act), and the 2020 COVID pandemic (which prompted record-setting levels of public welfare payments). We will do what we can to build our own political capacity and to ally with other advocates to increase our ability to proactively steer Congress's agenda, but we acknowledge that this level of power is still many years away, if we can achieve it at all. It is relatively more likely that we maintain or slightly improve our current level of influence, and then use that influence to improve the quality and odds of passage of AI safety legislation after an appropriate crisis.

OUR FUNDING SITUATION

Our Expenses & Runway

Unfortunately, our funding situation can only be described as "dire," which threatens to undo nearly all of the benefits described earlier in this post. Our bare minimum budget for continuing operations with our current team is about $1.6 million per year, or $133k per month. The vast majority of this funding goes to salaries and health insurance for our 10-

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

We currently have only about $150k in cash reserves, meaning that if we receive no additional funding, we can only operate through the end of May 2025 (this month) before laying off our entire team. The remaining money would go toward paying out accumulated PTO hours, handling unforeseen expenses, and paying our operations manager to arrange for an orderly shut down of our corporation.

No Good Way to Cut Costs

We have considered various plans for reducing our staff size, but none of them make sense in the context of our mission and our strategy. Right now we have two lobbyists covering all 535 members of Congress plus about 250 committees and subcommittees; cutting one of those two lobbyists would mean that there would be important Congressional offices who we don't get to meet with on a regular basis. In order to do their jobs effectively, our lobbyists need support from a public policy expert, a technical policy expert, a grassroots organizer, a grasstops organizer, and a communications director, all of whom need to be coordinated by some type of executive and who need some type of operational support. We have tried doing without some of these roles in the past, and it has hurt us politically; Congressional offices have asked us pointed questions about what part of the public we represent, or who we have talked to about our legislation, and if we cannot answer these questions persuasively because we have no one assigned to complete those tasks, then we lose credibility. We could, if necessary, cut our part-time junior research fellow's position, but this would not significantly affect our total budget. We could cut the director of development position, but this seems short-sighted; having a professional fundraiser on staff usually pays for itself and then some in the long run.

Similarly, we don't see a sustainable plan for cutting salary expenses while still accomplishing our mission. Our current burn rate of $133k/month assumes no raises, no performance bonuses, no adjustments for inflation, and no additional hires beyond a replacement for our communications director, who recently resigned to write a book. Our salaries are already well below market rate for our highly experienced team in a relatively high-cost-of-living area (Washington DC) and reflect everyone's willingness to take a pay cut in order to work on a mission that they believe in. We experimented in 2024 with

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Because our non-salary expenses are less than 10% of our total costs, we don't think we can fix our budget by skimping on, e.g., office supplies or our occasional work travel.

Our Revenue

We have no guaranteed sources of incoming funding. There are five places where we might get significant funding before the end of the calendar year: FLI, SFF, the Siegel Foundation, a private donor who has pledged support for a 501(c)(3) partner organization that we are in the process of spinning up, and a Manifund application we just launched. There are difficulties with all five of these funding sources. FLI's grant is extremely competitive, SFF has recently announced that they will be limiting their grants to 501(c)(4) organizations like ours, the Siegel Foundation is still in the early stages of preliminary talks with us and has not yet even agreed to accept an application from us, the private donor's contribution is held up by indefinite delays at the IRS, and our Manifund application, which is seeking donations for a specific project on reducing chem-bio AI risk, is competing with dozens of projects and we're not sure if we'll be able to reach our goal in time.

Even if all five of these funding opportunities come through toward the high end of what we might reasonably hope for, they would still just barely cover our expenses for the year, and they would leave us with no runway going into 2026. In practice, we have to expect that one or more of these grants will either not come through or will offer us less than we are hoping for. There are also cash flow issues – the FLI grant will not be announced until June, the SFF grant will probably not be paid out until September or October, and any support from the Siegel Foundation would arrive, if at all, toward the end of 2025.

Surprise Budget Deficit

The main reason we have such a large budget deficit is that our reputation in the AI safety funding community seems to have quietly plummeted. In 2024, we had a positive reputation among AI safety donors and evaluators. This led to several donors writing six-figure checks after only a half-hour Zoom call. Our bottleneck appeared to be hiring, rather than funding – I was hiring people as quickly as I could identify people who would be a

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

By contrast, in 2025, CAIP has a lukewarm or neutral reputation, which has led to us being ghosted or zeroed out by many of the parties who we reasonably expected would support us. We have unsuccessfully sought funding from Open Philanthropy, Longview Philanthropy, Macroscopic Ventures, Long-Term Future Fund, Manifund, MIRI, Scott Alexander, and JueYan Zhang. None of these parties had any specific critiques of our work that they were able to share with us – instead, they simply said that there were other funding opportunities that seemed more promising to them, or that politics no longer seems to them like an effective way of promoting AI safety, or that they had heard from third parties that CAIP was not a valuable funding opportunity.

In the second and third blog posts in this series, I will discuss why I believe this lukewarm assessment is mistaken and where the mistake is coming from. For now, I simply want to point out that our urgent request for your money is not based on a budgeting or planning failure. Rather, we chose a size for our team in 2024 that seemed modest and conservative based on the apparent amount of available funding at that time, and then in 2025, our funding suddenly dropped by over 50%, meaning that even our conservative level of spending is unsustainable.

The Bottom Line

What all of this boils down to is that we cannot continue our operations without support from people like you. None of us are independently wealthy, and we cannot afford to carry on our work for free, no matter how important it is. If we cannot get support from institutions, then we will need support from individual donors, or we will have to shut down. We urge you to donate – we can accept contributions of up to $10,000 at https://www.centeraipolicy.org/donate, and if you are considering a larger donation, then you can reach me at jason@aipolicy.us and I will be very happy to discuss it with you. If you prefer to support a specific CAIP project, you can donate to our Manifund project on reducing chem-bio AI risk and help us reach our fundraising goal there.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

The Need for Political Advertising (Post 2 of 7 on AI Governance)

No comments 61 karma

Log in to save where you left off

Mentioned in

122 **The Center for AI Policy Has Shut Down**
61 **The Need for Political Advertising (Post 2 of 7 on AI Governance)**
53 **Mainstream Grantmaking Expertise (Post 7 of 7 on AI Governance)**
45 **'GiveWell for AI Safety': Lessons learned in a week**

Comments 25

Sorted by New & upvoted

LeahC 1y 46 4 2 1

Unfortunately, there was an effective effort to tie AI safety advocacy organizations to their funders in a way that increased risk to any high-profile donors who supported federal policy work. I don't know if this impacted any of your funders' decisions, but the related media coverage could have been cause for concern (ie Politico). Small dollar donations might help balance this.

It seems very likely that the federal government will attempt to override any state AI regulation that gets passed in the next year. Jason put together a strong, experienced team that can navigate the quickly shifting terrain in Washington. Dissolving immediately due to lack of funding would be an unfortunate outcome at a critical time.

Context: I work in government relations on related issues and met Jason at an EAG in 2024. I have not worked with CAIP or pushed for their model legislation, but I respect the team.

MichaelDickens 1y 17 4 4

There is a related concern where most of the big funders either have investments in AI companies, or have close ties to people with investments in AI companies. This biases them toward funding activities that won't slow down AI development. So the more effective an org is at putting the brakes on AGI, the harder a time it will have getting funded.*

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Esben Kran 11mo 13 2 0

To me, it's an interesting decision to pull funding because of this type of coverage. There's a tendency in AIS lobbying to never say what we actually mean to "get in the right rooms" but then when we want to say the thing that matters at a pivotal time, nobody will listen because we got there by being quiet.

Buckling under the pressure of the biggest lobby to ever exist (tech) putting out one or two hit pieces is really unfortunate. Same arguments can be made for why UK AISI and the AI Safety Summits didn't become even bigger; simply that there was no will to continue the lobbying move and everyone was too afraid of reputation.

Happy to hear alternative perspectives, of course.

Neel Nanda 1y 37 9 1

I'm sorry to hear that CAIP is in the situation, and this is not at all my area of expertise/I don't know much about CAIP specifically, so I do not feel qualified to judge this myself.

That said, I will note on the meta level that there is major adverse selection when funding an org in a bad situation that all other major funders have passed on funding, and I would be personally quite hesitant to fund CAIP here without thinking hard about it or getting more info.

Funders typically have more context and private info than me, and with prominent orgs like this there's typically a reason, but funders are strongly disincentived from making the criticism public. In this case, one of the stated reasons CAIP quotes is "had heard from third parties that CAIP was not a valuable funding opportunity" can be a very good reason if the third party is trustworthy and well informed, and often critics would prefer to be anonymous. I would love to hear more about the exact context here, and why CAIP believes they are making a mistake that readers should ignore, to assuage fears of adverse selection

I generally only recommend donating this when you are:

- Confident the opportunity is low downside (which seems false in the context of political advocacy)
- If you have a decent idea of why those funders declined that you disagree with
- Or you think sufficiently little of all mentioned funders (Open Philanthropy, Longview Philanthropy, Macroscopic Ventures, Long-Term Future Fund, Manifund, MIRI, Scott Alexander, and JueYan Zhang) that you don't update much
- You feel you have enough context to make an informed judgement yourself, and grant makers are not meaningfully more well informed than you

I'm skeptical that the reason is really just that it's politically difficult for most funders to fund political

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

PabloAMC 1y 10 0 0

I think this makes sense, but it seems kind of disconnected from the presentation, which seemed to indicate CAIP proposes reasonable policy and has a strong team. Perhaps Jason can clarify why he thinks major donors have passed on this opportunity.

Jason Green-Lowe 1y 25 1 1 2

I wish I could! Unfortunately, despite having several conversations and emails with the various AI safety donors, I'm still confused about why they are declining to fund CAIP. The message I've been getting is that other funding opportunities seem more valuable to them, but I don't know exactly what criteria or measurement system they're using.

At least one major donor said that they were trying to measure counterfactual impact -- something like, try to figure out how much good the laws you're championing would accomplish if they passed, and then ask how close they got to passing. However, I don't understand why this analysis disfavors CAIP. Compared to most other organizations in the space, the laws we're working on are less likely to pass, but would do much more good if they did pass.

Another possible factor is that my co-founder, Thomas Larsen, left CAIP in the spring of 2024, less than a year after starting the organization. As I understand it, Thomas left because he learned that political change is harder than he had initially thought, and because he felt frustrated that CAIP was not powerful enough to accomplish its mission within the short time that he expects we have left before superintelligence is deployed, and because he did not see a good fit between the skills he wanted to use (research, longform writing, forecasting) and CAIP's day-to-day needs.

Thomas's early departure is obviously an important piece of information that weighs against donating to CAIP, but given the context, I don't think it's reasonable for institutional donors to treat it as decisive. I actually agree with Thomas's point that CAIP's mission is very ambitious relative to our resources and that we most likely will not succeed. However, I think it's worth trying anyway, because the stakes are so high that even a small chance of success is very valuable.

James Herbert 1y 21 3 1

If I were to bet on what's happened here, I'd bet it's something to do with Thomas leaving.

Looking at his LinkedIn and his Forum history, he seems very well connected in the field of AI safety.

I suspect it was easy to get funding because people knew and trusted him.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

should never be a reason to defund an otherwise functioning org, doing seemingly crucial work.

If the alternative theory is correct and the hit pieces are to blame, that still seems like an incorrect decision. When you're lobbying for something important you can expect some pushback, that shouldn't be a reason to pull out immediately.

James Herbert 10mo 3 0 0

I agree!

Jason Green-Lowe 10mo 2 1 0

Very well said! I think your first paragraph sums up the most important parts of the story of why CAIP was defunded -- Thomas lost interest, mostly for private reasons, and the x-risk funders relied far too heavily on this data point. In part this is because the x-risk funders appear to lack any kind of formal grantmaking criteria, as I write about in post 7° of this sequence.

Neel Nanda 1y 8 2 1

I do not feel qualified to judge the effectiveness of an advocacy org from the outside - there's a lot of critical information like whether they're offending people, if they're having an impact, whether they're sucking up oxygen from other orgs in the space, if their policy proposals are realistic, if they're making good strategic decisions, etc, that I don't really have the information to evaluate. So it's hard to engage deeply with an org's case for itself, and I default to this kind of high level prior. Like, the funders can also see this strong case and still aren't funding it, so I think my argument stands

Jason Green-Lowe 1y 19 1 1 3

I think these are great criteria, Neel. If one or more of the funders had come to me and said, "Hey, here are some people who you've offended, or here are some people who say you're sucking up their oxygen, or here's why your policy proposals are unrealistic," then I probably would have just accepted their judgment and trusted that the money is better spent elsewhere. Part of why I'm on the forum discussing these issues is that so far, nobody has offered me any details like that; essentially all I have is their bottom-line assessment that CAIP is less valuable than other funding opportunities.

PabloAMC 1y 4 3 2

I think we agree. Thinking out loud: Perhaps the community should consider a way to have a more

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

gergo 1y* 28 6 1 2

Disclaimer: I have a friend working at CAIP.

It's incredibly sad to hear CAIP might needs to shut down. They are one of the very few adults in the room in Washington. If their team doesn't meet the bar of funders, I don't know who else should.

(Of course, it's more complicated than that, see LeahC's comment°).

Marcus Abramovitch 11mo 13 3 1

I really wish donors would write up their reasons for/against funding things like Manifund does. It really doesn't take long, so I don't buy the usual "time constraints" argument, but I think the people at CAIP would waste a lot LESS time if they were told "we don't think your work is good/meets the funding bar. These other orgs are better at it for the cost and would rather scale them up than fund your organization." It'd waste a lot LESS grantee time. They can just add a disclaimer that this isn't any way an abdication of their character or something.

That said, I looked at CAIP, and while you say here what you've accomplished, it seems hard to verify that. I searched the EA forum and the internet, and only 2 people you list in your staff have posted on the forum here, and one more has an account. I didn't do the same for LW or AF, but I expect similar there, maybe a bit more. It's going to be very hard to get 7 figures per year for your non-profit if you aren't very public about what you are doing. I know that takes time, but if you aren't advertising the work you do, people won't know about it.

Here, in this post, you sort of just make a new account and say, "If we don't get $1.6M soon, we are shutting down". That's... a lot of money. It's hard to overstate how hard it is to make and then donate that amount of money per year. But more importantly, it'd be very weird if this worked. There needs to be a lot more constant communication and updates from organizations if they want to continue to receive funding. It begs the question, what does the "Communications Team" do? I know they are probably primarily responsible for communicating with congresspeople, but at least some of that needs to be communicating with the EA/AIS community.

Jason Green-Lowe 11mo 9 1 0 1

Marcus, I agree with your first paragraph -- part of my frustration is that it's not obvious to me that any AI governance donors have more than a vibes-based metric for what it means for "other orgs to be better at it for the cost." Donors did eventually tell us that they thought other orgs were more cost-effective, but they couldn't or wouldn't say which effects they thought were cheaper elsewhere or how they were measuring that.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

However, I object to your implication that until recently we haven't been communicating or sending updates about our work. You may recall that in October 2023, you and I met in person at EA Global Boston, where I pitched you about our work. At your request, I sent you a 15-page document in March 2024 in which I quantitatively estimated the microdooms averted by CAIP's work. You seemed to agree that our work was promising, but at EA Global Boston October 2024, you told me that you were donating only to animal welfare charities. You did not share any criticisms of our work at that time other than that it was not in your preferred cause area.

Our communications team sends out a weekly newsletter to 1,400 subscribers, which is also publicly posted on LinkedIn and X. We also sent out various documents to past and prospective donors, including a detailed annual report that included a dozen examples of us being cited favorably in mainstream media, along with links. We have hosted regular in-person events that are open to the public, where anyone can come and see for themselves which Congressional staffers are coming and speaking with us. If any of the donors had at any time expressed any concerns about our accomplishments being "hard to verify," we would have gladly taken them along to one of our meetings or events or sent them whatever additional information they might have wanted. However, I have never before heard any such complaints.

Marcus Abramovitch 11mo 7 1 0

Yes, I do recall meeting at EAG two years ago and yes, I wasn't giving any money to AI safety back then. I don't have any criticism because I don't know better about what policies to support and didn't want to add noise.

I don't want to be too critical of CAIP here and my original comment might have come off as too harsh. I'm not a big fan of the current funding ecosystem being so private with info and reasoning being so secretive. It's caused us harm before and I expect will cause us harm again and I expect even greater it is causing passive harm where organizations don't improve.

Throwaway81 1y 9 1 1

This contains several inaccuracies and misleading statements that I won't fully enumerate, but at least 2:

1. The Nucleic Acid Synthesis Act does not at all "require biolabs that receive federal funding to confirm the real identity of customers who are buying their synthetic DNA." It empowers NIST to create standards and best practices for screening.
2. It's not the case that "The particular bills that we edited did not pass Congress, but this is because almost nothing passed out of the 118th Congress." Lots of bills passed in the CR and other packages. But it was a historically dysfunctional and slow year

Personal gripe: the model legislation is overly prescriptive in a way that does not future proof the statute

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Your point about the Nucleic Acid Synthesis Act is well-taken; while writing this post, I confused the Nucleic Acid Synthesis Act with Section 4.4(b)(iii) of Biden's 2023 Executive Order, which did have that requirement. I'll correct the error.

We care a lot about future-proofing our legislation. Section 6 of our model legislation takes the unusual step of allowing the AI safety office to modify all of the technical definitions in the statute via regulation, because we know that the paradigms that are current today might be outdated in 2 years and irrelevant in 5. Our bill would also create a Deputy Administrator for Standards whose section's main task would be to keep abreast of "the fast moving nature of AI" and to update the regulatory regime accordingly. If you have specific suggestions for how to make the bill even more future-proof without losing its current efficacy, we'd love to hear them.

Throwaway81 1y 1 0 0

Sure, I'm not going to be able to respond any more to this thread but the methods of governance prescribed themselves are not future proof, as AI governance may need may change as the tech or landscape changes, and the definitions are not future proof.

Holly Elmore 1y 8 3 0

Then there should be future legislation? Why is it on CAIP and this legislation to foresee the entire future? That's a prohibitively high bar for regulation.

Throwaway81 1y 4 1 3

Most legislation is written broadly enough so that it won't have to be repealed because it's rare that legislation is repealed. For example, take their current definition of frontier AI model, which is extremely prescriptive and uses 10^26 in some cases. Usually to be future proof these definitions will be written broadly enough such that the executive can update the technical specifics as the technology advances. Those types of regulations are the sorts of things that would include such details, but not the legislation. I can imagine a future where models are all over 10^26 and meet the other requirements of the model act's definition of frontier AI model. The reason to even govern the frontier in the first place is because you don't know what's coming -- it's not like we know that dangerous capabilities emerge at 10^26 so there's no reason to use this threshold as putting models under regulatory scrutiny forever. Also (eventually) we might achieve some algorithmic efficiency breakthroughs in which the most capable (and therefore dangerous) models don't need as much compute anymore and so might not even qualify as frontier AI models under the Act anymore. So I see the risk of this bill first capturing a bunch of models that it doesn't mean to cover and then possibly not covering any models -- all because it's not written in a future proof way. The bill is written more like a regulation for for the executive level rather than the legislative level.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

statements that I won't fully enumerate" than nobody doing CAIP's work to get the beginnings of safeguards in place.

Jason Green-Lowe 1y 9 1 0

Our model legislation *does* allow the executive to update the technical specifics as the technology advances.

The very first text in the section on rulemaking authority is "The Administrator shall have full power to promulgate rules to carry out this Act in accordance with section 553 of title 5, United States Code. This includes the power to update or modify any of the technical thresholds in Section 3(s) of this Act (including but not limited to the definitions of “high-compute AI developer,” “high-performance AI chip,” and “major AI hardware cluster”) to ensure that these definitions will continue to adequately protect against major security risks despite changes in the technical landscape such as improvements in algorithmic efficiency." This is on page 12 of our bill.

I'm not sure how we could make this clearer, and I think it's unreasonable to attack the model legislation for not having this feature, because it very much does have this feature.

SummaryBot 1y 7 0 0

Executive summary: In this first of a three-part series, Jason Green-Lowe, Executive Director of the Center for AI Policy (CAIP), makes an urgent and detailed appeal for donations to prevent the organization from shutting down within 30 days, arguing that CAIP plays a uniquely valuable role in advocating for strong, targeted federal AI safety legislation through direct Congressional engagement, but has been unexpectedly defunded by major AI safety donors.

Key points:

1. CAIP focuses on passing enforceable AI safety legislation through Congress, aiming to reduce catastrophic risks like bioweapons, intelligence explosions, and loss of human control via targeted tools such as mandatory audits, liability reform, and hardware monitoring.
2. The organization has achieved notable traction despite limited resources, including over 400 Congressional meetings, media recognition, and influence on draft legislation and appropriations processes, establishing credibility and connections with senior policymakers.
3. CAIP’s approach is differentiated by its 501(c)(4) status, direct legislative advocacy, grassroots network, and emphasis on enforceable safety requirements, which it argues are necessary complements to more moderate efforts and international diplomacy.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

5. Green-Lowe argues that CAIP's strategic, incremental approach is politically viable and pragmatically impactful, especially compared to proposals for AI moratoria or purely voluntary standards, which lack traction in Congress.
6. He invites individual donors to step in, offering both general and project-specific funding options, while previewing upcoming posts that will explore broader issues in AI advocacy funding and movement strategy.

This comment was auto-generated by the EA Forum Team. Feel free to point out issues with this summary by replying to the comment, and contact us if you have feedback.

More from Jason Green-Lowe

63 **Shift Resources to Advocacy Now (Post 4 of 7 on AI Governance)**
Jason Green-Lowe · 11mo ago · 38m read 5

61 **The Need for Political Advertising (Post 2 of 7 on AI Governance)**
Jason Green-Lowe · 1y ago · 16m read 0

53 **Mainstream Grantmaking Expertise (Post 7 of 7 on AI Governance)**
Jason Green-Lowe · 10mo ago · 45m read 2

View more

Curated and popular this week

Impact Requires Implementation: What Dispensers for Safe Water Scale-Back Can Teach Us About Monitoring
Tony Senanayake · 13d ago · 17m read · 3

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

The AI people have been right a lot

Dylan Matthews · 4d ago · 7m read · 11

This post was crossposted from Dylan Matthew's blog by the EA Forum team. The author may not see or reply to comments. ---------- Subtitle: Try to keep an open mind as the world gets increasingl…

In Development Magazine Launch

Lauren Gilbert · 4d ago · 1m read · 4

We're announcing a new development magazine: In Development. Our goal is to promote discourse on development without dogma. In other words, we public stories doing deep dives into pressing development issues, backed up by evidence, argument, and experience. The first piece is by Paul…

Recent opportunities in AI safety

2 **Introducing the Labor Automation Forecasting Hub**
postreal · 8h ago · 1m read 0

80 **Help me launch Obsolete: a book aimed at building a new movement for …**
Garrison · 11d ago · 8m read 4

11 **Pivotal Research Fellowship applications are open (deadline May 3)**
Tobias Häberli, Tilman, James Nicholas Bryant, Ariel · 13h ago · 3m read 0

View more

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.