The Need for Political Advertising (Post 2 of 7 on AI Governance)

61

by Jason Green-Lowe May 20 2025

AI safety | Existential risk | Policy | AI governance | Political polarization | Frontpage

In my previous post° in this sequence°, I made a fundraising pitch for my organization, the Center for AI Policy (CAIP), asking readers to donate to us so we can continue lobbying Congress to pass strong AI safety legislation.

Here, in the second post, I want to zoom out from CAIP's particular concerns and argue that the AI safety movement as a whole should be investing far more resources into "political advertising." We need to seek out individual politicians and personally let them know about the risk from misaligned superintelligence and show them what they can do to mitigate that risk.

The core of my argument for this second post is that AI governance ideas are not self-enacting: even the best policy ideas need political champions, or they will not become law. Therefore, to win at AI governance, we have to actively promote and advertise our ideas to political decisionmakers, not just to academic experts. The expected value of the movement's AI governance efforts is something like Q × P, where Q is the quality of our ideas, and P is the power of our advocacy. If P is zero or close to zero, then it doesn't matter how good our ideas are, because our ideas will not steer the future.

If we fail to adequately invest in advertising our ideas, the most likely result is that we'll wind up with a glut of 'orphaned' AI governance ideas that have been thoroughly researched and vetted but that don't impact the real world because they have few or no champions in the political arena. AI governance researchers will discover great strategies for saving the world, but nobody will put them into practice, and misaligned AI will destroy civilization.

Good Technical Alignment Ideas Might Spread Automatically

"AI safety work" can be split into two very broad categories: (1) technical alignment research and (2) AI governance.

By "technical alignment," I mean anything that's directly aimed at making AI systems safer, better understood, more reliable, more compatible with human values, or more controllable.

Technical alignment is not my field, but from what I can see from the outside, AI scientists already want to design maximally safe AI. Once AI developers find out how to make AI safe and compatible with human values at an affordable price, they'll automatically want to do that, because the vast majority of computer scientists prefer to have an AI that will be more reliable (rather than less reliable) at following the spirit of their instructions. Executives and managers at these companies likewise want AIs that are better at following instructions.

The problem to be solved in technical alignment is thus primarily a problem about ***lack of knowledge.***

If you invent a useful new AI safety technique, you might have to do a little bit of work to promote the idea at, e.g., machine learning conferences, but for the most part a good technical idea will spread on its own if you let it. The inventors of the idea have a vested interest in making sure they get credit for their contribution, and other people have a vested interest in staying current in their field and using the best-available techniques. Thus, the hard part in technical safety research is discovering the new ideas, not spreading the new ideas.

I think it would be perfectly reasonable to hire a few people to promote and spread good AI alignment ideas, but it doesn't seem like that's strictly necessary – most of those ideas will probably spread on their own, even without any formal advocacy.

Good AI Governance Ideas Won't Spread Automatically

This is definitely ***not*** the case for AI governance. By "AI governance," I mean anything that's aimed at controlling what kinds of AI hardware and software gets built, who gets access to it, how it will be evaluated, and what should happen if an evaluation raises concerns about an AI system's safety.

The premise of AI governance is that AI developers often have bad private incentives that don't line up well with what would be best for the public, and so we need to find ways to correct those incentives. For example, Meta or DeepMind or OpenAI might want to release their latest AI model immediately, in order to get good PR, boost their quarterly earnings reports, capture market share from a rival, maximize their own chances of gaining power with transformative AI, etc. – even if there's some risk that this model will catastrophically malfunction. A company that releases a risky model captures most of the model's upsides (profit, glory, power, etc.) for itself, but the downsides (bioweapons, loss of human control, etc.) get shared across all of humanity. As a result, those developers will be strongly tempted to release a new AI system even when that would yield negative expected utility for the world.

The most effective way to fix this problem isn't by doing research about the flawed incentives – the basic way in which the incentives are flawed is already reasonably well-understood by all of the relevant parties. ***There's no mystery to be solved here***: AI developers have bad private incentives, and so sometimes they'll act on those incentives and take actions that are bad for the public.

This isn't the kind of problem that can be solved by acquiring ***better knowledge***. Instead, what AI governance activists need most is ***better incentives***.

Right now, society is encouraging AI developers to train and deploy AI models that are unreasonably dangerous, so we need to persuade society to stop doing that. Some amount of research might be useful at the margins in order to better understand how we might best convince a particular institution (e.g., Congress) to take action to change AI developers' incentives...but the research itself won't automatically solve the problem.

Unlike brilliant technical innovations, brilliant policy ideas won't be spontaneously adopted by the relevant decision-makers, for two reasons.

Spreading AI Governance Ideas Requires Overcoming Political Opposition

The first reason is that the policy space is ***hotly contested***. Very few policy ideas are truly win-win for all parties, and the parties that expect to lose will do everything they can to block a policy idea that they see as bad for them. This lines up roughly with "conflict theory" from Scott Alexander's conflict theory vs. mistake theory° dichotomy. I don't know whether the whole world runs on conflict theory, but parts of DC sure do.

The executives at private AI developers like Meta and OpenAI don't actually want to be regulated (at least not in a way that would effectively promote AI safety), because that would mean sharing some of their power with government officials who they don't know, like, or trust. Even if AI safety regulations would be better for society as a whole, they won't be seen as better by Sam Altman or Mark Zuckerberg, and those people have enormous political power that they can and will use to oppose any such regulations.

In addition to controlling armies of skilled and well-connected lobbyists, famous tech CEOs often have higher approval ratings and better name-recognition than Congressional leaders. For example, can you name the current chair of the House Science Committee? If not, then maybe you understand why he's reluctant to antagonize Satya Nadella – that's not actually a political battle that's likely to go well for the Science Committee Chair unless he has a powerful coalition backing him up, or hard evidence of very obvious misconduct from Microsoft.

The kind of evidence he would need to win that contest would have to be much less subtle than "some of Microsoft's products seem dangerous according to experts." Instead, it would have to look something more like "Microsoft is using Windows to secretly steal money from its customers' bank accounts." Otherwise, the media controlled by Microsoft (including Barron's, MSN, and HarperCollins Publishers) runs stories reporting on what an unfair bully the Science Committee Chair is, and those stories probably get repeated by other media outlets who either sympathize with Big Tech (e.g., the Washington Post, which is owned by Jeff Bezos), or who simply don't have time to run an independent investigation.

So, getting a good regulation enacted that will fix AI developers' bad incentives isn't just a matter of discovering a good regulation and gathering enough evidence to prove to a neutral observer that the regulation is good for society – it also requires overcoming political opposition from the stakeholders who have selfish reasons to oppose that regulation.

Being able to win a debate against Big Tech lobbyists is only the first and smallest step in overcoming their political opposition. In addition to convincing politicians that AI safety policies would be good for society, i.e., winning the merits-based argument, we also need to win some purely political arguments. We need to demonstrate that our preferred policies enjoy broad and deep support from the general public, or that they are acceptable to most of the relevant interest groups other than Big Tech, or that they are supported by relevant campaign donors, or by the relevant Congressional committee chairs, or by other key allies in or out of Congress. We need to be able to make the case that our bills or other policy proposals are more likely than not to succeed, because most politicians mostly want to back winning ideas – they don't want to invest their limited resources in supporting an idea (no matter how good it is) that is doomed to stagnate. If you can't show that your idea is likely to attract (or already has) a winning political coalition behind it, then many politicians won't even consider your idea.

This can be very frustrating for political outsiders, but it's not all that different from some of the principles motivating effective altruism: policymakers want to do the most good per unit of their precious time, and in order to do so, they need to filter ideas based on which ones seem like they have a chance of succeeding. Backing a politically doomed idea is typically useless even if, counterfactually, the idea would have helped society if it had been implemented.

Spreading AI Governance Ideas Requires Overcoming Political Inertia

The second reason why good AI governance ideas don't spread on their own is that the policy space is ***very noisy***. There are a lot of different activists and groups at any given time who are all loudly insisting that their issue is the most important and most urgent issue to address. Some of them are wrong, some of them are staking out a position based on subjective values, and some of them are right, but it's not obvious in advance to any particular politician that any particular activist is worthy of their time and consideration. Jack Clark had a great post about this a couple of weeks ago; I largely agree with his assessment.

A typical Congressional staffer is responsible for five or six different ***categories*** of policy: they might, e.g., make recommendations about which regulations to support in health care, energy, transportation, senior citizens' issues, telecommunications, and banking.

Within each of these categories, there are dozens of different sub-topics being considered at any given time – the "tech and telecom" category might include social media, access to fiber optic cables, satellite launches, TV advertising, e-mail spam, funding for national public radio, data privacy, and artificial intelligence. Within a "sub-topic" like artificial intelligence, there are dozens of different policies being proposed by various stakeholders – someone wants subsidies for American chip manufacturing, and someone else wants to combat deepfake pornography, and someone else wants the Social Security Administration to increase its use of AI while processing disability benefit claims, and then over in one corner you have a handful of effective altruists arguing for guardrails against AI's catastrophic risks.

Suppose a staffer works overtime and logs 3,000 hours per year. That leaves 500 hours per policy category (e.g. "tech and telecom"), which gets split across 20 different policy topics (e.g. "AI"), which gets split across another 25 policy proposals (e.g. "mandatory security audits for AI"), leaving a grand total of 1 hour per year to consider any given policy proposal.

During that 1 hour, the staffer won't be able to give the proposal their undivided attention – instead, they're juggling meeting requests, the latest instructions from their Senator, the breaking political news on CNN and X, and an angry email from an important constituent who needs to be pacified. So, more realistically, the staffer might have 20 or 30 minutes per year to think and write about a particular AI safety proposal.

For a long-time reader of this forum, that might be enough time – if you're already familiar with how AI works, why AI is dangerous, what other AI safety policy proposals have been made, etc., then you might be able to evaluate a new proposal in 20 minutes and write an email recommending it to your boss with the remaining 10 minutes. Similarly, if you're one of the handful of staffers in Congress who's an expert on computer science, you might do fine.

On the other hand, if (like the median Congressional staffer) you're a relatively recent political science graduate who's never tried to code and never watched a video about neural networks, 20 minutes probably isn't nearly enough time to make sense of an average AI policy proposal. Much of the jargon will be lost on you, you won't understand the core concerns that are motivating the policy, and if you try asking your coworkers (who are much like you) for a second opinion, they'll probably agree with you that it sounds like science fiction, i.e., speculative and unlikely to be an immediate problem.

So, most likely, you stop reading the policy proposal halfway through, and you move on to another task in your extremely busy day. This isn't the result of nefarious lobbying or corruption or political opposition – it's just that politicians have an enormous number of things that they're asked to think about, and they can't be experts on all of them – so promoting your policy idea to their attention requires active effort and a particular skill set.

Earning the Chance to Persuade

It's not enough to just articulate a good idea and provide evidence that it's good; you also have to do quite a bit of ***political advertising***. Before you can even make the case that your policy would be good for society, you first have to persuade policymakers that it would be a good idea for them personally to dedicate enough time and effort to *evaluating* your policy. This means distilling the arguments behind that policy into something that can be very comfortably absorbed by the extroverts who work in politics: possibly a two-page handout with color graphs and illustrations, but more likely a vivid anecdote or a single punchy analogy that's delivered as part of an in-person conversation.

In order to get a chance to have that in-person conversation, you might need to invite politicians to a happy hour, or hold a fundraiser for them, or hire lobbyists who know them personally, or present on a 'fun' topic that's adjacent to AI safety and that will attract staffers who are bored or tired or depressed or otherwise in need of some relief. You probably have to go to Capitol Hill in person, repeatedly, and not expect politicians to come to your office or download your papers. To craft an effective political advertising strategy, imagine you're an underpaid, overworked staffer who's constantly being bombarded by pleas for attention. You literally don't have time to meet with everyone who wants to talk to you. Which meeting invitations will you accept? Chances are, your choices will mostly be driven by who makes it easy and rewarding for you to talk with them.

That's why, in order to make sure an AI governance idea succeeds, it's not enough to simply research and document the best AI governance ideas – there is also quite a lot of advocacy work that needs to be done to market those ideas to actual policymakers.

The need for this advocacy doesn't line up neatly with either conflict theory or mistake theory – it's a third type of problem. This "theory of political advertising" is different from mistake theory because there's no amount of patiently researching clever ideas that will

cause those ideas to win – in addition to being correct, you also have to be reasonably good at self-promotion. But the theory of political advertising is also different from conflict theory, because there's no evil or selfish opponent to be overcome. We're not necessarily enmeshed in a battle; we just need to take the time to build a political coalition that can promote a bill to the top of Congress's extremely crowded agenda.

We Either Solve Bad Incentives, or We Probably Die

If we ***don't*** successfully promote our ideas to actual policymakers, or if those policymakers don't successfully change the incentives of AI developers, then sooner or later (probably sooner) the developers will recklessly train and deploy misaligned superintelligence, and most of us will die.

The private AI developers are a caught in a race with each other where none of them are able and willing to hold back a model by even three weeks to run all of the tests that they would need to run to make sure that their current models are reasonably safe. Instead, the people evaluating new AI models are forced to rush through their work with early-stage prototypes and inadequate benchmarks, and then when they warn that they can't be confident that the models are safe, the developers publish their models anyway.

There is no plan in place that we can trust to cause the developers to suddenly start exercising more caution as the stakes get progressively higher. Some developers have published "responsible scaling policies," but many of these policies aren't strict enough to reliably prevent a catastrophe, even if they were fully adhered to. For example, Anthropic appears to be conceding that it does not expect to be able to protect its model weights against a concerted hacking attempt by state-level actors, and most companies have not even set any kind of quantitative target, e.g., less than a 1% chance of causing more than 100 deaths per year, let alone provided evidence tending to show that they are on track to meet that target.

Of course, there's no good reason to think that all companies will fully adhere to their responsible scaling policies. The same pressures that are currently causing companies to skimp on the evaluations and safeguards they think are needed at, e.g., ASL-2 will probably also cause companies to skimp on the evaluations and safeguards that they think are needed at ASL-4. Although the risks of more powerful models will be somewhat more

apparent to tech executives, the benefits of an early deployment of those models will also be more tempting; instead of just capturing market share, an executive might believe that early deployment is needed to seize control of the entire future.

On balance, it doesn't seem likely that all companies will spontaneously decide to start being more careful in the future. As I write this, OpenAI is waging a legal campaign to free itself from as many of the constraints of its nonprofit management as possible, suggesting that it wants to pursue profits in the future even more recklessly than it is today. Even if OpenAI sees the light and puts solid guardrails in place, it only takes one company to defect and (accidentally or otherwise) release an extremely dangerous AI.

What this means in practice is that under the *status quo*, most of us probably die shortly after the first AI developer achieves superintelligence. Most types of poorly aligned superintelligence will instrumentally converge° on power-seeking behaviors that prompt them to seize the planet's resources; at a minimum, this means a future where humans are gradually disempowered and locked out of any important decisions. If the alignment is bad enough, then superintelligence could easily lead to total human extinction.

There are no laws in place to prevent this outcome, there is no credible effort to pass such laws in Congress right now, there is no mass boycott or industry consensus that could realistically cause all major AI developers to permanently adopt and rigorously adhere to adequate safety practices, and the handful of attempts to pass state-by-state AI safety legislation are at risk of being preempted by hostile federal legislation.

We have to change this *status quo*.

This means we have to be thinking about, measuring, and promoting AI governance activities based on their chances of changing this *status quo*. We don't have the luxury of choosing the activities that most interest us, or of choosing the activities that feel most comfortable. We have to identify the activities that are most likely to change the *status quo*, because otherwise there's an excellent chance that misaligned superintelligence will cause billions of deaths sometime in the nextdecade.

It's not obvious that we need any more surprising breakthroughs to get those billions of deaths. If we just continue scaling based on present-day trends, the expected outcome of that scaling is that sometime in the next several years, AI starts doing useful computer science research, accelerates its own growth curve, and becomes more powerful than the rest of humanity put together.

It's possible that bottlenecks in data or electricity will slow this process somewhat, but even a slowed-down process will still probably get us to superintelligence in the mid-2030s. Any given constraint can and most likely will be eroded given adequate time and financial incentives. If we "run out of electricity" for scaling in 2027, but scaling remains profitable, then people will probably build more power plants. Those new power plants won't be ready immediately, but they won't take decades to build, either.

If you're a skeptic of scaling laws and you want to reject them in favor of taking the median estimate of thousands of AI researchers, that still gives us a 50% chance of "unaided machines outperforming humans in every possible task" by 2047, i.e., in 22 years. At a minimum, that gets us widespread unemployment that makes the Great Depression look like a hiccup; more plausibly, it leads to billions of deaths. If we don't find a way to change AI developers' incentives before then, most of us are probably going to die young.

The horse population peaked in 1915 because after that, humans and machines were better than horses at every possible task; the 24 million horses we had then are down to only 6 million horses now. If we don't successfully take action to prevent it, humans will suffer the same catastrophic losses. That's the path we're on right now; that's the *status quo*, and that's our default future.

This sequence will continue with a third post estimating our current ratio of researchers to advocates and arguing that this ratio is very poorly optimized for getting policymakers to change the *status quo*. The fourth post will argue that we cannot afford to further delay the shift from research to advocacy, the fifth post will catalog the orphaned AI governance ideas that are currently available and waiting for someone to flesh them out and advertise them, and the sixth post will ask why our movement has been systematically underfunding advocacy and offer suggestions about how to correct that problem.

Previous:

Please Donate to CAIP (Post 1 of 7 on AI Governance)

25 comments 141 karma

Next:

We're Not Advertising Enough (Post 3 of 7 on AI Governance)

3 comments 54 karma

Log in to save where you left off

Mentioned in

122 The Center for AI Policy Has Shut Down

54 We're Not Advertising Enough (Post 3 of 7 on AI Governance)

Comments

No comments on this post yet.
Be the first to respond.

More from Jason Green-Lowe

141 Please Donate to CAIP (Post 1 of 7 on AI Governance)
Jason Green-Lowe · 1y ago · 40m read
25

63 Shift Resources to Advocacy Now (Post 4 of 7 on AI Governance)
Jason Green-Lowe · 11mo ago · 38m read
5

53 Mainstream Grantmaking Expertise (Post 7 of 7 on AI Governance)
Jason Green-Lowe · 10mo ago · 45m read
2

View more

Curated and popular this week

Impact Requires Implementation: What Dispensers for Safe Water Scale-Back Can Teach Us About Monitoring

Tony Senanayake · 13d ago · 17m read · 3

TL;DR Between 2022 and 2023, GiveWell and Open Philanthropy committed over $70 million to Evidence Action's Dispensers for Safe Water programme. An independent survey in 2025 found that in Uganda and Malawi the programme was reaching approximately 70% fewer people with detectable...

The AI people have been right a lot

Dylan Matthews · 4d ago · 7m read · 11

This post was crossposted from Dylan Matthew's blog by the EA Forum team. The author may not see or reply to comments. ---------- Subtitle: Try to keep an open mind as the world gets...

In Development Magazine Launch

Lauren Gilbert · 4d ago · 1m read · 4

We're announcing a new development magazine: In Development. Our goal is to promote discourse on development without dogma. In other words, we public stories doing deep dives into pressing development issues, backed up by evidence, argument, and experience. The first piece is by Paul...

Recent opportunities in AI safety

80 **Help me launch Obsolete: a book aimed at building a new movement for ...**
Garrison · 11d ago · 8m read · 4

11 **Pivotal Research Fellowship applications are open (deadline May 3)**
Tobias Häberli, Tilman, James Nicholas Bryant, Ariel · 13h ago · 3m read · 0

12 **AI Can Already Evade DNA Synthesis Screening. Congress's New Bil...**
Sophie Kim · 22d ago · 7m read · 0

View more