AN ACTIVIST VIEW OF AI GOVERNANCE

We're Not Advertising Enough (Post 3 of 7 on AI Governance)

54

by Jason Green-Lowe May 22 2025

AI safety | AI governance | Frontpage

In my previous post° in this series, I explained why we urgently need to change AI developers' incentives: if we allow the *status quo* to continue, then an AI developer will recklessly deploy misaligned superintelligence, which is likely to permanently disempower humanity and cause billions of deaths. AI governance research can potentially be helpful in changing this *status quo*, but only if it's paired with plenty of political advertising – research by itself doesn't automatically convince any of the people who have the power to rein in AI developers.

Executive Summary

Here, in the third post, I want to make it clear that we are not doing nearly enough political advertising to successfully change the *status quo*. By my estimate, we have at least 3 governance researchers for every governance advocate. This means that we are predictably generating more good ideas than we can hope to share with decision-makers. This is an unforced error on our part that we should correct.

A major reason why this 3:1 ratio is poorly suited to changing policymakers' minds is that you almost always want to have more staff working on a 'central' activity that can ***directly*** achieve your goals than on a 'supportive' activity that only ***indirectly*** supports those goals. This is mostly a matter of arithmetic: the benefit of an indirect support worker is multiplied by the size of your central team, so if your central team is very small, then it's hard for support staff to do much good. You don't need a full-time accountant to support a

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your cookie settings for more control, or review our cookie policy here.

Reject Accept all

accountant's talents to pay off. If you had three accountants for every taco truck, that would be an even worse use of resources.

The directness or indirectness of work isn't a binary variable – rather, it's possible for work to be one, two, or several layers removed from its ultimate goal. If your theory of change relies on person A supporting person B, who will support person C, who will support person D, then at ***each*** stage of that operation, the 'inner' team has to be large enough relative to the 'outer' team in order to justify investing in a specialist – meaning that the core team should be truly massive before you start investing in specialists who are three or more levels removed.

Unfortunately, most of the research being done in the AI governance community today is very abstract, and is several layers removed from the core goal of changing AI developers' incentives. The research typically does not even comment on any particular law or regulation – instead, most of the research tries to make general predictions about the future or to weigh the pros and cons of various types of policies. This kind of academic research could be useful to policy wonks who do offer constructive criticism about a specific policy, which could in turn be useful to the people trying to draft better versions of that policy, which could in turn be useful to the people selling and advertising that policy on Capitol Hill – but the effects are very indirect and heavily diluted.

We could afford to support a few pure academic researchers despite this dilution if we had many thousands of people working on oral advocacy – but by my estimate, there are less than 60 full-time advocates working on AI governance in the US. As a result, I think we should be looking to transition as many academic researchers as possible into advocacy roles and help the rest of them re-focus their research on topics that will be of direct relevance to advocacy.

There are 3 Researchers for Every Advocate

Methodology

To get my estimate that there are three researchers for every advocate, I looked at the

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

LinkedIn page or their website bio that they work on AI governance, then they're not included in the estimate.

I have been working and networking in the US AI Governance field for the last three years and I have been proactively trying to find out who else is in this field with me, so I expect that my list of organizations is reasonably (although not perfectly) complete.

I count anyone whose job title is "researcher," "fellow," "policy analyst," or a similar variant as a researcher, and I count anyone whose job title is "government relations associate," "communications director," "engagement specialist," etc., as an advocate. For managers, directors, and other supervisors, I try to estimate how many of their staff are working as either researchers or advocates; if necessary, I split their time and assume they are working half-time as a researcher and half-time as an advocate. Similarly, I try to divide up support staff, operations staff, and fundraising staff based on the relative proportion of research and advocacy that the institution is engaged in. I do not count government employees (because they are not supposed to be advocating for any particular policies), or professors who are not explicitly affiliated with a research institute, policy institute, or advocacy group (because they are presumably focused mostly on teaching and on purely academic scholarship). I *do* count PhD students who describe themselves as employed full-time in a particular research job, but I do *not* count other full-time students (because they presumably are focused more on learning the fundamentals of their field than on conducting original research or advocacy).

I do not count people working on technical AI alignment, because I am focused on AI governance. Similarly, I do not count people working for institutions that are focused on field-building, recruiting, fundraising, or organizing conferences – such work is valuable, but it is not the focus of my argument in this blog post.

To determine whether a person works on "US" issues, I rely on where their institution is based, unless their job title explicitly indicates otherwise. For example, the Future of Life Institute's primary address is in the United States, but I have excluded the FLI researchers whose job title is "EU Fellow." Similarly, the Future Society could be thought of as a primarily European organization, but I have included a few of their researchers who clearly focus on America, e.g., their Director of US AI Governance. I suspect that this

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

full-time equivalent (FTE) research positions, Apollo Research employs 20, and the Center for Long-Term Risk employs 12. I am not aware of any similarly sized group of UK-based AI governance advocates who lobby the US government.

To determine whether a person works on "AI Governance," I have reviewed their website or LinkedIn biographies and made a snap judgment. In some cases, I counted a person as "half-time" because it appeared that they spent part of their time working on AI Governance and part of their time working on other topics. Not everyone on the list would identify as a 'doomer' or even as someone who "works on AI safety" – the list intentionally includes people with a broad range of opinions and backgrounds. However, I have tried to limit the list to only those people who appear (based on their published work and their public bios) to support some type of proactive AI governance. I therefore do not include people who appear to believe that decisions about AI should be left entirely to individual private companies with no form of coordination (voluntary or otherwise), and I do not include AI accelerationists.

In the interest of making a conservative estimate, I did not include any of Epoch AI's 22 staff as either researchers or advocates. Although much of their research might wind up cited in other AI governance work, and although their website shows that they received $6 million in funding from Open Philanthropy, it's not obvious that their primary function is to conduct AI governance research – most of their research aims to predict factual trends about AI capabilities, rather than to analyze how those capabilities might best be governed.

I do not count the 10 people working for my own organization, the Center for AI Policy (CAIP), because, as I have explained in the first post in this series, we have been denied funding by most of AI safety's institutional donors, and it is very likely that all of our staff will soon need to find jobs in other fields.

Results

Using this methodology, I estimate that there are currently 202.5 full-time equivalent US AI governance researchers, but only 55.75 full-time equivalent US AI governance advocates. This is a ratio of approximately 3.6 researchers for every advocate.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

entirely overlooked some people. I am confident that the true ratio (using my methodology) is between 3 and 4 researchers per advocate.

A Note on Privacy

Out of an abundance of caution, I have not actually included the list of the particular advocates, researchers, and organizations in this public post.

The reason I am not publishing the list is because many elements of the AI policy community strongly value a certain kind of discretion or anonymity, for various reasons. Some organizations or programs that appear to me to be focused on catastrophic risks from AI take pains to maintain a public presentation as 'mainstream' or 'establishment' think tanks who cover a broad range of issues. Other organizations simply want to avoid controversy or avoid being attacked by, e.g., Senator Ted Cruz.

I believe this effort to maintain a low profile is misguided. It seems to me that portraying an AI safety organization as mainstream is likely to lead to bad results no matter whether your portrayal succeeds or fails. If you successfully convince politicians that you are a mainstream national security think tank that only incidentally covers AI policy, then those politicians have no particular reason to call you for advice when there is an AI emergency or an AI policy opportunity; you are likely to blend into the thousands of other policy professionals in DC, most of which have better connections or better credentials. Likewise, if you fail to convince politicians that you are a mainstream national security think tank, then those politicians are very unlikely to call you for advice, because they have no reason to trust you; they can see that you're misrepresenting your organization, and that you're not even particularly skilled at deception.

I don't understand how blending into the background is supposed to help. In order to spread ideas and convince others that those ideas have merit, you have to tell other people about those ideas. If you communicate clearly enough for third parties to understand why they should agree with you, then you will also communicate clearly enough for them to figure out that you're working on AI safety.

Moreover, I do not believe that sharing the list I used to create the estimate used here

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Nevertheless, the primary point I am trying to make is that the community should be shifting resources from research to advocacy, and I am concerned that if I publish a list of people working on AI governance, then the mere fact that the list was published will be a serious distraction from my main point. Rather than fight two different battles, I will keep the list semi-private.

If you have a particular concern about my methodology or results, and you need to see the list in order to resolve that concern, please send me an email at jason@aipolicy.us. In that email, please use your real name and explain what your question or objection is. If I do not already know you, then please schedule a video call at https://calendly.com/jason-aipolicy/15-minute-meeting so that we can establish trust. On a case-by-case basis, I will consider sharing the list with specific people who have a legitimate reason to see it and who I trust not to further distribute or publish the list.

A Note on Activist Researchers

Some people whose job title is "researcher" may also spend some time presenting their findings to policymakers, and in that sense, they may be doing their own advocacy. However, the impact of the advocacy conducted by researchers is diluted in several different ways.

First, and most obviously, part-time advocacy involves fewer hours per person. If 10 researchers all spend 10% of their time on advocacy, the benefit is probably no higher than the benefit you'd get from 1 advocate working full-time on advocacy.

Second, any part-time activity is typically less effective *per hour* than a full-time activity, because if you're only spending a few hours a week on something, then you don't get as much of a chance to improve your skills at it. If you're a full-time advocate, you can repeatedly practice your 'elevator pitch' for the most important pieces of legislation, you can get feedback (from politicians' facial expressions) about which versions of that pitch are most effective, and you can refine your pitch based on what you learn. You'll also just get more comfortable with DC jargon and norms and fashion as you spend hundreds of hours there, which will make you a more natural and therefore a more effective speaker. By

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Similarly, if you work full-time, you can invest in resources that require economies of scale, like BGov (an expensive service that gives you contact info for hundreds of Congressional staffers, organized by committee), or Politico Pro (which gives you detailed insider gossip about what's happening in DC), or a physical office near Capitol Hill that signals to lawmakers that you're a serious player in the game. By contrast, if you're only doing advocacy part-time, you might be flying in from out of state for short periods and taking whatever meetings you can get while recovering from jet lag.

The *kind* of people who are doing research 90% of the time probably don't have a set of talents and skills that are optimized for advocacy. Although anyone can learn to be a better advocate, the ***best*** advocates are typically extroverts who enjoy making and maintaining connections with lots of strangers and who can quickly and accurately figure out what those people are feeling and how to capitalize on those feelings. These are not traits that are commonly found among people who have chosen to do advanced intellectual research, which typically requires many hours of quietly sitting alone and reading and writing.

Finally, I suspect at least some researchers are overestimating how much time they're spending on activities that have a meaningful impact on policy. Publishing a paper on your organization's website that includes recommendations for policymakers is typically not an effective way of actually reaching those policymakers, especially if you don't have any statistics on how many times the paper is being downloaded or who is downloading it. Submitting a response to a federal agency's Request for Information is a little more effective, because then at least you know that one of that agency's staffers will have to skim your paper for long enough to catalog it, but you're still not guaranteed to reach anyone who has any decision-making authority. Showing up at a policy conference that includes some Congressional staffers or agency officials can be a good way to network, but if you are not personally interacting with those officials at the conference, then the benefit is limited.

If you're working in research and you think your research is impacting policy makers, it's worth asking: which government officials have I personally talked to this month? Did they say that they changed their mind about anything? Did they make any new promises to vote for or co-sponsor or edit or introduce any laws or regulations? Did they say that they

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

I do not want to entirely dismiss the advocacy being done by researchers; I'm sure it has ***some*** positive impact – but because it's diluted in all these different ways, I don't think it substantially changes my point that we have more researchers than advocates. Even if we make the somewhat generous assumptions that all 200 researchers are spending 10% of their time on advocacy, and that their advocacy is half as effective per hour as the work done by full-time professional advocates, that would add another 10 advocates to the pool, and it would subtract 20 researchers from the pool, meaning the effective ratio would be roughly 180 researchers to 65 advocates – still almost 3:1.

Put the Largest Team on the Most Central Problem

There is no good reason to have three times as many researchers as we do advocates. On the contrary, we should have more advocates than researchers, because advocacy is a more central problem than research in AI governance. The reason why advocacy is more central is that (as I've been arguing throughout this sequence) AI governance isn't a problem that can be solved with better ***knowledge***; we have to actually go and change AI developers' ***incentives***, or else we are probably going to die when the developers release misaligned superintelligence.

To the extent that academic research on AI governance steers the future at all, it does so by making AI governance advocates more effective – and in that sense, the benefit of academic research is indirect, not central.

Multiplier Effects and Staffing Ratios

The main reason why you want a relatively large group of people working on a central problem and a relatively small group of people who are indirectly supporting that first group is that the benefit of each indirect "support worker" is ***multiplied*** by the number of direct "core workers" on your team. The larger your core team, the more good each support worker does; the smaller your core team, the less good each support worker does.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

surveyor, so that you can lay the buildings out neatly and make the most efficient possible use of the land.

Alternatively, suppose your construction crew is only 2 workers: you and a friend. In that situation, it'd be crazy to assign your friend to do full-time surveying, because your limiting factor is not the land, it's your time: you'll be lucky to build even a couple of houses in one year with just the two of you working, so you'd better focus on putting up drywall and doors and windows and so on, even if that means your houses aren't neatly arranged in a perfect grid.

In both examples, surveying is a net positive activity – but it's a thousand times less effective in the second example, because there are a thousand times fewer carpenters to benefit from it.

Adding more surveyors to the second example would not be expected to add more than a tiny bit of marginal value. If you have one carpenter and five surveyors, the carpenter will be able to work on an incredibly precise and accurate grid...but there's still only one carpenter and they just can't swing their hammer enough times to put up more than a couple of homes in one year, no matter how well-organized the grid is. In this situation, if at all possible, you would want to reassign one of the surveyors to work as a carpenter, even if they were a professional surveyor who was brilliant at their job and who didn't know much about carpentry.

If they needed to spend six months learning carpentry before they knew enough to start doing work on the doorframes, that would be a perfectly reasonable trade-off: their six months working as the second carpenter will get more homes built than their twelve months working as the sixth surveyor.

For similar reasons, if your goal is to distribute good music as widely as possible, you don't want one pianist with four piano tuners. If your goal is to rapidly respond to fires, you don't want one firefighter supported by ten telephone dispatchers. If your goal is to run a successful restaurant, you don't want two cooks supported by a team of thirty customer service reps handing out customer satisfaction surveys.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

firefighters. It's not about which job is more important; it's about what staffing ratio makes sense given the marginal gains from adding an additional person in each role.

The marginal gains from adding another core worker are almost always moderate; they scale approximately linearly until you saturate the field. If you have 500 carpenters, you can build houses in a large neighborhood about five times as fast as if you had only 100 carpenters.

By contrast, the marginal gains from adding another person who makes core workers more efficient vary quite a lot depending on staffing ratios. The marginal value of adding another support worker is very high when the core team is large relative to the support team, and very low when the core team is small relative to the support team. If you have 100 carpenters and 0 surveyors, adding your first surveyor adds lots of marginal value. If you have 10 carpenters and 10 surveyors, adding an 11th surveyor adds very little marginal value.

Whichever activity is closest to your core goal almost always needs to have ***more*** people working in it than the activities that support that goal by making it more efficient, because otherwise the arithmetic starts working against you. It mostly doesn't matter how efficient the core activity is if barely anyone is performing that core activity.

How Many Levels Removed Are You?

The problem of indirect support teams needing to be small relative to the core teams they're supporting grows even more acute when there are ***multiple layers of separation*** between the support team and the ultimate goal.

Marathons and S'mores

To see what I mean by "multiple layers of separation," consider someone whose ultimate goal is to run a marathon. A level 0 activity for this person might be going on a training run – the training run directly contributes to your ability to run the marathon, so it's zero levels removed. Buying sneakers is a level 1 activity for this goal – it's one level removed from your goal of running the marathon. With a better pair of sneakers, you can train

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

this is only useful if you use what you learn to buy the best sneakers, and then use those sneakers to go on harder training runs. A level 3 activity is installing a new web browser that loads reviews faster – with a better browser, you can learn about the best sneakers more efficiently, but this is only useful if you actually use the browser to read several reviews, which is only useful if you then go out and buy the best sneakers, which is only useful if you then go on harder training runs. A level 4 activity is teaching yourself HTML – if you learn more about how browsers work, you'll be able to do a better job of selecting the best browser, which will let you load reviews faster, which will let you pick out better sneakers, which will let you go on harder training runs, which will help you run the marathon.

Hopefully it seems obvious that learning HTML is a terrible way of preparing to run a marathon – if your actual goal is to finish the marathon, then you would be much better off skipping the HTML lessons and reallocating that time to jogging practice or stretching or even just getting more sleep. Because learning HTML is so many layers removed from your actual goal, it's not an efficient way of accomplishing that goal.

Even if you were the coach of a large high school track team, you probably still wouldn't assign even one student to spend every practice learning HTML. The multiplier effect of learning HTML is so small (because it's so far removed from the core activity of running) that you would need to be able to reap that effect across hundreds of thousands of runners before that kind of specialization started to make sense. If you have fewer people than that, you would want to do things that are more closely tied to the core activity.

Similarly, if your ultimate goal is to roast some s'mores, the level 0 activity could be lighting a fire. To prepare to light the fire, you might perform the level 1 activity of gathering firewood. To help you gather firewood more efficiently, you might perform the level 2 activity of sharpening the axe that you use to chop wood. To help you sharpen the axe more efficiently, you might perform the level 3 activity of talking to friends and discussing which blade-sharpening techniques work best. However, if most of your friends aren't interested in this topic, then you might perform the level 4 activity of travelling to a blade-sharpening convention so that you can meet more people who like to talk about blade-sharpening.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

and in any case you can probably just buy some firewood at your local hardware store for less money than it would cost to go to the convention. Meeting blade-sharpening experts is so far removed from the core task of roasting marshmallows that it's not a reasonable way of pursuing that goal.

These are silly examples, but as I'll show in the rest of this subsection, they follow approximately the same structure as what we're doing right now in AI governance. Writing an academic paper about AI governance is about as far removed from the task of changing Sam Altman's incentives as learning HTML is from running a marathon. It's not that the paper can't be helpful – it's that the help is so far removed from the core activity of fixing AI developers' incentives that it would have to be spread across thousands of advocates before it was more cost-effective than just directly working on advocacy – and because we don't ***have*** thousands of advocates, we'd be better off focusing most of our efforts on lower-level tasks.

0 Levels Removed: Government.

The "level 0" task for AI governance mostly happens *inside* the government. For example, passing a law that penalizes reckless developers is zero levels removed – it directly changes AI developers' incentives. The people who can do level 0 work include legislators, regulators, judges, prosecutors, and police. To the extent that AI developers need insurance or private accreditation, this group might also include insurers and auditors. Similarly, if there are strong enough norms in place about following industry standards, then the people who sit on a committee that writes industrial standards (e.g., at IEEE, or at NIST) might also belong in this group.

Right now, almost no level 0 work is happening, because there are very few laws that require any particular behavior from AI developers. The closest recent example I can think of is Congress's passage of the Take It Down Act, although this deals only with one particular harm from AI, and not with existential risks.

1 Level Removed: Politics.

Activists and advocates who are trying to convince the government to change the rules

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

how many people are protesting outside his sound-proofed office. However, protestors can sometimes convince legislators to write new laws, which *would* change Sam Altman's incentives.

Examples of level 1 work include sending your Congressperson questions to ask at a hearing, meeting with agency officials, giving press conferences, writing open letters, filing litigation, campaigning for Senator Hawley or against Senator Cruz, and organizing boycotts. Right now there is not much level 1 work happening, because the AI safety movement employs very few professional advocates. One recent example that does qualify as level 1 work is the letter to the California Attorney General urging him to block OpenAI's conversion to a for-profit company. The letter prominently features a specific request for a specific politician to take a specific action, which, if agreed to, would immediately change an AI developer's incentives.

This is the kind of work that I would like most of the AI governance community to be doing.

2 Levels Removed: Drafting.

Policy wonks who are drafting specific, concrete policies for the activists to promote are doing "level 2" work, in the sense that this work is two levels removed from changing the incentives of AI laboratories. This could include writing laws, writing regulations, writing treaties, writing contracts, writing industry standards, or writing corporate governance documents. By itself, writing up the text of model legislation doesn't do anything to change Meta's incentives, but it's possible that a political advocate will be able to use that model legislation to make a more effective pitch to a Congressperson, who will in turn be more motivated to try to pass a law that *would* change Meta's incentives. I'll cover specific examples of the documents that need to be written in post 5 of this sequence.

Other activities that belong in this category include opposition research, polling, and legal surveys. To do their jobs effectively, advocates may want juicy gossip about the weaknesses of their political opponents, so that those opponents are easier to discredit. For example, if you can show that a supposedly independent medical clinic is actually funded by the tobacco industry, then politicians will be less likely to believe that clinic when it

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Finally, if, as an advocate, you know what other key players are doing in your field, then it's easier for you to convince lawmakers to trust you and to follow your advice, because you seem well-informed about things that politicians care about. A modest amount of background research on, e.g., what state legislatures are doing on AI governance, or what federal agencies are doing on AI governance, can enhance advocates' persuasive power, and is therefore only two levels removed from the core task of changing AI companies' incentives. A good example of this kind of legal survey is Convergence Analysis's 2024 State of the AI Regulatory Landscape.

It makes sense for a portion of the AI governance community to do this kind of work – 10%, or 20%, or something along those lines. We don't need as many drafters as we do advocates, but the drafters are still valuable.

3 Levels Removed: Commentary.

Any document that's offering any kind of commentary on someone else's policy proposal is at best "level 3" work. This includes papers that offer constructive criticism about an existing idea, that evaluate the likely impact of one or more policy proposals, and that muster arguments for or against a particular policy proposal. If you're arguing about how many FLOPs a threshold should be set at, that's three levels removed from the core work of changing AI companies' incentives. This work can only be effective if there's already a reasonably well-drafted concrete policy proposal in play, which is already being promoted by reasonably competent advocates, who are already being listened to by reasonably powerful politicians. If any of those steps are missing, then the commentary has no measurable impact on the real world.

Similarly, if the main contribution of a document is to suggest a new policy idea that someone else might someday want to flesh out into a concrete policy proposal, then that's three levels removed from the core task of influencing AI companies' incentives. Until someone actually goes and does the work of drafting up that policy proposal into something reasonably concrete, it doesn't have any impact at all on which models DeepMind will deploy.

For example, inventing the idea of "compute monitoring" or explaining why compute

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

pass it into law. However, if nobody ever writes up a sample compute monitoring bill, or at least fleshes the idea out enough to show how precisely how it could work, then compute monitoring remains stuck at the idea stage. Politicians are too busy to take an abstract idea and turn it into a law all by themselves. If nobody does that work for them, then nobody will actually go and monitor Amazon's chips, and Amazon will remain free to rent or sell their compute to anyone who can pay, even if compute monitoring is a brilliant idea with a mountain of evidence behind it.

If there are only 60 AI governance advocates, it's not obvious that any of them should be working full-time on creating this kind of commentary – whatever marginal gains can be had from improving the quality of theoretical policies (which might get drafted someday) are heavily outweighed by the need to go and actually draft any policies at all (right now) and then show those drafts to decision-makers (next month). I expect this point will seem more compelling after you read post 5 in this sequence, which catalogs all of the 'orphaned' policy ideas that have been sitting on arXiv for years, patiently waiting for someone to bring them to Capitol Hill for the first time.

Another reason not to stress over making marginal improvements to as-yet-undrafted policies is that whether any given policy gets passed has more to do with the news cycle and the force of the advocacy behind it than the exact quality of the policy. This is sometimes called the "multiple streams framework" or just "waiting for the policy window to open." If you have a strong political coalition supporting a mediocre policy during a month when the news won't stop talking about your issue, the mediocre policy is pretty likely to pass; if you have a mediocre political coalition supporting an excellent policy during a month when the news is ignoring you, the excellent policy is very likely to fail. If this doesn't seem obviously true, it might be worth reflecting on all of the mediocre policies that have already become law. How do you suppose they got passed?

Because of these dynamics, we're much more likely to succeed by taking a higher number of reasonable shots on goal than by trying to take one perfect shot. Besides, the quality of even our roughest policies is still much better than the policies we'll get under the *status quo*. Unless and until we get something done in government, the default policy is that companies keep recklessly developing and scaling unsafe models until one of them kills

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Many of the documents being created by AI governance researchers are still yet further removed from the core task of altering AI developers' incentives. They don't offer any significant new commentary on any particular policy proposal – instead, they consider abstract or academic policy questions, or they seek to categorize or organize existing material.

For example, "Considerations for Governing Open Foundation Models" is not advocating on behalf of any particular policy tool for coping with open source models, nor is it directly making suggestions about how to improve a policy on open source models. Instead, it catalogs some of the risks and benefits of open source models. This could be indirectly useful if, e.g., someone were trying to offer constructive criticism about an open source policy proposal and they weren't sure what categories or concepts they should be referencing – but if nobody is writing that particular piece of constructive criticism, or if the underlying open source policy proposal hasn't yet been fleshed out, or if no advocate is presenting that proposal to politicians, or if the politicians don't vote the proposal into law, then this research does not have the power to steer us into a different future.

Similarly, "An Overview of Catastrophic AI Risks" is not primarily aimed at supporting any particular policy proposal. Instead, the paper says that its "goal is to foster a comprehensive understanding of these risks and inspire collective and proactive efforts to ensure that AIs are developed and deployed in a safe manner." Most of the 55-page paper discusses how AI development could cause harm, with only about 6 pages going toward "suggestions" about how to respond to these risks. Because these suggestions cover several different policy ideas for each of four different types of risks, it is difficult for these suggestions to cover new ground or go into substantive detail. For instance, the suggestion on "legal liability for developers of general-purpose AIs" does not endorse any particular liability regime. Likewise, the suggestion about "adversarially robust anomaly detection" notes that "AIs could enable anomaly detection techniques that could be used for the detection of unusual behavior," but does not say anything about how to measure the effectiveness of such techniques or how to hold AI developers accountable for implementing such techniques.

As a result, the paper is what it claims to be: an overview. Policy wonks might find the

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

it's very likely that this chain of cause and effect will break somewhere in the middle. Even if the chain holds, the research probably won't drive a significant amount of policy change, because it's not obvious how any particular advocate is supposed to use this research. The overview is too long to be of interest to most policymakers, yet too vague to contain concrete new policies.

Another type of research paper tries to estimate a particular fact or make a specific prediction about the future, even if this prediction does not directly bear on any particular policy. For example, Daron Acemoglu's research predicts that AI technologies will increase total factor productivity by about 0.66% over the next 10 years. It's not immediately obvious how (if at all) this finding connects to any particular AI governance effort. Professor Acemoglu suggests that economic results will be better if "AI is used for generating new tasks for workers," but it is not clear who could ensure that this happens or how they would do so. As a result, his research is at least four levels removed from the core task of changing AI developers' incentives.

Some academic research tries to make qualitative predictions, rather than quantitative predictions. Earlier this month, RAND published the 73-page paper "On the Extinction Risk from Artificial Intelligence." The goal of this paper is not to pin down any particular fact or statistic, but rather to "see if we can describe a scenario in which plausible extrapolations from current circumstances could lead to an outcome that meets our definition of an extinction threat." The paper considers AI-related threats from nuclear war, bioweapons, and geoengineering, and eventually determines that "an extinction threat is at least possible" from each of these types of threats, but that AI would have to overcome significant barriers to literally trigger human extinction.

The paper concludes by recommending several avenues for further research, including research into cybersecurity, research aimed at identifying relevant "risk indicators," research into new types of threats, and research aimed at identifying "decision triggers." Although the paper briefly suggests that humanity should improve its general resilience by, e.g., reducing nuclear weapons stockpiles and preparing for future pandemics, the paper does not say who specifically should pursue these goals or how they should accomplish them. It is not clear that the paper is even attempting to recommend that any

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

One way of interpreting the RAND paper is as a call for national security experts to pay more attention to existential risks from AI. On this theory, the point of the paper is not to immediately advocate for any particular policy change, but rather to shift politically valuable experts into 'our' camp, so that future policy changes will be easier.

There are two problems with this theory. First, the paper seems to be aimed primarily at national security *researchers*, rather than national security *policymakers*. The 'action item' mentioned most often in the paper is the need for more research to be done on particular topics. There are no specific recommendations made in the paper that Marco Rubio or Steven Miller or Dan Caine could reasonably act on. Second, if the paper is meant as a recruiting tool, then it's not obvious why it wouldn't be more efficient to directly lay out arguments why national security experts should start working on existential risk from AI.

It's possible that there's some kind of 5-D chess game being played here where offering an academic discussion of AI problems will subtly attract national security policymakers in ways that are not apparent to me as someone who's only spent three years in DC. I don't object to having some experts play some 5-D chess...but I believe that in addition to the 5-D chess, we should also be trying the direct strategy of clearly explaining what we want and why we want it. That direct strategy seems to be woefully underemployed. If our intermediate goal is to recruit more national security policymakers, then I would like to see someone publish a paper with a title like "Why National Security Experts Should Address Existential Risks from AI."

As one final example of "level 4" research, consider the paper "Who Should Develop Which AI Evaluations?" The paper discusses four different methods that could be used to develop AI evaluations, suggests some of the advantages and disadvantages of each method, and then uses this analysis to propose criteria that could be used to determine which method to use for particular types of evaluations. The categories developed in the paper are quite reasonable, but the paper is almost entirely disconnected from any concrete recommendations. It does not say which politicians should act, or when, or what they should do. If AI developers choose to use whichever evaluations are best for their marketing, rather than use the evaluations that are best for public safety, how would the authors respond? It's not clear. The connection between this paper and the success of any

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

All of the examples I've cited here are, in my opinion, excellent academic research. For example, Professor Acemoglu's research won a Nobel Prize. My purpose in citing these examples is not to criticize them ***as academic research***, but rather to point out that even the best academic research is unlikely to have much marginal impact on AI developers' incentives. Just because a paper has "extinction risk" in the title doesn't mean that publishing the paper will reduce extinction risks. There's not much correlation between how academically important a paper is and how useful that paper is to advocates.

This can be difficult to accept, especially if you or your friends are currently working on AI governance research. It's tempting to generate reasons why research is valuable even if it doesn't straightforwardly lead to better laws. Perhaps research can add to the credibility of the AI safety movement as a whole, or perhaps research is a good way to make valuable connections with other experts, or perhaps research will help us identify the best possible policies so that if we (somehow) acquire some political power in the future, we'll be able to use that power as effectively as possible.

I don't actually disagree with any of these suggestions – I just think they're not intense enough to outweigh the benefit of having another advocate out there knocking on doors and directly telling politicians about existential risk.

We really need everyone we can get to spread the word in DC. I have been shocked and humbled to see how many Congressional offices were simply unaware of basic facts about AI safety. In December 2024, I met with at least five offices – including some on the Judiciary Committee – who were very surprised to hear that AI developers aren't covered by existing whistleblower laws. In February 2025, we met with a Representative who didn't know that large language models aren't naturally human-interpretable. In April 2025, I met with a district office director who asked me for informational materials to help explain what a data center is. If we don't send people to DC to personally tell politicians why misaligned superintelligence is dangerous, then most of them won't ever understand.

We don't need to do advanced research to cure this basic ignorance: we just need to talk to people, preferably over coffee. There is ***so*** much low-hanging fruit on the ground. I'm begging you to help me pick it up. I think many of the people who have been working in academic research could and should learn how to become effective advocates – even today

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

For researchers who can't or won't work directly as advocates, I strongly urge you to find a way to refocus your next paper on something that's ***fewer*** levels removed from the core task of changing AI developers' incentives. If you've been writing level 4 theory papers, consider writing a level 3 commentary on a specific policy instead. If you've been writing level 3 commentary, consider actually drafting a level 2 example of a policy and including that draft as part of what you publish.

After you publish a paper that took you four weeks of work to write, I challenge you to take 5% of that time – ***8 hours by the clock*** – and ***spend it trying to get that paper into the hands of someone with some decision-making authority***. If you don't know anyone with decision-making authority, spend the time calling or emailing your colleagues and seeing if they can introduce you. If your colleagues don't know any decision-makers either, then spend the time going to a conference or a meetup and keep meeting more people until you ***do*** know someone who can change a policy.

In the next post in this sequence, I'll talk about why ***now*** is the right time to shift resources from research to advocacy. A decade ago, we were absolutely right to be focused on research, because you need abstract research to found a new field. Now that the field of AI governance has been reasonably well-established, though, we need to work in that field, which means we need to put down some of our pencils and pick up some telephones. I will examine various reasons for further delaying that shift and show why none of them are defensible.

Previous:

The Need for Political Advertising (Post 2 of 7 on AI Governance)

No comments 61 karma

Next:

Shift Resources to Advocacy Now (Post 4 of 7 on AI Governance)

5 comments 63 karma

Log in to save where you left off

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

59 Career Advice for Policy Researchers

Comments 3

Sorted by New & upvoted

Marcus Abramovitch 11mo 12 1 0

This is a great post. Way too many people in EA want to be able to do remote work behind a computer and not get into the thick of things that actually change things. I agree with your marathons and smores analogy; I don't think I agree specifically with what you called the Levels for AI work but I digress since I like the main message.

Too often, I think people make the mistake that if they are working up the stack, they are getting a lot of leverage but we need 10x the amount of people at the lower trophic levels and without those, you don't have any leverage at all. There are few exceptions to this of course, like MATS who produce a bunch of AI safety researchers but that's directly increasing the number of "level 0" people.

Jason Green-Lowe 11mo 1 0 0

Yes, that's a great insight! People assume that if they're high up the stack, then they must have a lot of leverage -- and this can be true, sometimes. If you are the first person to run a study on which curable diseases are neglected, and there are a million doctors and nurses and epidemiologists who could benefit from the results of that study, your leverage is enormous.

However, it's not always true. If you're the 200th person to run a study on the risks of AI, but there are only 60 AI advocates who can benefit from the results of that study, then your leverage is weak.

I don't want to insist on any particular number of levels for any particular kind of work -- the key point is that on average, AI governance is way too high up the stack given our current staffing ratios.

Jonathan Jacqueminet 9mo 2 0 0

I agree with these arguments. Though part of me feels like research (both practical policy research and abstract research) is still very important because it seems we don't actually have sufficiently good answers just yet.

Obviously, the shorter your timelines, the more freestyle muddling through based on uncertainty about good solutions seems optimal

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your cookie settings for more control, or review our cookie policy here.

More from Jason Green-Lowe

141 **Please Donate to CAIP (Post 1 of 7 on AI Governance)**
Jason Green-Lowe · 1y ago · 40m read
25

63 **Shift Resources to Advocacy Now (Post 4 of 7 on AI Governance)**
Jason Green-Lowe · 11mo ago · 38m read
5

61 **The Need for Political Advertising (Post 2 of 7 on AI Governance)**
Jason Green-Lowe · 1y ago · 16m read
0

View more

Curated and popular this week

Impact Requires Implementation: What Dispensers for Safe Water Scale-Back Can Teach Us About Monitoring

Tony Senanayake · 13d ago · 17m read · 3

TL;DR Between 2022 and 2023, GiveWell and Open Philanthropy committed over $70 million to Evidence Action's Dispensers for Safe Water programme. An independent survey in 2025 found that in Uganda and Malawi the programme was reaching approximately 70% fewer people with detectable...

The AI people have been right a lot

Dylan Matthews · 4d ago · 7m read · 11

This post was crossposted from Dylan Matthew's blog by the EA Forum team. The author may not see or reply to comments. ---------- Subtitle: Try to keep an open mind as the world gets...

In Development Magazine Launch

Lauren Gilbert · 4d ago · 1m read · 4

We're announcing a new development magazine: In Development. Our goal is to promote discourse on

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Recent opportunities in AI safety

11 **Pivotal Research Fellowship applications are open (deadline May 3)**
Tobias Häberli, Tilman, James Nicholas Bryant, Ariel · 13h ago · 3m read
0

150 **AI Safety's Biggest Talent Gap Isn't Researchers. It's Generalists.**
Topaz, Agustín Covarrubias 🔶 , Alexandra Bates + 2 more · 7d ago · 7m read
9

148 **The Anthropic IPO Is Coming. We Aren't Ready for It.**
Sophie Kim · 8d ago · 9m read
16

View more

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.