AN ACTIVIST VIEW OF AI GOVERNANCE

Shift Resources to Advocacy Now (Post 4 of 7 on AI Governance)

63

by Jason Green-Lowe May 27 2025

AI safety | Existential risk | Policy | AI governance | Political polarization | Frontpage

In my previous post in this series, I estimated that we have 3 researchers for every advocate working on US AI governance, and I argued that this ratio is backwards. When allocating staff, you almost always want to have more people working on the more central activity. I argued that in the case of AI policy, the central activity is advocacy, not research, because the core problem to be solved is fixing the bad private incentives faced by AI developers.

As I explained, the problem with these incentives is less that they're poorly understood, and more that they require significant political effort to overturn. As a result, we'll need to shift significant resources from research (which helps us understand problems better) to advocacy (which helps us change bad incentives).

In this post, I want to explain why it's appropriate for us to shift these resources ***now***, rather than waiting until some future date to scale up our advocacy.

The arguments for why we should wait until later to start advocating can be divided into three broad categories: (a) we aren't yet sure that any of our policies will be robustly helpful, (b) we haven't yet learned how to advocate successfully, and (c) we don't yet have enough senior politicians in the movement to staff a large advocacy effort. What these objections have in common is that they seem to justify an emphasis on research for the next few years – if we're missing knowledge or experience that we would need in order to win at advocacy, then maybe we should bide our time while we gather more resources.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your cookie settings for more control, or review our cookie policy here.

Reject Accept all

can solve the challenge of recruiting effective advocates with willpower, flexibility, and funding.

Moreover, even if we were not self-confident, our chance of success will not be significantly improved by waiting or by doing more academic research. On the contrary, we face enormous time pressure to get better incentives in place before AI developers push us past the point of no return, so we need to start moving AI regulations forward now, using whatever policy tools and policy advocates are currently available.

OUR BEST POLICIES OFFER POSITIVE EXPECTED VALUE

There are several commonsense AI safety policies that would clearly offer positive expected value if they were enacted. I discuss them in detail in the first post in this sequence, but what they boil down to is that AI developers should have to check in with at least one other stakeholder before deploying billion-dollar models. This could be an insurer, an auditor, an internal safety researcher who enjoys whistleblower protection, a government regulator, and/or an in-house defense attorney who advises them on the (non-trivial) risk that deploying the model will lead to an expensive lawsuit.

The reasoning is that instead of letting CEOs immediately deploy a model whenever they're tempted to do so, it's clearly better if the CEO has incentives to get a second opinion and double-check their work. Erroneous decisions to train or deploy a model that turns out to be dangerous are high-stakes, predictable, and avoidable, so we should put some incentives in place that help CEOs avoid those errors.

As discussed earlier in this sequence, CEOs face enormous temptation to defect against society by releasing products that haven't yet been adequately tested. Executives capture a larger share of the upsides of early releases (profits, market share, glory, etc.) than they are forced to bear from the downsides of early releases (deaths, illness, war, loss of human control). A very similar set of incentives also gave us cars without seat belts and children's pajamas that catch on fire – there's no good economic reason why society as a whole

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Another way of framing this issue is that by default, the court system is less reliable than the marketplace. There is very little piracy of, e.g., ChatGPT, so if you as a customer want access to the benefits of the product, then you will have to pay for it, but if you as an innocent bystander are harmed by the product, then it's not obvious that OpenAI will have to pay for that. Most lawsuits are never filed at all, and of the ones that are filed, most of them are dismissed or settled for a small fraction of their value. In order to change that result, we need either upgraded liability laws, or some kind of regulation, or both.

We've Already Laid a Philosophical Foundation for AI Governance

One response to this call for AI safety policies is to say that they sound promising, but that we're radically uncertain about how AI works, so it's too soon for any type of regulation or even to adjust private incentives through, e.g., liability and auditing.

For example, back in 2021, Luke Muehlhauser of Open Philanthropy suggested that there is no consensus° about how to best respond to AI – at that time, he believed there was uncertainty about whether it is good or bad to hasten the development of AI. In 2023, Luke partially adjusted his stance on these issues, noting that "much has changed" since he articulated his uncertainty in 2021, and that there are now some policies that he "tentatively thinks would increase the odds of good outcomes from transformative AI." It's not clear to me to what extent Luke or Open Philanthropy have become less tentative between 2023 and 2025, but I would argue that today, a "tentative" stance on most AI governance policies is not really appropriate – we should instead be firmly supporting them.

Radical uncertainty is only appropriate at the very beginning of a field, when the field's core concepts are still in the process of being defined. It might have been premature for Gregor Mendel to argue in favor of specific laws to conserve biodiversity, because the entire concept of what it meant for species to be different from each other was still a topic of hot debate. Likewise, it would not have been very convincing if Sigmund Freud argued for universal mental health care, because even leading figures were still sharply and reasonably divided about what it meant to do psychoanalysis and who could be helped by

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

evidence showing that these techniques are helpful, and we have a good theoretical understanding of why the techniques would be expected to help, and the way that the techniques help lines up nicely with the relevant theories.

Likewise, we already know enough about AI governance to move forward with some basic policy prescriptions. Experts in the field are not confused about what reinforcement learning is, or what gradient descent is, or why these techniques will not automatically lead to anything like empathy or wisdom. We understand why deep reinforcement learning usually creates an uninterpretable set of artificial goals which are likely to diverge rapidly and unpredictably from human values when applied in contexts that are different from their original training environment. We know that an AI that appears to yield helpful, polite, accurate responses in training may yield unhelpful, hostile, or inaccurate responses after being deployed, and we know why this happens.

Similarly, we are not confused about instrumental convergence or situational awareness. We know that a wide variety of AI designs with a wide variety of initial goals are likely to adopt power-seeking behaviors. Indeed, we have been able to demonstrate this mathematically since 2021. We also know that as AIs continue to scale, they will spontaneously become better able to identify and pursue opportunities to gain power, even if no one has specifically attempted to train the AI to have power-seeking capabilities.

Whenever we have bothered to test these theories, we have found ample experimental evidence that AIs will at least occasionally lie, scheme, self-replicate, blackmail humans, ignore orders to shut down, break the law, and accumulate resources. We have found that AI will at least occasionally act against the interests of humanity when this seems to be useful for maximizing the AI's reward function. That is exactly what we should expect AIs to do – the idea that AIs will always be inherently kind or restrained or cooperative is naïve anthropomorphism that does not have any theoretical support from the structure of AIs or from our recorded experiences with them so far.

We don't have to wonder about whether more intelligent AIs will have more opportunities to maximize their reward functions at humanity's expense, because we've seen that even relatively modest increases in scale of the current architecture have repeatedly unlocked more and more capabilities, such as speaking multiple languages, generating real-time

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

AI beating the human world champion at chess, Go, Jeopardy!, poker, Diplomacy, and Starcraft, all of which came as a surprise.

At the very least, this means it is no longer possible to be justifiably confident that further capabilities will not be unlocked soon – whatever flaw in the experts' analysis prevented them from accurately forecasting the arrival of AIs that could generate useful software code or solve protein-folding problems could also be preventing them from accurately forecasting the arrival of superintelligence.

Similarly, we have no way to be justifiably confident that we will be rescued in time by, e.g., alignment research conducted by AI agents. Some forms of AI might accelerate technical alignment research, and that research might be decisive, but we won't know in advance that such a research program is going to succeed or even that it will be very likely to succeed. Part of the problem with delegating the problem of alignment to agents that are smarter than we are is that we won't know for sure whether they're telling us the truth.

The combination of (1) AIs that are expected to prioritize their own power over humanity's welfare and (2) AIs that could plausibly outcompete humans is all that is needed to generate a serious risk of serious catastrophe in the near future. It's not necessary for us to determine exactly how often these problems occur, because even if only 1 in 100 AI models display these behaviors, that's still an unacceptable risk – there are thousands of models circulating in the ecosystem, some of which will be deliberately fine-tuned to adopt such behaviors or which will face incentives structures that reward them for adopting such behaviors. There is much that remains uncertain about the future of AI – we don't know for sure whether superintelligence will arrive in 2027 or 2037 or 2047 – but we don't ***need*** to answer these questions before we can confidently determine that ***the correct level of regulation today is not zero***.

The fact that superintelligence is reasonably likely to be extremely dangerous is all we need to justify a wide swath of AI regulations – any of which would be better than the *status quo*. We can and should disagree about the exact type of regulations that would be best, but we must not let this blind us to the urgent necessity for passing *some* of them.

There's no room left for reasonable people to take refuge in radical uncertainty about

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

to at least suggest a specific cost or downside to those regulations that might outweigh their benefits.

Regulations Won't Backfire by Pushing Companies Overseas

In the first post in this sequence, I discussed why moderate AI regulation won't cause the US to lose an arms race with China. If anything, sensible regulation will help the US widen its lead by avoiding unnecessary technology transfers. A related concern is that overly strict regulations could drive AI developers overseas, to countries that have weaker laws.

However, as Anthropic CEO Dario Amodei points out, AI developers' threats to leave the US are not even slightly credible. In his words, "That's just theater, that's just negotiating leverage." Most leading tech companies still have their headquarters in Silicon Valley or Seattle or Manhattan even though these are crowded and expensive and high-tax places to live and work. If they're not going to shift their research buildings 100 miles away to save a few billion a year in rent, then why would they want to shift their research buildings halfway around the world just to escape from regulations whose purposes they at least broadly agree with? And even if tech companies wanted to move, what makes us think that their scarce, valuable AI experts would agree to pick up their families and follow them to China?

Regulations should be designed with some care to reduce the burden on AI developers: paperwork should be minimized, the costs of compliance should be borne by the government whenever possible, and rules should be no stricter than necessary to accomplish their objectives.

If specific tech companies nevertheless object to specific regulations and credibly threaten to move their laboratories if the regulations are passed, then we can weaken or adjust the regulations as necessary until those threats fade away. However, we should not preemptively abandon the entire project of AI safety regulation just because someday some company might respond by trying (and most likely failing) to move some of its staff to a rival country. If your actual goal is to win a competition for AI talent, you can do that

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Regulation Is Helpful Even If It's Not Fully Future-Proof

Another complaint that CAIP sometimes received about its model legislation was that the legislation isn't fully future-proof. Our proposed bill would rely in part on thresholds based on FLOPs, but AI evolves rapidly enough that the specific number of FLOPs specified in the bill might quickly become inappropriate. If there is a paradigm shift in how AIs are trained, then the entire concept of FLOP-based thresholds might become obsolete; perhaps in 2028, it will be just as nonsensical to talk about identifying frontier models based on FLOPs as it would be today to talk about identifying the most advanced models based on what fraction of the human experts who provided their decision-making criteria had PhDs.

Part of the answer to these complaints is to point out that the bill itself contains numerous provisions for keeping it up to date. For example, CAIP's model legislation calls for establishing an entire Office of Standards whose only job would be to stay informed about how AI is changing and to rewrite the technical thresholds in the bill as needed to keep pace with those changes. There are also sections in the bill that deal with catching and correcting both Type I and Type II errors (so that models can be subject to the correct degree of regulation even when the level of risk they pose does not match up well with the number of FLOPs used to train them), and that allow ordinary citizens to file requests for the agency to take action with respect to new types of threats, which the agency would be legally obligated to respond to.

However, the more powerful answer is that we should not let the perfect become the enemy of the good. Yes, it's possible that with enough technological change, some parts of AI safety legislation could become ineffective, and it's quite plausible that a future Congress would be too slow to adapt that legislation, leaving the public unprotected. This is ***not*** a reason to never pass any safety legislation at all! Even in this 'worst-case' scenario where the bill eventually becomes obsolete, we still get the benefits of its protection in the meantime. AI catastrophe is not solely an all-or-nothing affair; we might be quite grateful if safety legislation helps us escape, e.g., an AI-assisted engineered pandemic that would otherwise have caused 50 million deaths in 2028, even if that legislation becomes obsolete

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

not have this problem, or (2) the legislation will become so profoundly obsolete so quickly that it will do more harm than good.

It doesn't make sense to suppose that future advocacy efforts will have less of a challenge with future-proofing their proposals, because the rate of technological change is only expected to accelerate for the foreseeable future. In the past, it was a good heuristic that if you felt overwhelmed by a new technology, you should wait a few years before trying to regulate it, until the technology had become better understood and more predictable. Now, though, the next wave of technology is likely to arrive before the previous wave has been fully assimilated. We live in uncertain and rapidly fluctuating times, which is uncomfortable, but we can't escape that discomfort by waiting for things to stabilize, because the immediate future will only be even more unstable than the present.

It likewise doesn't make sense to assume that even a well-drafted and reasonably flexible bill will quickly become not just obsolete but actively harmful – at least, not if the harms you're thinking about include any type of existential risk. It's possible that some companies will face some degree of unnecessary or futile regulatory burden after an AI safety bill becomes obsolete, but the costs of complying with the policies CAIP has championed are relatively small – on the order of $60 million per year in government spending and $1 billion per year for the entire industry's compliance costs. This is less than 1% of the industry's 2025 revenue. These estimates are for a strict regime that imposes heavy requirements on all frontier AI companies – if you think that's too expensive, you could still get a useful amount of safety by having weaker requirements.

There is simply no comparison in terms of the economic value of even slightly reducing the odds or even slightly delaying the onset of the collapse of civilization. If you conservatively estimate the value of a human life at $1,000,000, then an 0.1% chance of avoiding 1 billion casualties is worth $1 trillion – far more than companies could ever lose by complying with requirements to, e.g., conduct a detailed audit of their largest models, or to file a report about their wholesale purchases of advanced chips.

Part of why it so obviously makes economic sense to do ***something*** to regulate advanced AI is that right now we have essentially zero regulation on American AI developers, so the marginal gains of the best available interventions are huge, and the marginal costs of the

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

AI safety regulations can provide enormous public benefits at a very low cost. Even if these benefits decay quickly over time while the costs slowly increase, the net present value is still very likely to be hugely positive.

Regulations Don't Have to Lead to Oligopoly

Another way AI safety policies might backfire is that they could somehow promote a monopoly or an oligopoly on frontier AI, which would further centralize power and lead to corruption or oppression.

Most of this argument rests on an unrealistic view of how expensive it would be to comply with commonsense safety regulations. AI is one of the least regulated industries in America; there are considerably more regulations on ice cream, beans, airplanes, and haircuts, even though none of these products pose any existential risks. Nobody suggests that the ice cream market is being cornered by one or two companies because of all the onerous regulations.

If it is very expensive to comply with a regulation, then it might be the case that only the largest or most-profitable companies can afford to do so, but the expense of, e.g., completing an audit is tiny compared to the expense of training a large AI model in the first place. If you can afford to spend a billion dollars on compute, you can also afford to spend a million dollars on running an audit; the odds are tiny that even one company will be driven out of the market because their profit margins are so thin that they literally would rather exit the market than pay for the audit. The odds that *most* such companies will be driven out by mild, commonsense regulations are nonexistent.

People are nevertheless concerned about allowing any type of regulation because they fear that it could lead to regulatory capture. According to this view, safety advocates might mean well and have the best intentions, but so did many other types of social reformers, and all of their reforms quickly wound up being subverted and serving as an excuse to give more power to large corporations who could bribe or pressure the regulators into doing their nefarious bidding. Instead of blocking companies that don't do things safely, regulators supposedly spend their time blocking companies that aren't part of a small

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

It's reasonable to be concerned about regulatory capture, which is part of why CAIP's model legislation includes anti-capture provisions such as a requirement for the agency to publish any exceptions it grants in the Federal Register, a right for citizens to formally protest agency actions or non-actions and receive a timely response from the agency, and a Deputy Administrator for the Public Interest who is empowered to investigate and report on internal misconduct.

However, it's unreasonable to assume that regulatory capture ***always*** kicks in promptly and always undoes all of the benefits of new regulations. For example, Jim Epstein at *Reason* claims that the government's shutdown of the Chinatown bus industry in 2012 was due to regulatory capture, saying that it was "largely based on fabricated charges" and was spurred on by "politically-connected corporate carriers" like Greyhound who wanted to shut down the Chinatown buses so that they would have less competition. That's one side of the story, but the NHTSA found that the Chinatown buses had seven times the normal rate of fatal accidents, based in part on practices like not requiring drivers to have valid drivers' licenses, not administering drug tests, allowing drivers to work double shifts without a rest, and not performing regular inspections and maintenance of their vehicles.

For the most part, regulatory capture is part of the ordinary business of politics – it's certainly a threat to be concerned about, but the exact extent of that threat is heavily disputed, and it's not a reason why government projects will automatically fail.

If it were easy and reliable to engage in regulatory capture, then we would expect the civil service to be one of the country's more lucrative professions, with status-conscious or insecure parents advising their children to go into the bureaucracy so that they can collect large bribes. This was indeed the advice that many children would get in, e.g., Qing China or Soviet Russia, but it's not anything I've heard in America.

Regulations Won't Significantly Promote Nationalization

A variant of the regulatory capture idea is the fear that promoting any type of regulation will prompt the US government to partially or wholly nationalize the AI industry, bringing

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Several years ago, people used to say that we shouldn't bring AI to the government's attention, because the government might be more interested in nationalizing AI if they knew how powerful it would be. That particular ship seems to have sailed; when Oprah and the Pope and Joe Rogan are all talking about AI, it's not really going to be a secret from the White House.

The only remaining question is whether we should expect that having modest, commonsense regulations like an insurance requirement or a whistleblower protection law will somehow increase the likelihood that the government will inappropriately nationalize AI. It seems clear to me that the answer is "no," because there is no precedent in American history for the full nationalization of such a large and private industry. America hasn't even attempted to ***maintain*** an incidentally arising government monopoly on government-sponsored technologies like nuclear energy and the Internet; we're certainly not about to try to purposefully ***establish*** a government monopoly over a privately-developed technology like AI.

The two closest parallels (which still aren't very close) were President Wilson's temporary nationalization of American railroads to solve supply problems during World War I, and President Truman's failed attempt to temporarily nationalize the steel mills in 1952 to end a labor strike.

President Wilson acted pursuant to explicit authority granted by an act of Congress, and that act specified that the nationalization would automatically end shortly after a peace treaty was signed, and that all railroad companies would retain ownership of the 'nationalized' assets and that the government would have to make regular lease payments on those assets and return the assets in the same or better condition as when they were borrowed. There was also no attempt to prohibit the use of railway technology by local entrepreneurs or researchers; Proclamation 1419 only applied to the larger "common carriers," and even they were encouraged to continue normal operations unless and until they received specific wartime orders from the government. The railways were returned to private control on schedule, three years later.

President Truman's attempt to nationalize the country's steel mills was decisively rejected by both the Supreme Court and the popular press. Several major newspapers ran front

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

unions got most of what they wanted from the strike, Truman's approval rating dropped to a historical low of 22%.

Steel mills were considered just as vital to national security in 1952 as AI might be in 2025 or 2030 – the strike directly endangered the supply of ammunition for American soldiers in the ongoing Korean War, as well as interfering with weapons deliveries to fragile American allies. The steel mill seizure was intended to be a temporary measure, and it was proposed in the context of a political environment that was far more comfortable with government intrusion into the economy than the environment of today – for example, Truman's Wage Stabilization Board made binding decisions about the prices and wages set by private companies. Nevertheless, the steel mill seizure was firmly defeated.

A permanent (or indefinite) seizure of AI companies' assets would be even more politically challenging and is extremely unlikely to take place under anything like the current political environment. Making nationalization happen would require a far more radical stimulus than a whistleblower protection law or a system of government auditing that parallels the audits that have been in place for decades for airlines and buses. It's possible we might nationalize AI someday – but if we do, that nationalization won't happen *because* of modest AI regulation. The scale of change needed to cause nationalization is just too large for modest AI regulations to be a significant contributing factor.

ADVOCATING FOR POLICIES NOW MAKES THEM MORE LIKELY TO PASS

Advocacy Won't Offend Politicians

Even if we are confident that our AI policies (if enacted) would be net benefits for humanity, one might think that openly advocating for those policies will somehow backfire, causing those policies to be less likely to be implemented compared to a world where we published our policy solutions in academic journals and otherwise remained silent. According to this argument, we should restrict ourselves primarily to research (which we know will increase the sum of human knowledge) and avoid spending resources

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

There's two different variants of this argument. The first variant relies on unusually incompetent advocates. I will grudgingly admit that political advocacy that's sufficiently clueless or offensive can sometimes be of negative value – if, for example, you go around telling Congressional staffers that only an idiot would support further AI development, or if you repeatedly interrupt them and raise your voice while they're trying to talk, then that might do more harm than good.

However, it's very easy to avoid this level of incompetence. Remain calm, ask some questions, listen respectfully to the answers, try to find out how you can be of help to the person you're talking with, and be reasonably polite. That's really all it takes to avoid insulting politicians to the point where you would have been better off not showing up at all. By the very nature of their work, politicians have to deal with and charm a wide cross-section of the public, not all of which are diplomatic or kind. If you can manage to have a civil conversation, you can be competent enough to avoid conducting advocacy that has literally negative value.

The second variant worries about partisan polarization. If your choice of arguments for AI safety or your way of presenting yourself feels too closely associated with one political party, then you might create or strengthen an association between AI safety and that party, which could lead the other party to reject AI safety on partisan grounds.

This is a real concern, especially if you're coming from a heavily progressive area like Berkeley (90% of which voted for Kamala Harris) and you aren't aware of how your political beliefs and vocabulary differ from the American median. In my experience, though, this is a problem that can be solved with an average of four hours per person of training. It is not ***that*** hard to get well-intentioned people to stop making analogies to global warming and COVID while they're at work.

Other key strategies for being and appearing to be a bipartisan advocate include insisting that the bills you endorse have co-sponsors from both parties, double-checking the open letters you sign to make sure they are not insulting either political party, and keeping an eye on the overall partisan balance of staff at your organization so that you can make a point of hiring some Republicans (or Democrats) if you're short on employees from that party. This last point pays double dividends in that when there are people from both

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

If you know that you're a rabid political partisan and that you will be unable to exercise enough self-control to mostly frame AI safety arguments in a mostly bipartisan way, then you should stay away from public advocacy. Otherwise, it's not that hard to frame your arguments in terms that have no particular partisan bias.

Advocacy Isn't Useless

Sometimes skeptics about advocacy seem to me to be confusing uncertainty about the results of any particular political campaign (which is valid) with uncertainty about whether political advocacy in general tends to help (which seems very hard to justify).

Any given political push might fizzle out if it lacks the resources to convince its target audience, or if that audience is distracted by breaking news. For example, in summer 2001, Congress was very seriously considering a Patients' Bill of Rights, and President Bush was actively negotiating a guest worker program with Mexico. Both of these initiatives were dropped when the World Trade Center was attacked on 9/11 as national attention shifted almost exclusively to national security. Anyone who donated money to support those two political initiatives essentially lost their investments – but they didn't lose their money because they made bad decisions; they just got unlucky.

Slightly earlier that same year, the Chamber of Commerce and the Heritage Foundation played a major role in successfully supporting President Bush's tax cuts (EGTRRA), and the NAACP and the Education Trust were important contributors to passing the No Child Left Behind Act.

There's no meaningful difference between the support offered by the Chamber of Commerce for EGTRRA and the support offered by, e.g., Families USA and the Kaiser Foundation for the Patients' Bill of Rights. One group was lucky enough to get a news cycle that allowed its preferred bill to pass, and the other wasn't. This is part of how politics works, and it would be really strange to suggest that all the various sophisticated, professional, well-funded organizations that are constantly supporting legislation in DC are all wasting their time and money. ***Sometimes*** a particular campaign ends up being wasted, but in expectation, the campaigns are all valuable to their supporters.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

More money means more chances to advertise, more ability to recruit higher-quality lobbyists, more ability to recruit volunteers, more ability to support those volunteers with high-quality materials and travel expenses, more ability to donate to relevant campaigns, and so on. Nobody suggests that a political candidate will be defeated in an election because people keep donating too much money to their cause; the situation is no different for political policies.

Political Opposition Can't Be Avoided by Delay

A related theory of political uncertainty is that revving up political advocacy at the wrong time or in the wrong way can galvanize or crystallize your political opposition, making it harder for future attempts to succeed. One of the most lucid exponents of this theory is Anton Leicht. Leicht isn't wrong to point out that, e.g., if even pro-regulation California can't sign a relatively minimal AI safety bill into law, then "policymakers elsewhere might believe the case for safety policy could not possibly be very strong."

Unfortunately, there weren't any good alternatives to taking a swing at something like CA SB 1047. If AI safety advocates had talked about the dangers of AI for years on end without anyone even attempting to pass an AI safety bill, then people would start wondering whether the rhetoric is disingenuous – if you're so worried about AI, why aren't you trying to do something about it? Is it because the topic is so complex that it's not even possible to draft good legislation? Is it because you're just a smoke screen meant to distract attention from AI ethics issues?

Or, if AI safety advocates carefully refrained from mentioning the dangers of AI in public for years, then people would ask why they remained silent. Why the conspiracy? How dare AI safety experts sit on the knowledge that a tragedy was coming and wait for people to die before they took action? Why should the American public listen to anyone so heartless?

In the abstract, it would be nice to wait until the perfect moment to launch a political campaign to maximize your odds of success. In practice, there's no way to simultaneously defuse all possible political attacks. No matter what political strategy you choose, someone will find a plausible way to attack you for it, because that's their job. Being attacked doesn't

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

alliances that you can absorb or deflect the attacks and carry on to a political victory. Both the Affordable Care Act and the Bush-era tax cuts came in for serious and prolonged attack and ridicule from many important stakeholders; nevertheless, both policies became law and have remained in effect for many years.

Another reason not to wait too long before beginning a political campaign is that there was never any serious possibility of getting meaningful AI safety legislation to "sneak through" a legislature without Big Tech noticing and opposing it. Meta alone employed 65 federal lobbyists in 2024, not including their support staff, policy researchers, and so on, and not including any work they do on state legislatures, on impact litigation, on general public relations, and so on. OpenAI employed 18 federal lobbyists. Alphabet employed 90 federal lobbyists. Amazon employed 126 federal lobbyists.

That's **299 lobbyists** just from those 4 companies. The vast majority of these lobbyists would have jobs even if there were never any AI safety advocacy; the companies are building up their political muscle in large part so they can win unrelated battles over copyright and electricity and privacy. At least one of these lobbyists is definitely going to notice if there's an AI safety bill slowly working its way through Congress.

Bills are designed to take months to turn into laws in large part to guarantee that all reasonably alert stakeholders have a chance to weigh in on whether those bills make sense. After a bill is introduced, it has to be referred to at least one committee, which holds a public hearing to markup the bill and consider amendments, and then it has to wait for a slot to be considered by the full House or the full Senate, and then the whole process repeats in the other chamber, and then there's another chance to kill the bill while the President or Governor decides whether to veto it. It can take a public relations firm a couple of weeks to develop a marketing strategy and execute on it, but the PR firm needs much less time to mobilize than a bill needs to become law.

From the outside, it can certainly appear that introducing an AI safety bill or making a speech about AI safety "caused" opposition to arise and therefore made future bills less likely to pass. The opposition chronologically follows the introduction of the bill, and if the bill were introduced later, then the opposition would occur later. However, the opposition is basically guaranteed to arise at whatever time you do introduce the bill. There's no tactic

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

responsible for changing the total amount of opposition. The opposition was always there, it's just that most people only notice the opposition while a political contest is actively unfolding.

Improving Advocacy Skills Requires Practice

A variant of the argument above is that while the opposition to any given policy proposal might be constant, our ability to overcome that opposition can vary over time. Perhaps if we wait until some later date, our advocacy will be more effective, so we should wait to arouse the opposition until after we are ready to defeat it.

The problem with this strategy is that the main way to improve the strength of an advocacy movement is to practice engaging in advocacy. A key reason why AI advocacy efforts are relatively weak right now is that very few people who deeply understand AI safety are skilled and experienced political advocates. Most of the people who have a rigorous understanding of the catastrophic risks posed by AI are either computer scientists (who discovered the problem based on their professional research), or philosophers (who discovered the problem based on asking abstract questions about which types of problems are likely to do the most harm). Neither group is likely to have more than a token amount of explicitly political expertise – for example, very few computer scientists run for electoral office or work for multiple years in government; most of the computer scientists who do work in government have technical responsibilities, rather than policy responsibilities.

This lack of political experience is not a problem that can be solved by writing white papers, any more than an aspiring pediatrician could expect to master her craft solely by publishing articles in medical journals. Being good at politics requires practice; you need to go out into the political arena, try to sell politicians on your ideas, try to organize volunteers and get them to show up for your cause, try talking to the media and seeing if they bother to publish your quotes, and so on.

You also need a robust set of relationships: in politics, as in many other areas of life, it's all about who you know. If you don't have pre-existing relationships with potential allies,

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

For the most part, research and writing is an extremely inefficient way to build the relevant relationships. Building a reputation as an academic expert in your field can get you a warm introduction to some politicians, but you still have to go and seek and use that introduction, and parlay it into additional connections to other stakeholders, and then contact those stakeholders often enough with information that's relevant to their day-to-day work so that they'll remember who you are and care about their connection with you. The vast majority of researchers don't have the time or the inclination to do all of that work on top of their written research.

If you're earnestly concerned that you're not yet skilled enough as an advocate to successfully cope with active opposition, then it might make sense to keep a somewhat lower profile – but that would mean keeping a lower profile as an advocate, not working in research. You can advocate for policies that are less controversial, or you can advocate in a less competitive forum (e.g., state or local politics instead of national politics). I have nothing against warm-up exercises. The problem is that policy research isn't efficient enough as an advocacy warm-up exercise to be worth investing in for that purpose.

If your goal is to increase your skill as a ***political*** advocate, and you insist on doing research to achieve your goal, then at the very least you ought to be engaging in ***political*** research, not policy research. Make a list of the most influential Senators; find out who their chiefs of staff are and what those people have said about AI policy in the past. Find someone who went to the same university as the chief of staff and see if they can get coffee. Pull up a list of bills that a Representative has voted for in the past, and see if you can convincingly connect some piece of AI safety legislation to one of their pet issues – if you're looking to support AI whistleblower protections, it helps to know that Senator Grassley loves whistleblowers; if you're looking to support mandatory audits, it helps to know that Senator Hawley doesn't trust Big Tech. Having the perfect policy won't help you if you have no idea how to pitch that idea politically.

Even then, though, don't do too much research. Politics inevitably comes with a small-n dataset: no situation is exactly like the one before it, it's not subject to any rigid physical laws that can reliably predict its outcomes, and there are sharp limits to how much you can learn about politics by studying it in a library. If you want to improve your advocacy skills,

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

WE CAN INCREASE THE SUPPLY OF ADVOCATES

It doesn't bother me that there isn't already a tailor-made supply of highly experienced AI safety advocates ready and waiting to start work, because I think we can rapidly hire and train enough new advocates to be helpful. Direct political advocacy is so much more effective on the margins than academic governance research that it's fine if it costs us a few months to gear up or if some of the newer advocates aren't quite as effective while they're learning their new craft. There are plenty of good ways to spend money right now to increase the effective supply of AI governance advocates.

Generously Fund Existing Advocacy Groups

First, we should be generously funding all of the existing advocacy groups. This includes the Center for AI Policy, Pause AI, the Center for AI Safety Action Fund, Americans for Responsible Innovation, Encode AI, the AI Policy Network, and a few others – if your advocacy organization isn't listed here and you'd like to be included, just let me know and I'm happy to add you in.

I believe all of these organizations can and should absorb new funding. All of us are skipping meetings that we would like to be having with important politicians because we just don't have the staff to cover all of them adequately. All of us are skipping out on methods of communication that would be useful because we have no one to develop them. Every major lobbying organization should have a video producer, an event manager, a social media expert, and a campaign finance arm that makes donations to AI safety champions. For the most part, we don't have any of those tools, because we don't have the funding for it. We're trying to do advocacy on the cheap, which is less likely to be effective.

If you ask an advocacy organization how much funding it can productively absorb, the answer you get will probably be too low, for two reasons. First, hiring high-quality mainstream government relations experts often requires a long-term funding commitment, which is comparatively rare in the AI safety funding space. The students and

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

By contrast, experienced political advocates often have families and mortgages to support, and these commitments can seem more important to them at first, before they've learned enough about AI safety to be truly terrified. Political experts who are earnestly curious about AI safety and who would like to learn more about it and do something to help will nevertheless shy away from joining a new organization that doesn't have a guaranteed financial runway. The solution to this problem is not to accept a sharp upper bound on advocacy funding based on the number of "true believers" who can be lured to work for a fly-by-night organization. Instead, the better solution is to offer generous multi-year commitments. As in many areas of life, we'll get better results if we dive in and commit hard to governance advocacy as a primary strategy than if we just keep dipping our toes in the water and then promptly yanking them out.

The second reason why funding requests from advocacy organizations are likely to be under-stated is that there's implicit pressure on leaders to make a bid that sounds reasonable and that respects the work being done by others in the field. If you make a request that's so large that it implicitly requires transferring resources away from other organizations, that feels antisocial and out of touch. The only exception would be if you have very firm evidence to support the results that you can deliver using your requested funding...but by its very nature, politics is unpredictable and poorly suited to delivering measurable results on a fixed schedule.

So, if we want to scale up political advocacy for AI safety – and we should – then one good way to start is to just double or triple the budget of every existing advocacy organization. The amount of funding these organizations have consumed in the past is not a reliable limit on how much funding they could productively absorb.

Use Headhunters

Similarly, the number of staff that advocacy organizations have been able to hire in the past is not a reliable guide to how many staff those organizations could usefully employ. In 2024, when CAIP was rapidly expanding, I had to hire people one or two at a time, running the initial phone screens and writing the work tests myself, because we had a limited and uncertain budget. This meant that it took literally all year to hire our 10-person team

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

screening out the people who are obviously unsuitable. Just because ***you*** don't know anyone who has the rare mix of skills to be a strong political advocate and to understand and care about AI safety doesn't mean those people don't exist.

With enough money, you can find and attract even people who have very rare skill sets. In addition to paying a professional headhunter to talk to their network, you can establish an advertising budget that will push the job description into the feeds of high-quality candidates, you can pay referral bounties, you can host a booth or a talk at high-prestige conferences that will put you in touch with more high-quality candidates, and you can have a full-time operations manager who clears away obstacles (e.g. visas, relocation, non-standard requests for benefits) and accelerates onboarding.

Part of why I belabored the point in my previous post about how much more valuable advocacy is than academic research is because I think these kinds of investments make sense for advocacy. If, on current margins, a high-quality advocate is ten times more valuable than a high-quality researcher, then we really ought to be paying for headhunters and job ads and operations managers in order to recruit that advocate, even if this soaks up funding that otherwise could have paid for two or three researchers.

Aggressively Train New Advocates

While I do believe that existing organizations have plenty of room to expand, I don't think we can or should rely entirely on existing advocacy organizations to get the job done, because the scale of the advocacy effort we need to produce is so large that we can't comfortably house all of it within the existing organizations. This means we need to be aggressively training both new advocates and new organizational leaders.

Blue Dot Impact has done an excellent job of putting together a basic curriculum in AI governance, and tech fellowships are doing a good job of placing early-career policy professionals in positions that require (and further develop) their scientific expertise. This still leaves a gap in resources that train people on how to be an effective advocate – you can know how AI works and how AI policy works, but that doesn't mean you know how to convince others to adopt good policies. We need something like "debate coaching," but far

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

be better prepared to win hearts and minds. This kind of training could be a good way to help AI safety researchers transition to a career in advocacy.

Similarly, we also need management training, because we need more people who are capable of serving as executive directors and as team leaders. Most of the skills involved in management can be taught – you can learn how to more effectively recruit, hire, train, evaluate, motivate, and terminate your staff. Similarly, you can learn how to create a budget, how to actively listen and accept criticism from subordinates, how to structure an effective board of directors, and so on. Training for some of these skills might be better housed inside the effective altruist or rationalist movement to take advantage of our unique approaches to measuring success, and some of them are probably best learned in mainstream management courses, to take advantage of what the rest of the world already knows about how to run an organization.

People sometimes cite a lack of high-quality leadership as a bottleneck on donors' ability to fund new advocacy organizations – but if that's the true objection, then we can't just sit back and wait for leaders to emerge – instead, we need to be investing in a crash effort to create competent new leaders.

WE DON'T HAVE THE LUXURY OF WAITING

Finally, it's worth remembering that we probably don't have the luxury of choosing the perfect moment to begin seriously investing in advocacy. The time remaining until superintelligence arrives could be very short, and the 'point of no return' beyond which no plausible political campaign can affect what happens with superintelligence could be even shorter than that.

Most major legislation takes years to pass and implement from the time that you start campaigning for it. If you introduce a bill today that has passionate and bipartisan support, Congress will still probably take a minimum of six to nine months to pass it – possibly much longer than that if the bill is defeated or simply deprioritized; in any given year, many bills that are important and well-liked will nevertheless not get called up for a vote by the full House and the full Senate, simply because Congressional leadership is busy with

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

that businesses have a fair opportunity to bring their operations into compliance with the new regulations. So, as an optimistic assessment, if you introduce a bill with massive support today, in May 2025, then you might get effective regulations by, say, October 2027.

Of course, we don't ***have*** massive political support lined up right now, and it takes years to build that support. If we funded several large advocacy groups today, then they might need three years to build an effective mass political movement, meaning that you'd be looking at regulations that would start changing corporate behavior sometime in 2030.

Now, you might not agree with Leopold Aschenbrenner or the AI Futures Project when they say that we should expect superintelligence in 2027 or 2028. Maybe you are very confident that we have more time than that. But effective regulations need to come into being not just before superintelligence is invented, but also while AI companies are still small enough to be truly subordinate to the federal government as a matter of power and *realpolitik*. Mechanize, Inc. recently launched a startup with the goal of automating the entire workforce. If they're even partially successful, how could the federal government credibly threaten them? Other AI companies are working on software that would be vital to, e.g., maintaining command and control over the American military's key weapon systems. It's hard enough even today to get support for legislation when people raise the abstract prospect that the legislation might someday interfere with American military dominance. What happens when the companies to be regulated are currently, actively supporting US military operations?

That's why I don't think we have any time to lose. We should have started massive lobbying efforts a few years ago, but the second-best time to start lobbying is right now.

Previous:

We're Not Advertising Enough (Post 3 of 7 on AI Governance)

Next:

Orphaned Policies (Post 5 of 7 on AI Governance)

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Mentioned in

122 The Center for AI Policy Has Shut Down

Comments 5

Sorted by New & upvoted

LeahC 11mo 3 0 1

In 2024, several lobbyists from the same firm represented both OpenAI and the Center for AI Safety Action Fund at the same time. I am not suggesting any conflict of interest on their part. However, I don't think the "299 corporate lobbyists vs. scrappy AI safety groups" is an effective framing, given that at least some of the money is flowing to the similar places.

I wouldn't compare external lobbyists to full time advocacy staff because (1) external lobbyists and lawyers typically cover many clients and are unlikely to be particularly committed to a single issue like AI safety, and (2) firms often register anyone who does outreach on the project, regardless of whether they meet federal lobbying disclosure thresholds. It's also pretty congenial in general; Amazon lobbyists have helped me with EA-related side projects without payment because they were being nice.

This isn't to say that advocacy couldn't absorb more funding. But the conflict framing doesn't seem to represent the facts on the ground, at if organizations want funding to hire mainstream government relations people who rarely see it that way. Upskilling people already committed to AI safety would be different.

You run an AI safety org full time and have a better idea of the field. I'm just throwing in my two cents re: representation disparities.

Jason Green-Lowe 11mo 4 1 0

So the 299 corporate lobbyists is less about measuring total influence and more about how many pairs of eyes there are on the playing field who are potentially able to notice a bill moving forward -- the odds that all 299 of them miss the bill for months on end are essentially zero.

You're right to be skeptical of "number of lobbyists" as a measure of influence; a better metric would be the total amount of money spent on in-house government relations experts, outside lobbyists, advertising, PR firms, social media experts, and campaign donations. I don't have access to those figures for tech companies, but I still feel confident that the total industry budget for DC influence is much higher than the total AI safety budget for DC influence, especially if we discount

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

if they're currently being employed to work against AI safety, then I would think we're in conflict with them. Do you see it differently? What kinds of conflicts (if any) do you see in the political arena, and how do you think about them?

LeahC 11mo 1 0 0

I saw this and I agree with your main points. I will be offline for a bit due to travel, but I am happy to have a longer conversation with more nuanced responses.

Policy teams at private companies are more well-resourced while, as you mentioned, working on issues ranging from antitrust to privacy and child protection. I may be wrong, but I think the teams focused specifically on frontier AI (excluding infrastructure work) seem to be more balanced than the provided numbers suggest. This observation may be outdated, especially since SB-1047. You likely have a better idea of the current landscape than I do, and I'll defer to your assessment.

Regarding "conflict framing" - I should have phrased this differently. I did not mean the policy conflicts which come up when a new or potentially consequential industry is facing government intervention. I meant a situation when groups and individuals become entrenched in direct conflict on almost all issues, regardless of the consequence. A recent non-AI example would be the philanthropically funded anti-fossil fuel advocates fighting carbon capture projects despite the IRA funding and general support from climate change-focused groups. The conflict has moved beyond specific policy proposals or even climate goals and has become a purity test that seems impossible to overcome through negotiation. This is a situation that I would *not* want to see, and I am glad it is not the case here.

Jason Green-Lowe 11mo 1 0 0

OK, let me know when you're back, and I'll be happy to chat more! You can also email me at jason@aipolicy.us if you like.

SummaryBot 11mo 2 0 0

Executive summary: This post argues that the AI governance community must urgently shift resources from research to advocacy, asserting that sufficient understanding already exists to support beneficial policies, and that delay—whether due to uncertainty, strategic caution, or capacity concerns—risks squandering a narrow and fast-closing window to meaningfully influence AI development before superintelligence or entrenched industry power makes regulation infeasible.

Key points:

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

2. Common objections to immediate advocacy are weak: Arguments that we lack robust policies, skilled advocates, or sufficient political influence underestimate our readiness and the diminishing returns of further delay.
3. Basic safety policies are clearly beneficial and feasible: Measures like audits, whistleblower protections, and liability mechanisms offer positive expected value and are analogized to historically successful safety interventions like seatbelts.
4. Regulation fears (e.g., backfiring, oligopoly, nationalization) are overstated: Thoughtfully crafted policies can avoid unintended consequences, and the post rebuts concerns around driving companies offshore, excessive cost burdens, regulatory capture, and governmental overreach.
5. Advocacy skills and infrastructure must be built through action: Effective political influence requires on-the-ground experience, relationship-building, and training—none of which can be developed passively through more research.
6. There's no time to wait for ideal conditions: Given long legislative timelines and the accelerating pace of AI development, waiting for more favorable advocacy conditions may render future efforts moot as corporate power and technological capability outstrip regulatory leverage.

This comment was auto-generated by the EA Forum Team. Feel free to point out issues with this summary by replying to the comment, and contact us if you have feedback.

More from Jason Green-Lowe

141 **Please Donate to CAIP (Post 1 of 7 on AI Governance)**
Jason Green-Lowe · 1y ago · 40m read
25

61 **The Need for Political Advertising (Post 2 of 7 on AI Governance)**
Jason Green-Lowe · 1y ago · 16m read
0

53 **Mainstream Grantmaking Expertise (Post 7 of 7 on AI Governance)**
Jason Green-Lowe · 10mo ago · 45m read
2

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Curated and popular this week

Impact Requires Implementation: What Dispensers for Safe Water Scale-Back Can Teach Us About Monitoring

Tony Senanayake · 13d ago · 17m read · 3

TL;DR Between 2022 and 2023, GiveWell and Open Philanthropy committed over $70 million to Evidence Action's Dispensers for Safe Water programme. An independent survey in 2025 found that in Uganda and Malawi the programme was reaching approximately 70% fewer people with detectable...

The AI people have been right a lot

Dylan Matthews · 4d ago · 7m read · 11

This post was crossposted from Dylan Matthew's blog by the EA Forum team. The author may not see or reply to comments. ---------- Subtitle: Try to keep an open mind as the world gets...

In Development Magazine Launch

Lauren Gilbert · 4d ago · 1m read · 4

We're announcing a new development magazine: In Development. Our goal is to promote discourse on development without dogma. In other words, we public stories doing deep dives into pressing development issues, backed up by evidence, argument, and experience. The first piece is by Paul...

Recent opportunities in AI safety

80 **Help me launch Obsolete: a book aimed at building a new movement for ...**
Garrison · 11d ago · 8m read · 4

11 **Pivotal Research Fellowship applications are open (deadline May 3)**
Tobias Häberli, Tilman, James Nicholas Bryant, Ariel · 13h ago · 3m read · 0

12 **AI Can Already Evade DNA Synthesis Screening. Congress's New Bil...**
Sophie Kim · 22d ago · 7m read · 0

View more

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.