AN ACTIVIST VIEW OF AI GOVERNANCE

Orphaned Policies (Post 5 of 7 on AI Governance)

46

by Jason Green-Lowe May 29 2025

AI safety | Building effective altruism | AI governance | Criticism of effective altruist causes | Public communication on AI safety | Research agendas, questions, and project lists | Frontpage

Show all topics

In previous posts in this sequence, I laid out a case for why most AI governance research is too academic and too abstract to have much influence over the future. Politics is noisy and contested, so we can't expect that good AI governance ideas will spread on their own – we need a large team of people who are actively promoting those ideas. Unfortunately, we currently have at least 3 researchers for every advocate, so many policy ideas have been "orphaned," i.e., nobody is taking those ideas and showing them to decision-makers who have the power to implement them.

The best way of addressing this imbalance would be to shift funding and jobs from research to advocacy. However, as a practical matter, I don't expect many funders to heed my arguments, nor do I expect many researchers to spontaneously quit and look for new jobs in other fields. So, in this post, I offer tips on how researchers can make their work more relevant to advocacy by drafting actual policy documents and by making concrete proposals in their white papers.

I also catalog eleven orphaned policies that I'm aware of and make suggestions about how you can "adopt" them. This isn't meant to be a perfect list -- there are good orphaned policies that I haven't included, and I might have included one or two policies that you find unimpressive. My hope is that collectively, these policies illustrate the size and scope of the backlog.

The important thing isn't that we agree on exactly which policy is best; the important

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your cookie settings for more control, or review our cookie policy here.

Reject Accept all

DRAFT ACTUAL POLICY DOCUMENTS

The most important thing individual researchers can do to make their work more relevant is to shift some of their efforts from general academic exploration to ***drafting actual policy documents.*** We have so many good policy proposals that have never gotten beyond the idea stage, and that need to be fleshed out.

What I mean by "fleshing out" a policy proposal is to actually draft an example of that proposal, ideally with specific, well-justified numbers, proper nouns, and mechanisms. Don't just say that we ought to have compute monitoring; write a bill that implements it. Don't just propose a windfall profits clause; write the corporate governance documents that would implement it. This is what CAIP did with our model legislation.

This is not as hard as you might think – there are $50 textbooks that will teach you the appropriate writing style for legislation or for corporate charters, there are free how-to guides available online, you can poach other companies' documents to use as templates, and you can ask LLMs to review your work and offer advice about where you might need to make edits. Policymakers will also be forgiving if you make small formatting errors, especially since they often have professional help for such editing. For example, when Congressional staffers draft bills, those bills usually get sent to Congress's Office of Legislative Counsel, which puts the finishing touches on them.

One of the advantages of drafting actual policy documents is that they're much shorter than a typical academic paper – you can have a very helpful bill that's only three pages long. This makes it easier to get feedback, and you should ideally be getting feedback from a wide variety of sources (and make appropriate improvements based on that feedback) before you show your policies to neutral policymakers.

Another advantage to drafting actual policies is that it forces you to make choices about who should do what and how they should do it. This helps you make sure that your ideas are politically feasible and logistically sound.

For example, one of the most frequent questions CAIP got when discussing our model legislation was "which Cabinet department should run your proposed new office?" This

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

safety enforcement efforts. This is not something you'd pick up from doing a literature review, because NIST's preferences aren't cited much in the academic literature. Instead, we caught this issue because drafting actual legislation prompted us to think about where the office should go, and because we talked to NIST employees to get their perspective. We eventually concluded that the Commerce Department would need to create a new regulatory office, which is exactly the solution that Senator Romney proposed in December 2024.

MAKE YOUR WHITE PAPERS SPECIFIC

If you're not going to draft an actual policy, you can still be of some help by writing a white paper that goes into enough concrete detail about what a policy should look like. The format of the document is less crucial than the fact that it (1) contains concrete details, and (2) takes a stand in favor of a specific outcome or makes a specific prediction.

Instead of discussing the pros and cons and desiderata of various categories of proposals, you need to be identifying one proposal you think is best and explaining why that proposal is best. Instead of saying "the government could require owners of advanced AI chips to register their location," a concrete policy proposal would say "the Department of Energy should hire three full-time agents to collect and analyze annual reports on the location of all H100 and A100 chips."

The latter proposal requires more work, especially if you're worried about getting the details right – and that's exactly the point. As discussed in the second post in this sequence, policy proposals need to be easy for politicians to implement, or else they'll usually move on to something that's easier. Politicians are insanely busy and overscheduled. Every minute you ask them to spend sifting through policy alternatives and figuring out the details of how a hypothetical bill should work is a minute that they can't spend pushing their colleagues to co-sponsor and vote for a bill that you've already written.

As a result, you can get much more traction with Congressional staffers by saying "here's the policy we think you should champion" than by saying "gosh, these issues seem like

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Admittedly, most research work is funded by 501(c)(3) donations that cannot pay for more than a small amount of direct political advocacy. However, there are ways to word the conclusion of a research paper that provide clear guidance without crossing the line into inappropriate political work. True, you might not endorse a bill that's currently being debated by Congress, and you certainly shouldn't be endorsing political *candidates*, but you can still truthfully say that one type of policy has better consequences than another. The law clearly states that "nonpartisan analysis, study, or research may advocate a particular position or viewpoint so long as there is a sufficiently full and fair exposition of the pertinent facts to enable the public or an individual to form an independent opinion or conclusion."

It's therefore not "political" to point out that banning A100 chip exports while permitting A800 chip exports is ineffective; that's a technical conclusion that a neutral researcher can reasonably draw. It's not "advocacy" to express an opinion that the next generation of LLMs will most likely uplift the capabilities of bioweapon designers to a degree that poses risks that would be considered unacceptable in any other industry. You can have a firm opinion on an issue without being a politician; nothing about having a 501(c)(3) tax-status requires you to drown every opinion you offer in a sea of "maybe" and "could" and "warrants further research." The point of tax-exempt research is to come up with scientifically informed opinions that politicians can draw on to inform their work; if your organization is too timid to firmly express those opinions, then it's not upholding its part of the social contract.

A good example of an adequately specific white paper is FAS's *An Early Warning System for AI*. They've got a budget, an estimate of the number of staff needed, an administrative home for the new office, and a step-by-step explanation of who would do what and how it would work. True, the paper doesn't literally include the text of model legislation – but that's fine, because the paper provides enough detail that a legislator who was convinced by the paper's arguments could easily turn the paper over to the Office of Legislative Counsel and have them translate it into a bill.

CATALOG OF ORPHANED POLICIES

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

additional work on one of these policies and I've overlooked it, please let me know, and I will update the list accordingly!

Windfall Profits Clause

GovAI proposed in 2020 that AI firms should commit ahead of time to redistribute most of the profits of transformative AI from its inventor to the rest of humanity. Their paper includes a possible series of 'tax brackets,' which is a useful detail, but they do not include sample language showing how to add a windfall profits clause to a corporate charter or corporate bylaws.

A windfall profits clause is an excellent idea, but to date, no one has drafted a sample windfall profits clause, let alone tried to persuade any particular corporation to adopt one. You can help by figuring out which corporate document(s) need to be amended, drafting an amendment that would have the appropriate effect, and then writing a letter to corporate social responsibility officers asking them whether they will pass that amendment, and, if not, why not.

Antitrust Waiver

LawAI proposed in 2021 that the government should issue a waiver to AI developers promising not to prosecute them under the antitrust laws for meeting to discuss minimum safety standards. Competitors are normally not allowed to meet with each other to agree on changes to their business practices, but such a meeting ***would*** be allowed if there were an explicit government waiver, or if the meeting was hosted by a *bona fide* trade association that has non-commercial purposes.

The consensus seems to be that companies are not actually that afraid of being prosecuted for antitrust violations based on negotiating industry-wide minimum safety standards, but because some companies (e.g. Google) are being actively investigated or prosecuted for other antitrust violations, the fear that such talks could endanger the parent corporation is a convenient excuse that can be used by executives to tell an engineer not to pursue safety agreements. Getting a waiver granted could eliminate that excuse, making such safety

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Attorney General and asked them to sign it.

It looks like the Frontier Model Forum is at least a plausible candidate for a non-commercial trade association that could serve as a protected forum for hashing out safety agreements, but it's not clear whether this forum has the desire or ability to help companies enter into binding negotiations, rather than just identify voluntary best practices.

You can help by writing the waiver letter or by investigating the Frontier Model Forum and seeing what if anything else they need in order to be more active in setting safety standards for existential risk from AI.

Strict Liability

Applying a less forgiving set of tort law to harms caused by AI has been discussed for several years; the Brookings Institute proposed using products liability law in 2019, and CAIP board member Gabriel Weil published a detailed analysis of several possible liability reforms in January 2024. California SB 1047 would have made some minor changes or clarifications to existing tort law, and a few other state legislatures have also considered modifications.

However, this should ideally be a 50-state project. Specific language in every state should be available for legislatures to adopt, and we should also be filing impact litigation that would give judges a chance to incorporate strict liability for AI into the common law.

You can help by writing a strict liability law for your state and by submitting comments or articles to a local law review journal arguing in favor of strict liability. Such laws are often more likely to pass (or be adopted by judges) when they have some academic support, but the support needs to be registered inside the legal community for it to be noticed.

Visa Reform

Changes to immigration policy that made it easier for AI researchers to move to the US could reduce the urgency of an 'arms race'. Especially if you think that other countries will

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

slowdowns would no longer seem to carry as sharp of a risk of losing an arms race. To date, no one has drafted a bill or regulation to allow for more visas for this purpose, let alone lobbied Congress or USCIS to adopt it.

You can help by figuring out which specific kinds of people need to be welcomed into the country, what kinds of visas they would need to be eligible for, who could grant those visas, what program or authority they could be granted under, and who should be in charge of making this happen. Ideally, those conclusions would then be written up as a proposal for an appropriate officer at the State Department and/or the Department of Homeland Security. You might also write to an appropriate official at the White House who works on immigration policy.

Because the current administration is relatively skeptical of immigrants, it will be important to narrowly tailor the visa reform to cover only the people who we need to have here in America in order to widen our AI tech lead against rivals like China. It would also be helpful to include policies that will reduce the risk that people admitted under the new program will conduct espionage or otherwise transfer tech to rival countries.

Insurance Requirements

A team of researchers led by the AI Objectives Institute recently put out a paper praising the benefits of AI safety insurance, arguing persuasively that "insurance has the potential to create a more favorable incentive structure by making practices such as safety-washing or underestimating AI-related risks less appealing." CAIP's co-founder, Thomas Larsen, was a strong proponent of requiring frontier AI developers to carry a minimum amount of insurance.

However, we never satisfactorily answered the question of what this minimum amount should be. What are a reasonable set of policy limits? How large can the deductible be? How large can co-insurance payments be? What is the scope of harms that would be covered by such policies, and what if any exclusions would be permitted? What kinds of re-insurance requirements would insurers have to meet to make sure that policies will be paid out even if the primary insurer is bankrupted by an unusually large claim?

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

sending the sample insurance policy to a real insurance company or to an actuary and seeing if they'd be willing to come up with a price estimate for it.

Public Grant Funding

The US federal government currently distributes half a billion dollars a year in AI research funding; some of this funding could be explicitly earmarked for AI safety. There was a $20 million allocation for safety research in 2023, but since then, no one has drafted a proposal for more such earmarks, let alone lobbied Congress or the National Science Foundation to adopt it.

You can help by figuring out how much money could usefully be absorbed by technical AI safety researchers over a variety of different timelines (e.g. 1 year, 2 years, 5 years, etc.), what kinds of programs they might work on, how these programs could be made legible to policymakers who have only a general interest in the topic (i.e., talk mostly about interpretability and reliability, not microdooms and superintelligence explosions), and how we might measure the success of such grants to avoid wasting money.

Global Crisis Hotline

Lawfare and Brookings published a thoughtful piece in 2024 calling for an "AI incidents hotline," similar to the legendary Red Phone that symbolized efforts to connect the US and the USSR to avert nuclear accidents during the Cold War.

There are a number of interesting questions that always need to be answered about such a hotline: who should staff it? What authority should they have? What kinds of incidents should the hotline be used for? What actions will the operators take when a major crisis is reported, and to whom will they delegate less critical emergencies in order to keep the line clear? What, if anything, can the US constructively do if China refuses to pick up the phone during a particular crisis?

The US currently seems to lack even a basic bilateral military hotline with China, so adding an AI-specific hotline will take a significant amount of work. If we do, we will also have to

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

To the best of my knowledge, nobody has laid out a detailed plan for how to create an appropriate international incident hotline, let alone proposed a treaty or other document that would cause one to come into being. You can help by researching how past hotlines have worked and what we can learn from their successes and failures, and then using what you learn to propose a specific plan to launch a new hotline.

Compute Monitoring

There has been much discussion of how the government could attempt to track large clusters of computing power with the goal of knowing who is doing large-scale training runs so that the government could intervene in an emergency. Yonadav Shavit's 2023 paper "What Does It Take to Catch a Chinchilla?" provides a useful amount of detail about how often inspections would need to take place, but there is still much work to be done in terms of figuring out who would do these inspections, what the penalties would be for noncompliance, and how the hardware innovations required would be paid for.

To date, there has been only limited legislative action on this topic. Senator Cotton has recently introduced the Chip Security Act, which would require export-controlled chips to have a location verification mechanism within 6 months of the bill's passage, and CAIP has sincerely endorsed it, but it's not totally clear how the location verification would work or what else might need to be done to monitor compute.

There are many details that remain to be worked out in terms of what specific hardware features could and should be placed on chips to make them easier for the government to monitor. Should advanced AI chips have GPS locators? Should they include proof-of-work features that allow others to identify what types of computations they were used on and roughly how many of those computations were performed? Should chips have a 'kill switch' that allows them to be remotely deactivated, or, more aggressively, a dead man's switch that automatically deactivates them if they do not receive the correct password at periodic intervals?

How much would it cost to develop each of these features, and how quickly could they be developed and manufactured? There are several academic papers that discuss these

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

LAWS Boycott

Groups of AI researchers have occasionally banded together to oppose the development of lethal autonomous weapons systems ("LAWS"), but there is no national or global policy that specifically bans such weapons, and weapons that are at least semi-autonomous continue to be manufactured and deployed (e.g., in Ukraine).

Section 1638 of the FY2025 National Defense Authorization Act (NDAA) included a statement that "it is the policy of the United States" to avoid letting AI compromise "the principle of requiring positive human actions" before firing nuclear weapons, but this falls slightly short of a binding legal requirement to always keep a human in the loop even for nuclear weapons, and conventional LAWS remain unregulated.

Although at least 30 countries have called for a ban on LAWS, it seems that none of them have actually written the text of a proposed treaty or a proposed new protocol to the Convention on Conventional Weapons that would create a legally-enforceable ban (for countries that sign it) on the use of LAWS. You can help by doing that work for them.

Similarly, to date, no one has written a specific policy that would make it illegal for the US military to use some or all types of LAWS – you can help by figuring out which types of LAWS would be most practical to do without and writing a policy that would require that we avoid their use. Stuart Russell has suggested that lethal autonomous weapons below 400 grams could be banned, as part of an effort to avoid the worst kinds of anti-personnel swarms while still allowing for LAWS that can destroy, e.g., enemy tanks and fighter planes. However, it's not obvious that the 400-gram figure is technically accurate; it seems to have been imported from the St. Petersburg Declaration of 1868, which dealt with exploding musket balls. It might be perfectly possible to design anti-personnel LAWS that weigh only 300 grams. You can help by updating this standard and coming up with a more technically justified red line for lethal autonomous weapon systems.

It could also be helpful to try to organize a consumer boycott against companies that are designing or manufacturing LAWS, or against companies that are doing so without adequate safeguards, or against companies that are designing specific LAWS such as anti-personnel swarms. You could help with this type of effort by making a list of which

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Industry Standards

NIST has already published an AI Risk Management Framework that contains voluntary best practices for coping with the risks of AI. It is not obvious that all major AI developers are actually complying with these voluntary best practices, or even that they have promised to do so. There is much work to be done lobbying companies to publicly promise to abide by the NIST AI RMF and then preparing checklists, scorecards and other tools to evaluate how well they are living up to this promise.

SaferAI's work in creating a rating system for AI companies is a good first step, but the ratings need to be broadened and aligned with the NIST criteria, and someone needs to convince the companies to commit to using these criteria and to publishing enough data often enough that third parties can meaningfully assess the extent to which they are successfully complying with those criteria.

You can help by developing checklists or scorecards that assess compliance, by developing tools, wizards, charts, and templates that make it easier for companies to comply, and by drafting sample agreements that companies could sign to show that they agree to follow these guidelines.

Structured Access to Research

Researchers need access to AI models in order to assess their safety, but corporations may be reluctant to share their source code for fear of losing a competitive advantage, and publicly sharing source code may be undesirable in any case because it may accelerate general AI capabilities research. To solve this problem, some companies are offering "structured access" to their models through an API.

However, it is not clear that anyone is lobbying for this trend to continue or for other companies to adopt structured access plans. You can help by defining the minimum class of researchers who should have structured access, the minimum amount of access that they should have, and the maximum amount of usage restrictions or other requirements that companies can impose on such researchers. For example, a non-disparagement

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Once we know more about what a reasonable structured access plan would look like, you can help by drafting an open letter that companies can sign to pledge that they will always provide this access, and then sending that letter to appropriate departments at AI developers and encouraging them to sign it.

CONCLUSION

All of the topics above need more people who are working on drafting concrete policy proposals and lobbying governments and corporations to implement those proposals. I am not aware of any AI safety issues that are "saturated" with drafters or lobbyists, in the sense that they already have about as many drafters as they need and adding more staff would face diminishing marginal returns.

By contrast, the movement as a whole is already saturated with general policy ideas, in the sense that we have far more good policy ideas than we can currently hope to act on. In addition to the eleven ideas discussed in this post, there are many more good ideas summarized at the bullet-point level in response to OSTP's call for comments on its AI Action Plan. You can sift through the 10,068 official comments here, or you can read CAIP's highlights in issues #67 and #73 of our weekly newsletter.

What this means is that if you're spending your time generating even more policy ideas, that's not likely to be helpful – unless your new idea is much better than existing ideas, and that improvement is obvious and visible to the handful of advocates who are trying to get those ideas implemented, then it's profoundly unclear how your new idea is going to translate into real-world change.

If you're spending your time musing about general considerations that might possibly improve the quality of other general policy ideas that other people might come up with someday, that's even less likely to be helpful.

I have to imagine that even the researchers themselves don't want to see their research go to waste. If you dedicate your professional career to seeking out new knowledge, don't you want someone to benefit from that knowledge? Don't you want someone to,

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

policymakers to adopt. I truly believe this will be useful on the margins.

If we want to change things at scale, and not just on the margins, then the only way I can think of to make sure that most of the research we produce actually gets read by the people in charge is to shift some of our funding from research to advocacy. In my sixth and seventh posts, I'll explain what institutional changes I think would need to happen within the AI safety funding environment to make that possible.

Previous:

Shift Resources to Advocacy Now (Post 4 of 7 on AI Governance)

5 comments 63 karma

Next:

Political Funding Expertise (Post 6 of 7 on AI Governance)

1 comments 37 karma

Log in to save where you left off

Mentioned in

122 The Center for AI Policy Has Shut Down

Comments 3

Sorted by New & upvoted

Chris Leong 11mo 9 2 0

Great post. Listing concrete examples of orphaned policies makes it much easier for folks to evaluate how much of a priority drafting orphaned policies should be.

That said, my belief is that actually it's not just fine for the AI governance community to propose far more policies than policies actually drafted, but this is exactly the way that it should be.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your cookie settings for more control, or review our cookie policy here.

Here's my hot-takes (disclaimer: by hot-takes, I really do mean hot-takes and I'm not a policy professional!)

• Windfalls clause: robustly good, but not on the critical path
• Antitrust waiver: seems robustly good
• Visa reform: hard to determine the sign of due to espionage concerns
• Insurance requirements: hard to determine the sign of due to moral hazard
• Public grant funding: quite hard to make sure this goes to anything useful, that said, UK AISI is distributing grants and I'm quite optimistic about their judgement
• Global crisis hotline: seems robustly good
• Compute monitoring: seems robustly good
• LAW Boycott: hard to determine sign due to unstable equilibrium
• Industry standards: I'm a lot more pessimistic about this than most folks. Very easy for this to create a false sense of security. Unfortunately, at the end of the day, if a company doesn't care, they don't care
• Structured access to research: I suspect that the companies will either give it voluntarily or it'll not be worth the political capital to try to mandate

So 3/11 seem robustly good, with another robustly good but not on the critical path.

Question: Are there any organisations focused on taking general policy proposals and developing them into specific proposals for legislation? I could see value in having an organising specialising in this stage if the majority of governance organisations are just throwing rather general proposals over the wall and hoping someone else will fill in the details.

Jason Green-Lowe 11mo 10 2 0

I'm not aware of any such organizations! This is an example of one of the 'holes' that I'm trying to highlight in our ecosystem.

We have so many people proposing and discussing general ideas, but there's no process in place to rigorously compare those ideas to each other, choose a few of those ideas to move forward, write up legislation for the ideas that are selected, and advertise that legislation to policymakers.

I don't object to the community proposing 5-10x more ideas than it formally writes up as policies; as you say, some filtering is appropriate. I *do* object to the community spending 5-10x more time proposing ideas than it spends on drafting them. The reason why it makes sense to have lots of ideas is that proposing an idea is (or should be) quick and easy compared to the hard work of drafting it into an actual policy document. If we spend 70% of our resources on general academic discussion of ideas without anyone ever making a deliberate effort to select and promote one or two of those ideas for legislative advocacy, then something's gone badly wrong.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

Executive summary: This post argues that the AI governance field suffers from a surplus of abstract research and a shortage of advocacy, leading to a backlog of promising but unused policy ideas ("orphaned policies"); to remedy this, the author recommends that researchers make their work more actionable by drafting concrete policy documents and adopting one of eleven specific, underdeveloped proposals detailed in the post.

Key points:

1. Advocacy bottleneck: The field has a researcher-to-advocate imbalance (~3:1), meaning many good AI policy ideas lack champions who can bring them to policymakers—creating a large backlog of "orphaned" proposals.
2. Drafting real policies is tractable and impactful: Researchers should draft actual legislative or governance documents, which are often shorter and more influential than academic papers, and easier for policymakers to act on.
3. Make white papers concrete and directive: Even without drafting full legislation, researchers can increase their work's utility by including specific recommendations, estimates, and implementation details.
4. Legal and funding constraints are navigable: Despite 501(c)(3) limits on lobbying, researchers can still advocate specific policies if framed as nonpartisan analysis with a clear evidentiary basis.
5. Catalog of eleven underdeveloped ideas: The post outlines eleven "orphaned" policies—ranging from compute monitoring and AI insurance to visa reform and LAWS regulation—each with specific missing components that a researcher could fill in.
6. Call to action and teaser for next post: The author urges researchers to "adopt" a policy and help develop it for real-world use, and suggests broader institutional reforms (to be discussed in the final post) are needed to shift funding from research to advocacy.

This comment was auto-generated by the EA Forum Team. Feel free to point out issues with this summary by replying to the comment, and contact us if you have feedback.

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

63 **Shift Resources to Advocacy Now (Post 4 of 7 on AI Governance)**
Jason Green-Lowe · 11mo ago · 38m read
5

61 **The Need for Political Advertising (Post 2 of 7 on AI Governance)**
Jason Green-Lowe · 1y ago · 16m read
0

View more

Curated and popular this week

Impact Requires Implementation: What Dispensers for Safe Water Scale-Back Can Teach Us About Monitoring

Tony Senanayake · 13d ago · 17m read · 3

TL;DR Between 2022 and 2023, GiveWell and Open Philanthropy committed over $70 million to Evidence Action's Dispensers for Safe Water programme. An independent survey in 2025 found that in Uganda and Malawi the programme was reaching approximately 70% fewer people with detectable...

The AI people have been right a lot

Dylan Matthews · 4d ago · 7m read · 11

This post was crossposted from Dylan Matthew's blog by the EA Forum team. The author may not see or reply to comments. ---------- Subtitle: Try to keep an open mind as the world gets...

In Development Magazine Launch

Lauren Gilbert · 4d ago · 1m read · 4

We're announcing a new development magazine: In Development. Our goal is to promote discourse on development without dogma. In other words, we public stories doing deep dives into pressing development issues, backed up by evidence, argument, and experience. The first piece is by Paul...

Recent opportunities in AI safety

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.

10

AIxBio Hackathon - April 24-26, $4,000 in prizes (online + SF + Berlin)

Kamil Alaa · 15d ago · 1m read

0

View more

We and our partners use cookies, including to review how our site is used and to improve our site's performance. By clicking "Accept all" you agree to their use. Customise your **cookie settings** for more control, or review our cookie policy **here**.