Political Funding Expertise (Post 6 of 7 on AI Governance)

37

by Jason Green-Lowe Jun 19 2025

AI safety | AI governance | Frontpage

INTRODUCTION

The Story So Far

In my first three posts in this sequence, I estimated that we have 3 researchers for every advocate working on US AI governance, and I argued that this ratio is backwards. Even the best research offers only modest and indirect support for advocacy, and research alone has negligible political power. Without political power, we can't change the bad incentives of AI developers that are very likely to lead to the collapse of human civilization.

In the fourth post of this sequence, I acknowledged that some amount of initial research to lay out the philosophical foundations of a new field might be needed before advocacy can begin, but I also showed that this initial research has been thoroughly completed. We know why unregulated AI is bad and we know at least some harmless ways that we can make progress toward fixing that problem.

In the fifth post of this sequence, I illustrated that point by listing eleven examples of 'orphaned' policies. Each of these policies was proposed by academic researchers – often several years ago – but the policies have not been drafted in any detail by policy wonks, let alone presented to decision-makers by political advocates. We clearly have an over-supply of academic ideas and an under-supply of political elbow grease.

This means that it's very strange that so much of the AI safety movement's funding has gone toward academic-style research and is continuing to go towards more research. Presumably, the funders can see the same trends that I've laid out in this sequence: they

should know as well as I do that funding more researchers than advocates is deeply suboptimal.

So, why do they keep doing it? My best guess is that they're biased by their own backgrounds: the staff of major AI safety funding organizations are overwhelmingly drawn from academic-style research environments. This may be causing them to fund other researchers even when that's not strategically optimal, simply because they're more comfortable with research or they are better able to understand the benefits of research.

In this sixth post, I will argue that rationalist and effective altruist funders should repair this flaw in their staffing so that they can more accurately evaluate the usefulness of future grant proposals. To its credit, Open Philanthropy hired an actual governance expert in April 2025, shortly after it finalized its decision not to fund CAIP. Similarly, in June 2025, Longview Philanthropy advertised an opening for an AI policy expert. However, these efforts are too modest and too recent to fully address the problem. It takes more than one or two experts to adequately evaluate an entire field's worth of advocacy proposals, and grant-makers should have been aware of this problem and taken action to address it well before 2025. It is deeply concerning that with the politics of AI regulation making front-page headlines all through 2023, the movement's grant makers are only just now hiring people who are qualified to evaluate grant applications to conduct political advocacy.

Response From AI Safety Grantmakers

I shared a draft of this post with all five of the major AI safety grantmakers (Open Phil, Longview, Macroscopic, LTFF, and SFF) a week before it went live to allow them a chance to respond to the criticisms made here. The general nature of this post has also been public for some time, since the table of contents for the entire sequence went up on May 22, 2025. Several of these grantmakers helpfully and graciously engaged with the draft by pointing out factual errors and misunderstandings; I am grateful to them for their thoughtful comments, and I believe that this post is stronger because of their feedback.

Nevertheless, I have not received any information from the grantmakers that satisfactorily addresses the core problems I am discussing here. To the extent that the lack of experienced advocacy staff among the major AI safety funders was an accidental oversight, I would like the AI safety movement's grantmakers to explain how they will avoid similar mistakes in the future. To the extent that this was an intentional decision, I

would like an explanation of why funders feel that they can evaluate political grantees without having political experts on staff.

So far, neither explanation has been provided. Until that changes, I would urge private donors to consider making their own grant evaluations going forward for political projects, rather than deferring to the judgment of large institutional donors like Open Philanthropy – if they cannot or will not explain their methods, then it is not obvious that their judgments are worthy of deference.

I will go into more detail about this last point in my seventh and final post, which shifts gears from a lack of ***political advocacy*** experience among AI safety grantmakers to a lack of ***mainstream philanthropic*** experience among AI safety grantmakers. In the seventh post, I will argue that although EA grantmakers are unusually good at choosing ***which*** cause areas to fund, they have no special advantage at choosing which grants to fund ***within*** a given cause area. If anything, EA grantmakers are likely to be somewhat worse than average at choosing specific grants within a cause area, because they have not hired people who are familiar with the best practices for this particular task, and as a result their grant evaluation methods appear to be excessively informal.

AI SAFETY GRANTMAKERS HAVE VERY LITTLE POLITICAL EXPERIENCE

Context, Purpose, and Methodology

My purpose in reviewing the credentials of the people discussed in this post is to demonstrate that the AI safety grantmaking ecosystem as a whole is seriously lacking in political advocacy experience. I cannot figure out how to make this point effectively without commenting on the experience levels of individual people, but this section is ***not*** intended as an attack on the usefulness of any particular person, much less as an attack on their character. On the contrary, my experience with the handful of people whom I have worked with from this list is that all of them are thoughtful, intelligent, hard-working, and well-intentioned.

However, when it comes time to evaluate whether the AI safety grantmaking field ***as a whole*** is well-equipped to discharge its responsibilities, these positive personal qualities do not outweigh a profound lack of relevant experience. Each individual person can be good

and talented and helpful, but if none of them are professional advocates, then as a group they will make predictable and avoidable mistakes when they try to evaluate advocacy proposals. No amount of purely academic understanding can fully compensate for the missing personal experience of how politics and government operate on a day-to-day level.

In my own role as the executive director of an advocacy organization, I often had to defer to the superior experience of the government relations and media relations experts who I hired, because they knew more about their fields than I did. I have a degree in Political Science from Yale and in Law from Harvard, and I've worked for several years in fields that have some relevance to CAIP's work, such as product safety litigation and nonprofit regulatory counseling. I have interned for a Senator and for the Justice Department. Nevertheless, I knew that I did not know enough to set myself up as the sole judge of which advocacy tactics would be effective, and so I hired true professionals who each had at least a decade of experience working on Capitol Hill and then regularly sought and followed their advice. It is my hope and my expectation that large AI safety funders will demonstrate this same humility.

To emphasize the fact that I am not intending to criticize any particular person's credentials, I have moved all of the individual reviews out of this post and into a separate Google Doc. Moreover, even within this separate Google Doc, individuals are not directly named – instead, each entry in the census links to that person's bio on their institution's website. This should make it harder for individual people's names to appear in connection with this review as part of, e.g., a Google search, and easier for people to look past the records of individual people to see the broader point that I am making about staffing ratios and missing expertise.

In order to assess the professional expertise of AI safety grantmakers, I am relying primarily on their public LinkedIn profiles and their public bios from their institutional websites. I may be unaware of or have overlooked some relevant professional experience. If you can add additional information about someone's political or philanthropic expertise that I have not mentioned here, please do so in the comments or by emailing me at jason@aipolicy.us, and I will be very happy to correct the record.

For me to classify someone as a political advocacy expert, they must have the equivalent of several years of full-time, post-college work experience in politics or government, including at least one experience where they were directly engaged in trying to persuade

other people to adopt a political position. If you have read the rest of this sequence, you will not be surprised to hear that I am not including time spent working as an AI governance researcher toward the definition of "advocacy expertise." My fundamental point is that academic-style research about policy options – even if it is at a think tank that happens to be located in DC – does not use the same skills or involve the same problems as direct advocacy for policy change.

If it is unclear whether someone is a political advocacy expert, or if it seems clear that they have spent about two to four years working on political advocacy, then I classify them as a "close call" and count them as the equivalent of 50% of a full political advocacy expert.

If someone has worked at a few political internships or on a few political side projects that seem to add up to roughly one year of full-time equivalent experience (i.e., roughly 2,000 hours of total work), then I list them as having "minor or incidental advocacy experience." I count each of these people as 15% of a full political advocacy expert, on the theory that they've spent less than one-seventh as much time working in that field compared to a mid-career professional whose primary career has been in political advocacy.

In the census, I describe several people as having work experience that primarily involves "research" or "academic-style research." The point of this category is not that these people are literally professors affiliated with a university. Rather, my point is that they seem to be interacting with the world primarily by reading and writing long documents that are usually abstract, conceptual, and/or exploratory. In other words, the questions they ask and the tasks they fill their day with are similar to the questions and tasks you'd see in academia. I also include a few recent graduates in this category, on the theory that most of their experience has also been colored by academia. I believe both researchers and recent graduates have a relatively distinct approach to the world, and that this approach tends to bias these grantmakers toward funding other researchers.

The scope of this census is meant to include all of the people who appear to work on AI safety grantmaking at Open Philanthropy, Longview Philanthropy, Macroscopic Ventures, the Long-Term Future Fund, and the Survival and Flourishing Fund. Together, these organizations account for the vast majority of AI safety funding. When enough information is available to do so, I focus on the AI governance staff within each organization (as opposed to technical AI safety research). If there is not enough public information to tell which staff work on which grants, I review the credentials of all of the publicly listed employees who appear to have any role in evaluating grants. Some grant

recommenders (e.g. at the Survival and Flourishing Fund) are not publicly listed; I have not attempted to find out or report anything about the credentials of anonymous recommenders.

Summary of Results

Based on my census, I count 1 person who definitely has political advocacy expertise, 5 people who might have political advocacy expertise, and 10 people who have some incidental political experience. Counting the close calls at 50% and the people with an incidental background at 15%, you could say that after Longview completes its current hiring round, we will have the equivalent of 1 + 2.5 + 1.5 = 5 political experts in AI safety grantmaking organizations.

Using that same census, I count 15 people whose background is primarily in academic-style research. To these 15 people, I add the equivalent of 2.5 more researchers based on the 5 'close calls.' I think this makes sense because the people named in the 'close calls' section have also done significant research work, and because the positions that Longview is hiring for might turn out to be filled by people whose primary background is in research. Similarly, of the 10 people who I classify as having incidental political exposure, 3 of them seem to have a significant background in academic-style research: Person K, Person L, and Person M. I therefore count each of them as 50% of an academic researcher. Adding the two and a half academic equivalents from the "close calls" section to the one and a half academic equivalent from the "minor political experience" section to the fifteen people from the "academic researchers" section means that we have the equivalent of 15 + 2.5 + 1.5 = 19 researchers in AI safety grantmaking organizations.

19 divided by 5 is 3.8 – so **we have almost 4 academic researchers for every advocacy expert on the largest AI safety grantmaking teams.**

A Note on Informal Consultations

I am aware that most of these funders at least occasionally consult with people who they see as outside political experts as part of their grant evaluation process. However, there are two problems with these informal consultations, and because of these problems, I think it's still valid to say that academic researchers are about 4 times as influential as advocacy experts on our AI safety grantmaking teams.

First, I believe that at least some of the people being consulted are not true experts in the field. Based on private conversations I have had, many (although not all) of the outside consultants seem to be, e.g., first-time Congressional fellows who have only recently moved to DC. People who are relatively young or in relatively junior roles might not be given enough access to candid conversations among powerful actors to accurately assess what is about to happen next on Capitol Hill. Moreover, they might not have seen a wide enough variety of circumstances in DC to have fully calibrated their predictions about how politicians will behave. For instance, an intern who has only ever seen how Congress behaves under a divided government might not be able to accurately predict how Congress will behave when one party has unified control.

Because there are very few EA-affiliated senior political advocates, these more junior advocates are the people who are most likely to have close personal connections with staff at EA's institutional donors. I acknowledge that when advisors disagree, grantmakers will *somewhat* discount the advice of these more junior advocates based on their relative inexperience. However, in the absence of any formal procedure for deciding how much weight to place on each consultant's advice, I am concerned that grantmaking staff will nevertheless often choose (perhaps subconsciously) to put excessive weight on personal friendships. As I will discuss in more detail in my seventh and final post, I believe it is too easy for grantmakers to value the advice of relatively inexperienced friends who work closely with them over the advice of deeply experienced experts who are seen as outsiders.

This ties into the second problem: there is not enough institutional incentive for grantmakers to ***heed*** the advice of the people being consulted. Being able to truthfully say that such people were consulted provides a reputational benefit to the grantmaking organization whether or not their advice is followed, and it is difficult for outsiders to observe how much influence any one consultant has on a final funding decision. There does not appear to be any formal process that would require grantmakers to make a note of instances when they are rejecting expert advice and explain their reasons for doing so.

If managers or other influential people inside a grantmaking organization prefer to fund academic research, and an informal outside consultant prefers to fund direct advocacy, then recommending that the manager's advice be disregarded in favor of the consultant's advice will often be costly in terms of internal office politics. Overruling the manager's preferences may have negative consequences for the grantmaker's career, or even for their social life, since many of the people who work in these funding organizations and in these research organizations are friendly and spend time together in social contexts. By

contrast, ignoring the consultant's preferences is psychologically painless, since grantmakers do not have to work alongside the consultant on a daily basis.

We cannot rely on feedback from reality to pressure grantmakers into following expert advice, because AI safety governance has an essentially binary set of outcomes that do not lend themselves to rapid feedback: either civilization will end in an AI apocalypse, or it will not. By the time we observe whether AI governance grants have been successful, it will be too late to change course. As I will discuss in more detail in my seventh post, we might be able to strengthen this feedback loop by aggressively collecting and analyzing more data, but the process for doing so has not yet been put into place.

This is one of many reasons why I recommend that we should bring as many political advocacy experts as possible *inside* the grantmaking organizations as full-time employees. If people with political expertise had full social 'status' within a grantmaking organization, then that would help reduce the artificial discrepancy in influence between research experts and political experts, making it easier to heed whichever advice is most objectively compelling.

WHY THE LACK OF POLITICAL EXPERIENCE MATTERS

I can't read the minds of grantmakers or find out exactly why they made their funding decisions; the reasoning behind such decisions is typically not made public and is often not even provided in full to grantees. However, **this census seems like a very promising explanation for why AI safety organizations are funding so much academic research: it's probably because they're largely staffed with academic researchers.**

There are about four times as many academic researchers as there are advocacy experts, and they have collectively decided to fund about three times as many academic research grants as they have political advocacy grants. It's human nature to like and cooperate with people who share your professional background. Even if these grantmakers are trying their best to be unbiased, they might have a strong tendency to appreciate and understand arguments in favor of funding research, while having a strong tendency to misunderstand or discount arguments in favor of funding advocacy.

As Open Philanthropy itself conceded in 2016, their approach "can lead to a risk" that many of their grant decisions will be made by "an intellectually insulated set of people who reinforce each others' views, without bringing needed alternative perspectives and counterarguments." These grantmakers are often personally friendly with the people they are choosing to fund, so Open Philanthropy acknowledged that "it sometimes happens that it's difficult to disentangle the case for a grant from the relationships around it. When these situations occur, there's a greatly elevated risk that we aren't being objective, and aren't weighing the available evidence and arguments reasonably."

Open Philanthropy would argue that their network is larger and more diversified than it was in 2016. However, in my opinion, the risk they warned of in 2016 has now come to pass in the context of AI safety grantmaking: they aren't objectively weighing the need to fund direct advocacy, because too much of their network is caught up in academic-style research.

Similarly, a fund manager for LTFF describes° what he sees as an "equilibrium" between the need to make use of local information and the need to avoid conflicts of interest that feature "a) definite recusal for romantic relationships, b) very likely recusal for employment or housing relationships, c) probable recusal for close friends, d) disclosure but no self-recusal by default for other relationships." Another way of framing this norm is that most of the time, when the main reason that two researchers know each other is because they work together on publishing papers that they each find valuable, they are comfortable making official recommendations that institutional donors should fund each other's work. Naturally, this would tend to increase the strength of the internal EA research bubble and make it harder for them to fairly consider non-research grant proposals.

I've done my best earlier in this sequence to rule out all of the "objective" reasons for preferring to fund research. It seems clear to me that funding more research than advocacy is an objectively suboptimal strategy for achieving the goal of preventing an AI disaster. Thus, I am driven to the conclusion that the movement's preference for funding academic research is primarily a "subjective" preference, i.e., that it is based on the personality and identity and relationships of the people making the funding decisions, and not on any rational grounds.

This conclusion seems especially compelling to me because funders have been unable or unwilling to explain precisely why they judge research programs to be more valuable than specific advocacy projects, such as the Center for AI Policy. As I will discuss in the final

post of this sequence, there does not seem to be any formal metric for evaluating the efficacy of AI advocacy projects or comparing their efficacy to that of research projects.

If the funders are subjectively biased toward funding research, then that would be very bad, because it would mean that we're taking scarce funds that could be used to try to avert the collapse of civilization, and instead spending it on projects that make funders feel comfortable. To help correct this bias, we should add additional advocacy expertise within the funding organizations themselves, so that advocacy projects will no longer suffer from an implicit handicap based on the average preferences of the grantmakers.

One grant manager who reviewed this post noted that, ironically, airing public complaints like this one can make it harder for EA grantmakers to recruit politically experienced grantmakers. I admit that this is a possibility, and that is part of why I have screened the individual reviews of credentials behind two different layers of hyperlinks. I hope and expect that any minor harm that might be done to recruiting efforts by this relatively mild and anonymized criticism will be strongly outweighed by the positive effect of building up a critical mass of political experts within AI safety grantmaking teams.

One of the most important obstacles to hiring political experts to join EA-affiliated grantmaking teams is that the EA movement has a reputation for being politically unsavvy and for staffing its grantmaking teams almost exclusively with academic-style researchers. The best ways to make EA grantmaking teams a more attractive place for politicians to work are to (a) hire more politicians, and (b) make it clear through job descriptions, outreach campaigns, website content, social media, etc. that politicians are affirmatively welcomed.

Thus, I urge AI safety grantmakers to aggressively recruit as many political advocacy experts as possible. Concretely, this means that Longview should make sure that its new positions are filled by people with an advocacy background, as opposed to people with a background in AI governance research. SFF is above average among AI safety funders in terms of its general political expertise, but they should try to recruit at least one grant recommender who is deeply familiar with DC politics in particular. Open Philanthropy should be posting at least two new positions for AI governance advocacy experts in addition to the one it has already hired, and LTFF and Macroscopic should each be posting at least one new position for AI governance advocacy experts. All of these openings should be widely advertised in places where advocates will see them, such as the Public Interest

Tech Job Board, Daybook, the Internet Law & Policy Foundry, Tom Manatos Jobs, Roll Call Jobs, and the Public Affairs Council.

If you even partially agree with the argument that I've been laying out in this sequence that political advocacy is far more impactful on the margins than academic research, then it's irresponsible to continue to maintain grantmaking teams that are primarily staffed with academic researchers. The risk that this imbalance in the staffing ratio will push grantmakers toward funding too much academic research is unacceptable and unwarranted.

Previous:

Orphaned Policies (Post 5 of 7 on AI Governance)

3 comments 46 karma

Next:

Mainstream Grantmaking Expertise (Post 7 of 7 on AI Governance)

2 comments 53 karma

Log in to save where you left off

Mentioned in

122 The Center for AI Policy Has Shut Down

Comments 1

Sorted by New & upvoted

SummaryBot 10mo 4 0 0

Executive summary: This evidence-based analysis argues that AI safety grantmakers are significantly underqualified to evaluate political advocacy projects due to a strong staffing bias toward academic researchers, and calls for a strategic overhaul in hiring practices to include more professionals with direct political experience to avoid suboptimal funding decisions that could jeopardize the effectiveness of AI governance efforts.

Key points:

1. The author's census shows nearly 4 academic researchers for every 1 political advocacy expert in major AI safety grantmaking organizations, leading to a bias in funding decisions.
2. Despite clear needs and opportunities for advocacy, grantmakers disproportionately fund academic research, potentially due to their own research-oriented backgrounds rather than objective impact considerations.
3. While grantmakers occasionally consult external political experts, these consultations are informal, inconsistently influential, and often involve junior personnel, failing to substitute for in-house advocacy expertise.
4. The lack of formal procedures and incentives to balance perspectives within teams increases the likelihood of decisions based on social comfort and internal relationships rather than strategic need.
5. The author urges funders to aggressively recruit seasoned political advocacy professionals into grantmaking teams and to advertise these roles in mainstream political job markets.
6. The piece critiques broader Effective Altruism practices, warning that without reform, EA's grantmaking processes risk reinforcing epistemic bubbles and undermining high-stakes efforts like preventing AI-driven existential risks.

This comment was auto-generated by the EA Forum Team. Feel free to point out issues with this summary by replying to the comment, and contact us if you have feedback.

More from Jason Green-Lowe

141 **Please Donate to CAIP (Post 1 of 7 on AI Governance)**
Jason Green-Lowe · 1y ago · 40m read
25

63 **Shift Resources to Advocacy Now (Post 4 of 7 on AI Governance)**
Jason Green-Lowe · 11mo ago · 38m read
5

61 **The Need for Political Advertising (Post 2 of 7 on AI Governance)**
Jason Green-Lowe · 1y ago · 16m read
0

View more

Curated and popular this week

Impact Requires Implementation: What Dispensers for Safe Water Scale-Back Can Teach Us About Monitoring

Tony Senanayake · 13d ago · 17m read · 3

TL;DR Between 2022 and 2023, GiveWell and Open Philanthropy committed over $70 million to Evidence Action's Dispensers for Safe Water programme. An independent survey in 2025 found that in Uganda and Malawi the programme was reaching approximately 70% fewer people with detectable...

The AI people have been right a lot

Dylan Matthews · 4d ago · 7m read · 11

This post was crossposted from Dylan Matthew's blog by the EA Forum team. The author may not see or reply to comments. ---------- Subtitle: Try to keep an open mind as the world gets...

In Development Magazine Launch

Lauren Gilbert · 4d ago · 1m read · 4

We're announcing a new development magazine: In Development. Our goal is to promote discourse on development without dogma. In other words, we public stories doing deep dives into pressing development issues, backed up by evidence, argument, and experience. The first piece is by Paul...

Recent opportunities in AI safety

11 **Pivotal Research Fellowship applications are open (deadline May 3)**
Tobias Häberli, Tilman, James Nicholas Bryant, Ariel · 13h ago · 3m read 0

150 **AI Safety's Biggest Talent Gap Isn't Researchers. It's Generalists.**
Topaz, Agustín Covarrubias 🔶 , Alexandra Bates + 2 more · 7d ago · 7m read 9

148 **The Anthropic IPO Is Coming. We Aren't Ready for It.**
Sophie Kim · 8d ago · 9m read 16

View more