Mainstream Grantmaking Expertise (Post 7 of 7 on AI Governance)

53

by Jason Green-Lowe Jun 22 2025

AI safety | Effective giving | Existential risk | AI governance | Coefficient Giving

Show all topics

Frontpage

Previously in this sequence, I estimated that we have 3 researchers for every advocate working on US AI governance, and I argued that this ratio is backwards – we need the political power provided by advocates to have a chance of preventing misaligned superintelligence. A few researchers might be useful as a 'multiplier effect' on the power of many advocates, but the converse is not true: there's no "magic bullet" that AI governance researchers can hope to discover that could substitute for an army of political foot soldiers. Even the best political ideas still need many political activists to spread them, because the political arena is noisy and contested.

Unfortunately, we have very few political activists. This means the good ideas that our governance researchers have been proposing are mostly "orphaned." In other words, these good ideas have no one to draft them into shovel-ready language that legislators could approve, and they have no one to advocate for their passage.

It's a bit of a puzzle that the AI safety movement nevertheless continues to fund so much academic-style research. Why do we need more good ideas if we already have a dozen good ideas that nobody is acting on?

In the sixth post in this sequence, I suggested that part of the answer to this puzzle is that we have 4 academic-style thinkers for every political advocate on our AI safety grant-making teams. Because the culture within these grantmaking teams is heavily tilted toward research, grantmakers may feel more comfortable funding research grants or may be better able to appreciate the merits of research grants. They may also have a warm collegial relationship with some of the researchers, which could subtly bias their funding decisions.

Ideally, these kinds of biases would be caught and corrected by formal grantmaking procedures, such as the use of written scoring rubrics. Unfortunately, such procedures are notably missing from major AI safety funders. A week before publishing this post, I gave an advance copy of it to the five largest institutional donors so that they could address any factual errors. Although several people offered helpful corrections, none of these corrections change my fundamental point: AI safety funders are evaluating grants using an informal and subjective method. ***No formal criteria are being applied to most AI governance grant applications.***

In this seventh and final post, I explain why these formal grantmaking procedures are helpful, why they are not used by AI safety funders, and how this can be fixed.

AI SAFETY GRANTMAKERS HAVE VERY LITTLE MAINSTREAM GRANTMAKING EXPERIENCE

The Effective Altruist Bubble

Most of the people who are running AI safety grantmaking teams have significant grantmaking experience – but essentially all of that experience is from inside the effective altruist (EA) bubble. A typical career might involve doing AI safety research for a couple of years, then doing some EA field-building or running a small EA organization, then working as an EA grantmaker, and finally managing an EA grantmaking team. It's reasonable to promote most of our leaders from within our own movement, but we've gone beyond that and cultivated an extremely insular bubble. ***At no point has any AI safety grantmaker ever worked as a grantmaker at any traditional philanthropy to see how those organizations make decisions about awarding and managing grants.*** We're not just talking about a general tendency to prefer internal candidates – we're talking about a level of detachment from the mainstream philanthropic world that is likely to result in profound ignorance about mainstream practices and techniques.

You might wonder what we have to learn from traditional philanthropies. Part of the premise of effective altruism is that it's possible to do much better than traditional philanthropies by using utilitarian philosophy and rigorous quantitative reasoning: because effective altruists think hard about what it means to do good and how to measure

that, we really do select better cause areas than traditional philanthropies. I firmly approve of the way the EA movement prioritizes, e.g., shrimp welfare over adopting puppies, or prioritizes curing malaria over supporting victims of breast cancer. I would not want to undercut any of the types of reasoning that underlie these broader strategic decisions.

However, just because the EA movement is better than mainstream philanthropies at selecting ***which broad cause areas to support*** doesn't mean that everything that traditional philanthropies do is wrong or pointless. In particular, traditional philanthropies have well-developed techniques for selecting ***which individual grants to fund within a cause area***, and (as I will argue throughout this post) these techniques are often much better than the comparable techniques used by EA grantmakers. Thus, rather than trying to reinvent the entire field from scratch, we should be able to learn from mainstream philanthropy's experiences and import some of their best practices when these practices are better than ours.

I conducted a quick census of the AI safety grantmaking staff at the five major institutional funders (Open Philanthropy, Longview, Macroscopic, LTFF, and SFF) based on their LinkedIn profiles and institutional bios. As in my previous post, I have lightly anonymized these staff by offloading their individual descriptions into a separate Google Doc and referring to them as "Person 1" and "Person 2," rather than directly using their real names. If you want to see who they are, you can just click on the relevant hyperlinks, but my purpose here is to point out a deficit in the overall bank of skills held by the movement as a whole, not to criticize individual staff.

Unfortunately, as you can see from my census, there are only a few examples of people working on AI safety grantmaking who have any non-EA philanthropic experience at all. ***None of this experience appears to be related to grantmaking.***

We have one person who worked for a decade as a mainstream Chief Development Officer, where she most likely would have focused heavily on raising new funds (as opposed to figuring out how to spend those funds wisely). We have three people who each have a few years of experience working as team leads or front-line staff at mainstream nonprofits – this would teach them about what charities need and about how to fill out grant applications, but not necessarily about how to design or evaluate grant applications. Finally, we have two people who briefly worked in roles at mainstream charities that seem unrelated to grantmaking.

That was all the non-EA grantmaking experience on AI safety grantmaking teams that I could find. To the best of my knowledge, everyone else at Open Philanthropy, Longview Philanthropy, Macroscopic Ventures, the Long-Term Future Fund, and the Survival and Flourishing Fund either has grantmaking experience from solely within the EA bubble, or has no previous grantmaking experience at all, or works on grantmaking for other cause areas besides AI safety.

If I've missed someone with relevant experience, please let me know, and I'll be happy to add them. Otherwise, I am forced to conclude that there is essentially zero mainstream grantmaking experience among AI safety grantmakers.

Ineffective Bootstrapping

Most of the techniques that AI safety grantmakers are applying to evaluate the quality of grant applications are therefore likely to be techniques that they invented themselves or that they learned from each other as they worked their way up the ranks of EA-affiliated philanthropies.

These techniques are very unlikely to be adequate.

In most cause areas, we would be able to "bootstrap" up from a very thin knowledge base by seeing which grantmakers were funding grants that tended to succeed. For example, if you're trying to fund malaria eradication, you can look and see what happened to malaria rates among the populations served by the grants you funded. Over time, you'll build up a reasonably clear picture of which grants were most successful, which will in turn allow you to identify which grantmaking techniques were most successful. Thus, even if your team starts out without knowing much about proper grantmaking procedure, your team could gradually converge on some of the best grantmaking practices through trial and error.

However, in the field of AI governance, it is impossible to do a significant amount of bootstrapping, because we get very limited real-world feedback, and because we have a vague and unclear theory of change.

AI governance provides very limited feedback because much of the "outcome" of AI governance work is long-term, speculative, and binary. To put it bluntly, either we will wind up in a future where there is an AI apocalypse, or we will not – and we will not necessarily observe which future we are in until it is too late to meaningfully change our

grantmaking decisions. Unlike malaria, misaligned superintelligence is not a problem that claims one victim at a time.

As a result, we should not expect to get much objective data about which AI governance grants were most successful. We should especially not expect to get this data with a casual or incidental effort – if we want meaningful data about which governance programs are succeeding, then we'll need to carefully identify reasonable proxies for success and then make a special effort to measure those proxies. As I will argue later in this post, it does not appear that this work has been accomplished. Because we do not know much about which governance projects have been most successful, our grantmaking teams will not have been able to learn much about which grantmaking techniques were most successful.

Suppose you apply a particular flavor of grantmaking to fund my AI governance research paper. Your methods lead you to be very confident that my research paper is worth funding. Later, you want to find out whether my research paper was successful so you can decide whether the methods you used led you to make an accurate prediction about the value of my grant.

What would it even mean to decide that my governance paper was "successful?" It is not clear what (if any) theory of change can be applied to the paper. You can perhaps observe whether I published a governance paper, whether other researchers agreed that the paper seemed insightful, or whether I had many citations, but none of these metrics are very well-correlated with how much (if at all) the paper helped save the world. It's profoundly unclear who is supposed to be reading these academic-style papers, why they will behave differently as a result of reading them, and why those behavioral changes will reduce the risk of an AI catastrophe. Not only are we not measuring any quantitative data about such papers – we don't even have a good explanation for why or how we should expect such papers to be helpful.

As a result, there's simply no viable feedback loop for AI safety grantmakers. Most AI grantmakers are very diligent, thoughtful people, but no matter how diligent you are, you can't improve your skills through deliberate practice unless you're given access to some kind of data about how well you've been succeeding, and in the AI governance field, that data mostly hasn't been available.

Therefore, our movement has a staff of professional grantmakers who were not taught how to manage a grantmaking program by mainstream philanthropies, and who are

unlikely to have learned much about that task based on their years of experience within the EA bubble.

HOW GRANTMAKING IS SUPPOSED TO WORK

The main reason why the AI safety funders' lack of non-EA grantmaking experience matters is that it tends to create excessively informal grant evaluations that are more likely to lead to suboptimal decisions. To explain what I mean by this, I want to first lay out how professional grantmaking is supposed to work and then lay out what my experience has been with grantmaking in the AI safety field.

The Core Process for Traditional Grantmaking

The core process for professional grantmaking consists of the following steps, in approximately this order:

1. A funder settles on a budget and timeline for achieving a particular goal.
2. The funder communicates their goals, budget, and timeline to potential grantees so that they can prepare to apply for funding.
3. The funder decides how grant applications will be evaluated and explains this process to potential grantees so that they know what kinds of grants to apply for.
4. The funder collects and evaluates grant applications based on the process.
5. The funder notifies all applicants about the decisions that were made on their grants and the reasons for those decisions.

This is the process that I used for five years at HomeBase, which refereed grant application processes every year for homelessness services at each of about 30 different counties spread across the US Pacific region. Nobody saw the basic steps of this process as overly fussy or unnecessary – on the contrary, participants used to badger us to release information about each of the five steps as soon as possible. People would drive for an hour to come to our briefings, because they wanted to find out how they could earn a grant and how they could keep it.

As part of this process, we always had a written rubric that we used to assign numerical scores to each grant application. An example of this kind of work can be seen starting on page 22 of this handbook. To the maximum extent possible, we tied scoring factors to measurable quantitative outcomes – for example, a rapid rehousing program could get up to 3 points for enrolling its clients in health insurance, up to 8 points for providing at least as many apartments as they pledged to provide, up to 2 points for helping clients increase their income, and so on. The larger the percentage of clients who get their health insurance, the larger the share of the 3 points for health insurance you could claim. We took the mystery out of funding; for the most part, you could look at your own performance and figure out whether you would be eligible to have your funding renewed.

When it wasn't possible to use quantitative estimates, we described the qualitative outcomes we were looking for in great detail, breaking these outcomes down into multiple subparts and awarding a specific number of points for satisfying each subpart. We then recruited a panel of independent judges to determine which subparts had been satisfied, and we averaged their ratings.

The resulting final scores were used to rank all of the grant applications in order, and then those applications' scores were made public so that everyone could see which applications were funded and why.

This procedure was required by federal law, but the level of transparency and rigor we offered was seen as typical in our industry, including by people coming from state and local government, churches, hospitals, and private foundations.

For example, Chris Kabel, a senior program officer at the Northwest Health Foundation, which manages $68 million in grants, writes that his organization hosts "a grantee forum where we invite anybody who is interested in applying for a particular program to learn about what the program is trying to achieve and what we're hoping to see in competitive proposals. We also answer questions they have that are relevant to their particular programs or initiatives. I'd say almost all of our grantees probably already know how they fit into our program's goals and strategies by the time they get a grant from us."

Here's a similar quote from Anne Warhover, the CEO of the Colorado Health Foundation, which manages $1 billion in assets. "Everybody has to be accountable in this world. We start by thinking about who is our customer, who is our stakeholder, to whom do we owe our accountability, and that is the people of Colorado. This isn't our money. This is their money. And so how can you possibly be accountable without showing results and having

some objectivity to your results? You can't just fling your money in all different directions and hope that some of it sticks. You have got to have strategies that will help you get to those results."

Likewise, Ken Thompson, a program officer at the Bill & Melinda Gates Foundation, writes that "the single most helpful thing you can do to make the reporting process useful to everybody is to be clear up front about what the project intends to accomplish. The other thing that is particularly helpful for grantees is to identify a set of reasonable goals to measure."

For some more examples of how foundations put this kind of grantmaking technique into practice, you can check out the scoring rubrics provided by:

- The Rising Foundation
- Ability Central
- The Tallmadge Foundation
- The John & Deborah Gillis Foundation
- The Gates Foundation – RISE Education Fund

Advantages of Using Formal Rubrics

These rubrics have two important features: (1) they assign a specific number of points to specific criteria, and (2) they explicitly say what is necessary to earn points in each category. Together, these features make it much more likely that grants will be awarded based on the foundation's true values. It is still possible for an evaluator to 'cheat' and assign a project an artificially high score simply because they like it, but doing so will require explicitly lying about whether it meets the written criteria. Most people who are altruistic enough to work at a nonprofit find this kind of explicit lying very distasteful and will usually avoid it.

Obviously, effective altruists can and should object to some of the specific goals being pursued by these grantmakers – grants to rural American schools or to fund career development for people with disabilities are honorable causes, but they are probably not reasonable choices if you are trying to do the most good per dollar. Effective altruists might also object to some of the specific criteria used in these rubrics, or at least to their weighting – for example, if I were writing the rubric for Ability Central, I would not award

five times more points for having "people with disabilities in leadership" than for having a clear "needs assessment" that shows why the grantee's project is necessary.

However, the ***process*** used by these mainstream philanthropies is absolutely vital, because it forces grantors to make decisions about which criteria are most important to them and to communicate those decisions to prospective grantees. I can criticize Ability Central's grantmaking criteria and try to improve them *because* those criteria are written down. If Ability Central had instead used an informal evaluation process that allowed their staff to recommend whichever grants seem best to them, then nobody would be able to tell that, e.g., Ability Central was missing out on opportunities to do more good because they were paying too much attention to the identities of their grantees' leaders.

Thus, even though Open Philanthropy's ***substantive goals*** are better than Ability Central's goals, Ability Central's ***grant evaluation procedures*** are better than Open Philanthropy's procedures.

If you haven't written down the criteria that you're using to evaluate grant applications, then nobody (not even you) can be sure that you're funding the projects that best support your mission. You might instead be funding whatever projects subjectively appeal to your grantmaking staff for reasons that you would not reflectively endorse. As Luke Muehlhauser of Open Philanthropy wrote in 2011, "Cognitive science shows us that humans just are a collection of messy little modules like anger and fear and the modules that produce confirmation bias. We have a few modules for processing logic and probability and rational goal-pursuit, but they are slow and energy-expensive and rarely used...our brains avoid using these expensive modules whenever possible."

If you allow grantmakers to make subjective personal decisions about which grants to fund without a formal process that forces them to connect their decisions to community-approved criteria, then at least some of the time, some of their brains will slip into cheaper processing modes that default to whatever conclusions feel comfortable – that's just how humans work.

Clarity vs. Rigidity

It's not necessary for these grantmaking criteria to be utterly rigid – if there's a good reason to make an exception to your criteria, then you can and should allow a grantmaker to write up a defense of why they are making an exception.

In other words, my proposal is not that donors should mechanically total up a list of points and then feel bound by the outcome. Rather, my proposal is that donors should go through the exercise of specifying concrete goals, assigning weights to those goals, and then totaling up the resulting points in order to (1) get clear in their own minds about what types of achievements are most important, and (2) make sure that they are not accidentally allowing themselves to place too much weight on criteria that are less objectively important.

If, after totaling up the points and taking the time to write up an explanation of why the points are misleading for a particular grant, a donor remains convinced that the points lead to an incorrect recommendation, then the donor should ignore the points and act on what they believe to be the correct recommendation.

However, if you don't even insist that people explain how their decisions relate to their principles, then much of the time, their decisions ***won't*** relate very closely to their principles, and as a result those decisions will predictably do less good.

HOW GRANTMAKING WORKS IN AI SAFETY

This seems to be what is happening in the AI Safety grantmaking space – there is no process that would require grantmakers to connect their decisions to their principles, and as a result, grantmakers are often making deeply suboptimal funding decisions.

Blundering in the Dark

I attended a talk at Less Online this year where Lydia Nottingham attempted to create "GiveWell for AI safety," asking participants "how might we estimate the cost-effectiveness of marginal donations in AI safety?" The session was fruitful in that we considered various alternative metrics and had a thoughtful discussion of their pros and cons. Should we measure the number of citations to a research paper funded by a grant? Which citations matter? Can they be weighted appropriately? How should an academic citation be weighed against a citation in popular media?

In one sense, this was a very interesting discussion. In another sense, it was absolutely outrageous that amateurs could be having such a conversation for an hour on end without

anyone stopping and saying, "Wait a minute; this has all been worked out already by the professionals; here are the metrics they use and here's why they use them!"

Hundreds of millions of dollars have already been allocated to AI safety projects. What was the basis of those allocations? The horrifying answer is that there is no such basis – or, at least, no basis that grantees are allowed to see.

How many pieces of legislation does Open Philanthropy think CAIP should have been able to edit for $2 million? How many citations in earned media do they think we should have had? How many times do they think we should have met with Congressional officials? What kind of language do they think those officials should be using when they talk about CAIP or about AI safety? If these aren't the correct metrics, then what ***are*** the correct metrics?

I did not have this information when I set out to build up CAIP and started making new hires, and I still don't have this information today, and that's ridiculous. If I had known exactly what our donors' expectations were in 2023, there's a good chance that I would have turned down the work and said, "I'm sorry, but I can't achieve those results on the budget you're offering." That would have saved the movement quite a bit of money. Alternatively, perhaps I would have been able to refocus our work so as to actually achieve the goals, thereby generating a lot of value for the movement. Instead, I was left to blunder in the dark.

A Note on Goodhart's Law

One grantmaker I spoke with defended this ambiguity by noting the potential that explicit performance targets could be abused by self-interested grantees. He is concerned that if, e.g., advocacy groups are scored based on the number of Congressional meetings they attend, then groups might, e.g., artificially increase their apparent performance by scheduling hundreds of low-quality meetings with less relevant Congressional staffers.

In my opinion, this conflates the bare possibility that Goodhart's Law might have some effect on AI safety advocates (which is true) with the conclusion that Goodhart's Law applies so strongly as to render it infeasible to have any formal performance criteria at all (which seems false).

To continue the example from above, if you are worried about advocates having too many low-quality meetings, then you can add a scoring factor that evaluates the quality of the

meetings. You can carry on adding as many factors as seems necessary until your scoring factors collectively arrive at a reasonable approximation of the thing you are trying to measure.

Bureaucracies engage in formal, multi-part evaluations all the time, not all of which are completely pointless or corrupted. Nobody suggests that, e.g., food safety inspectors are creating an artificial set of incentives for restaurants to sweep their floors, wash their dishes, and throw out expired food, even if these tasks are on a checklist that is distributed to restaurant owners in advance. The reason why these tasks are on the checklist is that they really do increase food safety, and there is no cheap and easy way for most restaurants to appear to have completed these tasks without *actually* completing these tasks. Similarly, when I was evaluating the performance of homeless shelters, I both checked to see how many people were sleeping in beds, and to make sure that those beds were clean and adequately sized and made of appropriate materials. If a shelter somehow found a way to increase the number of people who were sleeping indoors on appropriate beds, that would genuinely reflect well on the shelter.

For the most part, there's no reasonable way to cheat on a well-designed grant evaluation process. If you find yourself suffering badly from Goodhart's Law, that's usually a sign that you're not investing enough time or enough talent in designing the formal rubric. For instance, if you're handing out a billion-dollar grant, then you shouldn't expect a one-page checklist to adequately constrain the grantees; for that much money, even well-intentioned people will often fall prey to corrupt incentives to find extremely clever ways to optimize for satisfying the formal criteria. However, if you've got a million-dollar grant backed up by a five-page rubric, then it's usually not worth a grantee's time to try to artificially inflate their results. This is especially true when the grantmakers ultimately remain free to use their judgment to overrule the formal results of a rubric.

As a result, I think it is very ill-advised to deliberately obscure the grantmaking process as part of an effort to prevent artificial optimization. Grantmakers need to be very clear in their own minds about what they're trying to fund and why so that they can make optimal decisions, and grantees also need to know what kinds of grants will be funded, so that they can make responsible long-term plans and focus their efforts on activities that the community is willing to support. The problems that are currently being caused by a lack of clarity in these areas are far more severe than any minor difficulties that might arise based on Goodhart's Law.

Open Philanthropy's Evaluation

To illustrate what I mean by a lack of clarity, I am sharing a full copy of the final email that I received from Open Philanthropy regarding their evaluation of CAIP. I have other communications from them, but none of them are substantially more detailed. There is no appendix or similar document that unpacks their reasoning.

The point of this example is solely to illustrate the vagueness of AI safety grantmaking evaluations. Even if you happen to agree with all of Open Philanthropy's conclusions here, you should ***still*** want them to use more rigor than this so that you can be confident that their grantmaking decisions will reliably match up with their stated values.

"Dear Jason,

I hope this email finds you well. I wanted to follow up on our conversations regarding CAIP's funding application to Open Philanthropy.

After careful consideration and thorough evaluation, we have decided not to recommend CAIP for funding.

This was a challenging decision in light of the helpful information you've sent over the last few months. We continue to hold CAIP's thinking on policy development in high regard, and we appreciate CAIP's traction in policy engagement and communications. We concluded, however, that the counterfactual impact does not clear our especially high bar for recommending 501(c)(4) advocacy organizations, for which heightened sensitivities and less favorable tax treatment make donor dollars scarce.

This is a reflection in part of the improved overall state of the AI safety advocacy field, where there are now several organizations doing great work and not enough donors to fund them all. But, in the interest of transparency, it also reflects our continued sense that CAIP's strengths and weaknesses are an imperfect match for its activities. For example, we think CAIP has shown a capacity for developing and refining detailed policy ideas, and to some extent educating policymakers and (in the early stages) building grassroots support – all activities that could be done at a standard nonprofit. But we still have doubts about CAIP's political aptitudes, and the lobbying track record, while better than we initially thought, still does not seem strong enough to create a sufficiently compelling case.

Given our limited funding capacity for advocacy work, we need to direct resources where we believe they'll have the greatest marginal impact. We know that this decision is likely to create challenges for CAIP, and we sympathize. We genuinely wish you and the CAIP team success in your future endeavors, and we appreciate your hard work to improve the outcomes from transformative AI.

Best regards,

[Person N]"

Notice that this email does not include any scores, any benchmarks, or any quantitative information at all. Although the email alludes to an "especially high bar" for recommending 501(c)(4) organizations, it does not provide any information about where that bar sits or how it is calculated. The email mentions some of the criteria that Open Philanthropy considered, such as CAIP's "counterfactual impact" or its "political aptitudes," but does not say how these factors were evaluated or weighted against each other. It is not clear from the email what it would mean to have high political aptitude or how Open Philanthropy determined that CAIP's political aptitude is low or medium.

This is particularly troubling because without this kind of information, it is unclear how Open Philanthropy could compare or evaluate the relative good done by a research organization versus an advocacy organization. There are three times more researchers than advocates in the space, so on the margins, we can say that Open Philanthropy chose to fund some additional researchers instead of funding CAIP. Did those researchers have strong research aptitude? If so, does that mean that those researchers would do more to reduce the risk of an AI catastrophe than CAIP? Why? How do we know, or what is the basis of that opinion? It is not clear that Open Philanthropy is even attempting to answer these questions.

To its credit, Open Philanthropy does have published criteria for its AI governance grants. However, Open Philanthropy warns that such grants are evaluated "holistically," and it is not clear that most of these criteria were consulted when evaluating CAIP's application. According to Open Philanthropy's published criteria, "key considerations ... include:"

1. whether the proposed activities are justified by the applicant's "theory of change,"
2. the applicant's "track record" of success at similar projects,
3. the applicant's "strategic judgment" in making "well-thought-out decisions under uncertainty,"

4. an indication that the applicant is "aware of potential risks and can prioritize how to respond to them,"
5. the "cost-effectiveness" of the proposed budget in light of "the project's goals and planned activities," and
6. whether the grant application has a financial "scale" that is suitable for this particular fund.

Open Philanthropy's "sense that CAIP's strengths and weaknesses are an imperfect match for its activities" could be a reference to criteria #1 and/or #5, although it is difficult to tell, because the email does not explicitly reference any of the criteria. Similarly, Open Philanthropy refers to CAIP's "lobbying track record" as "not...strong enough to create a sufficiently compelling case," which could be a reference to criterion #2.

None of these three criteria are discussed in a clear enough way to allow Open Philanthropy to easily compare CAIP's performance to the performance of other organizations. For example, what precisely does it mean for CAIP's strengths and weaknesses to be an "imperfect match for its activities?" Is this like getting 2 out of 5 points? 3 out of 5 points? 4 out of 5 points? How well do other grantees' strengths and weaknesses match their activities, and how is this evaluated? Open Philanthropy doesn't say. Even if Open Philanthropy's grantmakers have an intuitive idea in their minds about what these phrases mean, such intuitions are highly vulnerable to subjective bias and to being misremembered or subconsciously edited from one grant to the next. The point of writing down a numerical score for each grant on each criterion is to mitigate such biases.

Open Philanthropy's other three criteria are not mentioned at all in their evaluation of CAIP. I cannot find any discussion of CAIP's strategic judgment, risk management, or scale.

Even if, for some reason, only half of Open Philanthropy's criteria were relevant to its evaluation of CAIP, it would still be helpful for all parties if they explicitly noted this. As I argue above in the section on "Advantages of Using Formal Rubrics," the process of thinking through how each criterion applies to each grantee and writing down your thoughts helps make sure that you honor your core values and that you are not accidentally substituting other considerations.

I do not believe the informality of Open Philanthropy's evaluation techniques is limited to CAIP. For example, Apart Research recently wrote° that their need to engage in last-minute fundraising was partly caused by "somewhat limited engagement from OpenPhil's GCR

team on our grants throughout our lifetime," and that "with OpenPhil, I think we've been somewhat unlucky with the depth of grant reviews and feedback from their side and missing the opportunity to respond to their uncertainties." This is a very polite way of saying that they don't know why Open Philanthropy isn't giving them more funding and that they haven't been able to find out.

I also spoke with Soren Messner-Zidell, a senior director at the Brookings Institution who was hoping to launch an AI safety communications project. He has over 15 years of experience as a public advocate. After he negotiated with Open Philanthropy for several weeks, his application was also denied without any type of detailed explanation. Like me, he does not know what criteria they are using to evaluate AI safety grants.

Longview's Evaluation

CAIP received a similarly unstructured final report from Longview Philanthropy. Again, there was no appendix or more detailed letter that provided significantly more information about Longview's process. If they conducted any more formal evaluation, they did not share it with us. Below is the full text of the most relevant email:

Hi Marta and Jason,

Thank you so much for engaging with us throughout this investigation. Sadly, Longview will be unable to find donors for CAIP's work in the near future.

This conclusion is based on our completion of several other grant investigations and discussions with our most reliable donors. The unfortunate reality as we see it is that there are more promising opportunities in AI policy than our donors are able to fund, and despite CAIP's strengths, we will not be able to provide funding for it.

We think CAIP is doing good work. We hope you find the support needed to continue as an organization, and if you do, we'd be happy to consider recommending CAIP to our donors in the future. But at least for the next few months, and likely until the end of the year, we will not be able to find additional funding for CAIP.

Thank you for all the effort you've put into this important cause. If there's anything beyond funding that we can do to help going forwards, please let us know.

Sincerely,

[Person V]

This email is even vaguer than the email from Open Philanthropy because it does not even gesture at the factors that might have been considered – this email provides only the raw decision that CAIP is less valuable than other projects.

Longview's website provides a list of some of the questions that they consider when evaluating grants, but it is not clear how the answers to these questions are measured or which questions are considered most important. Longview does not say anything about which of CAIP's answers (if any) were deemed unsatisfactory.

Longview's website also says that when they recommend a grant for funding, they "include quantitative predictions, so that philanthropists can clearly see what we expect in terms of key outcomes, and how likely we think they are."

This is praiseworthy, but it is unclear whether Longview follows similar procedures for rejected grants, or if they only do this for grants that they have already decided to recommend. If the former, it is unclear why they are unwilling to share even an outline of their predictions with grant applicants. If the latter, it is unclear how they know which grants to reject – if, as Longview writes, CAIP "is doing good work," then how was Longview able to decide not to recommend CAIP even before Longview made any quantitative predictions? If the rejection isn't based on a quantitative prediction, then what ***is*** it based on?

I find the lack of detail disturbing, and I think you should too. If there is some particular scoring factor (e.g. one donor's private preferences) that must be kept secret in order to maximize total funding, then that particular scoring factor could be censored. However, to share no details at all about the grant evaluation process with potential grantees sets them up for failure and wastes everyone's time. If we don't know how we're being measured or what standard we're being held to, then we can't make realistic plans to achieve those goals.

LTFF's Evaluation

The Long-Term Future Fund rejected a CAIP grant application in March 2024, and sent only a standard form rejection. They wrote, "We have reviewed your application carefully and regret to inform you that we have decided not to fund your project... Please note that we are unable to provide further feedback at this time due to the high volume of

applications we are currently receiving. We know that this is something many applicants in your position want and I hope that when we have more capacity we will be able to give more feedback."

A year later, in March 2025, I wrote to [Person AA], the manager of the Long-Term Future Fund, to see if they had developed additional capacity; I mentioned that I was travelling to California and would be happy to meet with him in-person to learn more about which projects they were interested in funding. My email wound up in his spam folder by mistake. We noticed this problem while LTFF was reviewing an advance copy of this post, and Person AA mentioned that Person K and Person L both have more DC experience than he does, implying that they would be better positioned to discuss how CAIP's grant application was evaluated.

I repeatedly attempted to get feedback from Person K about why he felt that CAIP was not a funding priority, both while he was working at CAIP and after he left in February 2024. However, he did not have any specific criticisms or criteria to share other than his general feeling that political advocacy as a whole was unpromising and that CAIP did not seem to be generating the kind of results he was looking for. He repeatedly offered to write up his objections in more detail, but he never did so.

I also reached out to Person L, and he wrote, "I'm not sure that I'd be able to provide a ton of useful info. I'm not actually personally that aware of what CAIP has been up to, and if anyone were to ask me what my thoughts are on CAIP, I'm pretty sure their main takeaway from my response would be that I don't have enough context to have a strong view either way."

I was able to chat in April 2025 with Person MM, another LTFF fund manager. Person MM did share some of his concerns about funding projects like CAIP. His main objection seemed to be that – so far as he could tell – CAIP was publicly conflating the issue of "AI being unreliable, or being used in critical systems" with "AI becoming agentic and deceptive and eradicating all of humanity." He expressed his worry that CAIP had sold out or would sell out by trading away the clarity and honesty of its message in exchange for political influence. I promised that CAIP was firmly opposed to any such trade, and I cited three of our recent public articles in which CAIP openly and clearly states its intense concern about existential risk. However, Person MM did not directly respond to this information.

I defend CAIP's choice to include at least some mention of less-catastrophic risks in our conversations with Congressional staffers in the first post° in this sequence. Even if LTFF does not find this defense convincing, I would still hope that they make their funding decisions based on some type of formal criteria, rather than informally rejecting applications from groups like CAIP based on a general opinion that most groups who work in Washington, DC will inevitably lose their focus on existential risk. For example, LTFF might have a rubric that they use to evaluate the quality of an advocacy group's public-facing communications, including both the level of focus on x-risk and presumably several other factors, and then assign a certain number of points to the grantee based on how well the grantee's actual communications actually satisfy each of those factors.

When I asked Person MM about whether LTFF used these kinds of formal written criteria for evaluating grants, he said, "I don't think there exists a single functional fund in the world with 'formalized grantmaking criteria,' so of course not." Person MM pointed out that his thoughts on grantmaking principles are spread across many comments on LessWrong and the EA Forum, and that these would be "kind of a pain to stitch together." Person MM added that LTFF intentionally "has people with different principles and perspectives" on grantmaking, so LTFF does not have any "universal grantmaking criteria."

As I showed in the previous section, many funds do in fact use formal written criteria, for the very good reasons that (1) this helps them ensure that the grants they fund are actually optimizing for their core values, and (2) this helps prospective grantees submit grants to organizations that are more likely to fund them, thereby making better use of everyone's time. If Person MM himself would find it unreasonably inconvenient to stitch together his own comments on grantmaking principles, then surely it would also be unreasonably inconvenient for prospective grantees to do so, especially since they do not know which LTFF reviewer will be assigned to their application – in order to intelligently plan a grant that is likely to receive LTFF funding, they would have to hunt through old comment sections to find comments from several different possible reviewers and then decide which comments (if any) are most relevant.

Similarly, if a grantmaker's principles are spread out across many decentralized comments, then it is harder for that grantmaker to use those principles as a tool to focus their own thinking – instead of reliably following their most important principles, they may sometimes be distracted by secondary considerations. Most grantmakers should be using written grant evaluation criteria for the same reason that most people should be keeping a

written calendar – it's important to get it right, and writing things down in a central place helps you remember and act on your priorities.

To their credit, LTFF staff have put together a sample of different portfolios of hypothetical grants° that they would fund depending on the current size of the LTFF budget, and these portfolios could help guide prospective applicants. Most of the twenty grants in the portfolios relate to academic-style research in some way, either by directly funding such research, or by funding people who want to continue their academic education. A few grants discuss AI governance research, but none seem to be about direct political advocacy. The closest approach is a hypothetical grant to hire a communications firm to teach best practices to professionals in longtermist organizations, which presumably might then lead to some of those professionals engaging in advocacy or improving the quality of their advocacy. The range of grants that LTFF is even willing to consider thus seems to be heavily biased toward research and away from advocacy.

As part of their explanation of how they created these portfolios, LTFF notes that "At the LTFF, we assign each grant application to a Principal Investigator (PI) who assesses its potential benefits, drawbacks, and financial cost. The PI scores the application from -5 to +5. Subsequently, other fund managers may also score it. The grant gets approved if its average score surpasses the funding threshold, which historically varied from 2.0 to 2.5, but is currently at 2.9."

While this quantitative scoring is at least a step in the right direction toward rigor, it does not seem to include any explicit scoring factors or rubrics. Any given investigator is essentially still making a subjective judgment; it's just that they're then reducing that subjective judgment to a number. There is no way for grantees to predict what that number might be or to design grants that are more likely to earn a higher number. To the extent that LTFF is suffering from a systematic bias toward research and away from advocacy, nothing about the process of writing down a single number for each grant would tend to correct that bias.

Other Evaluations

I had similar trouble getting any kind of concrete explanation of the funding process from other AI safety funders. SFF provides anonymous feedback from its recommenders, and allows for one round of follow-up questioning between the grantees and its

recommenders, but the recommenders did not appear to have any type of rigorous process for scoring their applications.

For example, one SFF recommender wrote "I'm not a DC insider, I was mainly deferring to what DC people told me." Another SFF recommender wrote, "The main reason I was unable to fund you was applications by more excellent US policy organizations than we had room to fund." Some of the recommenders mentioned that they thought other organizations had done more to promote "counterfactual policy changes," but did not say which changes they had in mind or why the changes at other organizations were more valuable. None of the recommenders cited any numbers in support of their funding recommendations.

Macroscopic Ventures had even less information. An email we received from Macroscopic's President, Ruairi Donnely, said only that "we discussed this with some of our grantmaking advisors and unfortunately we have decided not to make a grant. Sorry for not having better news. I hope you manage to secure funding and wish you all the best with your work."

WHY ISN'T THERE MORE RIGOR IN AI SAFETY GRANTMAKING?

A lack of detail in post-evaluation debriefings is excusable if there are clear, publicly available criteria for evaluating grants. Most grantmakers don't have time to explain their conclusions to all of the hundreds of people who apply for their funding, especially since many of those applications will be low-quality or otherwise not a plausible fit for the grantmaker's particular fund. If AI safety grantmakers had published quantitative criteria for scoring grants and then shared the resulting scores with applicants, I would not be expecting any further information.

What I find outrageous is that AI grantmakers are ***neither*** publishing formal criteria before applications are evaluated ***nor*** writing detailed feedback after grants are awarded. This means that nobody has any way of telling why some grants are funded and others are denied.

Reasonable Overhead

This lack of rigor in AI safety grantmaking would make more sense if AI safety funders were only distributing $1 million per year. If grantmaking budgets are small enough, then the extra time it would take you to develop a formal process and explain it to grantees is not cost-effective, because too much of your budget would be going to overhead.

However, this is not a reasonable explanation in the context of AI safety grantmaking. Open Philanthropy alone gives out more than $100 million per year just on AI governance. An average level of administrative overhead for a foundation is 11% to 15%. If you spend 5% of the overhead budget (so, about half a percent of the total grant volume) on designing and communicating formal rubrics, you'd have a budget of at least $600,000 per year, i.e., enough to pay a couple of full-time employees just to work on AI governance grantmaking rubrics. Obviously, this has not been happening.

It ***should*** be happening, because it's very easy to reap large savings or other large benefits from clearly explaining your grant procedures. If CAIP was an unnecessary expense, then the movement could have saved $2 million by clearly communicating that in advance – more than enough to cover whatever staff time was required to develop and communicate the relevant criteria.

Linchuan Zhang of LTFF once wrote° that most LTFF grantmakers are "*very* part-time," because they work long hours for their day jobs and only attend to LTFF grant applications with the limited time that they have remaining. He speculates that "full-time grantmakers at places like Open Phil and Future Fund" might "have similarly packed schedules as well, due to the sheer volume of grants."

If there is a shortage of staff time, then AI safety funders need to hire more staff. If they don't have time to hire more staff, then they need to hire headhunters to do so for them. If a grantee is running up against a budget crisis before the new grantmaking staff can be on-boarded, then funders can maintain the grantee's program at present funding levels while they wait for their new staff to become available.

Instead, AI safety funders appear to be just trusting their instincts and giving out money based on their intuitive preferences – even though there's every reason to believe that this results in deeply suboptimal awards.

Micro-Dooms Per Dollar

One of the grant managers at LTFF, [Person K], once told me that he thought it might be helpful for AI safety funders to evaluate projects in terms of a standard number of "micro-dooms" per dollar that each project is expected to avert. For example, if grantors believe that a project will reduce the risk of catastrophe by 0.01%, then they would give it credit for preventing 100 micro-dooms. This would provide a fair and reasonable way of comparing the value of AI safety projects that operate in relatively different specialties, such as advocacy targeted at politicians, abstract governance research, and communications aimed at the general public.

However, none of the funders that CAIP spoke with said anything to us about how many micro-dooms they expected CAIP to avert, or about their targets for how many micro-dooms per dollar they expected to avoid when funding a successful project, or how they would calculate the number of micro-dooms averted by a project.

I followed up with Person K while writing this post, and he wrote that, "In the past year or so, I've updated towards the position that it's very difficult to do correct cost effectiveness evaluations, and so we mostly just have to make best guess qualitative judgements. I believe this is basically the standard grant maker take."

I agree that it is difficult to make accurate cost-effectiveness evaluations, but I would be very surprised if the effort of trying were not worthwhile. When you attempt to quantify the value of an outcome, even if you make an error, your attempt will probably be more rigorous and more accurate than an attempt made without using any numbers at all. As long as you're honest about your error bars, you should be able to at least compare the benefits offered by two different grants in terms of their likely orders of magnitude.

If AI safety grant evaluations are so challenging that grantmakers can't even arrive at an approximate order of magnitude for the micro-dooms that each grant is likely to avert, then that strongly suggests that grantmakers don't know enough about specific AI safety projects to be usefully picking and choosing individual grants. In that case, grantmakers may as well just assign a high funding priority to every plausible grant in the most effective ***category*** – and the most effective category, as I have been arguing for the last 40,000 words, is advocacy, not research, because only advocacy has the power to change the incentives of the private AI developers who are likely to carry us with them into ruin.

CONCLUSION

Each year, most of the hundreds of millions of dollars that are donated to AI governance are spent on academic-style research. There does not appear to be any clear idea of what this research is supposed to accomplish, how those accomplishments will improve the world, or how much value this research will generate per dollar. The process that results in awarding grants to researchers appears to be based primarily on intuition and subjective judgments, rather than on any formal criteria.

If our AI safety funders were staffed by grantmaking experts who had worked for prestigious mainstream philanthropies, then I would be more inclined to give them the benefit of the doubt and assume that they ***do*** have rigorous rubrics hidden away somewhere, and that they've just chosen not to share them for some unspecified reason. However, instead, our AI safety donors are staffed by people who have very little grantmaking experience outside the EA bubble. ***It seems entirely plausible that most of them are simply unfamiliar with proper grantmaking procedures, or unfamiliar with the reasons why such procedures are useful.***

Similarly, if our AI safety funders were staffed by experts in political advocacy who had worked for many years in DC, then I would be more inclined to give them the benefit of the doubt and assume that they do have some good reason for directing most of their funding toward academic-style research, even though – as I have argued throughout this sequence – further academic research doesn't seem at all likely to prevent an AI disaster.

Unfortunately, we do not have grantmakers who know from deep personal experience what things are really like in Capitol Hill or Sacramento or the Supreme Court or the White House. Instead, we mostly have grantmakers who have perhaps completed a political internship at some point, and who are making decisions primarily based on their experience as academic-style researchers who have published papers and talked to other researchers. As such, ***it seems entirely plausible that most of them are simply biased toward funding other researchers because they find research to be a familiar and comfortable topic.***

We are therefore getting suboptimal results: instead of funding the best available projects that are most likely to positively impact the future, AI governance donors are primarily funding the projects that subjectively appeal to their staff.

My point here is not to criticize AI governance grantmaking staff – I believe they are all doing the best they can with the tools they have available. My point is that our grantmaking teams have not been given the right tools: we need to hire additional

grantmakers who have the political and mainstream philanthropic experience we need to invest wisely in the most effective AI governance grants.

Similarly, I do not want to offer any harsh criticism to AI governance donors simply because they have not hired the optimal teams for their grantmaking staff: at least the AI governance donors are donating to a good cause. Thousands of donors pick far less important causes to support, and hundreds of millions of well-off people in developed countries do not bother to make any significant donations at all. By comparison with their peers, people like Cari Tuna and Jaan Tallinn are doing an enormous amount of good. In general, I appreciate their efforts, and I truly am grateful for their support, even though I have spent most of this sequence complaining. If I met the donors, I would offer to shake their hands and buy them a drink.

Nevertheless, I cannot hide my sadness that so much of the money that has been donated to the urgent and important cause of AI safety is on track to be spent in ways that will have only a tiny marginal impact on the future. It is a tragic, avoidable, and enormous waste. If there is any chance that I can ameliorate some of this waste by speaking out about it, then I feel a duty to do so.

RECOMMENDATIONS

1. Hire staff to fill the gaps in political and philanthropic expertise

If you work at a large AI safety donor such as Open Philanthropy or Longview, I urge you to aggressively hire (1) experienced political advocates, and (2) experienced grantmakers from mainstream philanthropies. Jobs for these positions should be constantly open and widely advertised until at least ten new people with the relevant expertise have been hired in each of these two categories.

In addition, grantmakers should use the handful of experts they already have to help find more people with the experience they need. Open Philanthropy did very well to hire Melanie Harris, who is a genuine political expert; they should be asking her for referrals both for Open Philanthropy and for other grantmakers. Similarly, Carl Robichaud at Longview Philanthropy spent a decade running grantmaking in nuclear security at the Carnegie Corporation; if they are not already doing so, then Longview's AI team should be

trying to tap his expertise to improve their AI safety grantmaking processes, and they should be trying to tap his network to hire additional staff with mainstream philanthropic expertise.

If you don't have time to manage that hiring process, then I urge you to hire a headhunter to do so for you – the McCormick Group and NRG Consulting both have relevant networks in DC. If you don't have the funding to hire the additional grantmakers, then I urge you to try to redirect some of your organization's external grant funding to free up the resources to hire more internal staff. The lack of expertise in these categories is wrecking the average effectiveness of AI safety grants, so a small cut in the total size of those grants could easily be 'paid for' by investing in a better process that would make the remaining grants more efficient.

As I write this, in June 2025, it is remarkably easy to hire high-quality political and philanthropic staff. Washington, DC is full of Democrats who are looking for work because their party is out of power, full of moderate Republicans who don't fit in well with the MAGA agenda, and full of civil servants who have been fired or otherwise pressured to leave their jobs. Similarly, the widespread budget cuts and budget uncertainties around, e.g., USAID, Harvard research programs, and government contracting mean that traditional philanthropies are less able to attract and retain their top talent. We should take advantage of the ready supply of talent by hiring new kinds of experts.

To do so, we should be editing our job descriptions to stress our appreciation for political and philanthropic talent, and we should be posting those job descriptions on several mainstream job boards.

For example, one of Longview's current openings for a new grantmaker specifies that the grantmaker should be a "US AI policy specialist," but that opening allows for the possibility that the position could be filled by someone whose primary policy experience is as an academic-style policy researcher. Editing these types of job descriptions to more firmly emphasize the organization's interest in and appreciation for political skills and experience would go a long way toward making political advocates feel that they will be welcome inside the EA movement.

Similarly, if it has not already been added there, then both that job description and other new positions should be posted to a wide variety of job boards that specialize in political or policy work, such as Tom Manatos Jobs, the Internet Law & Policy Foundry, the Public Interest Tech Job Board, Daybook, Roll Call Jobs, and the Public Affairs Council.

Aggressively advertising a few explicitly-political positions to job seekers with relevant political experience would help create a critical mass of political advocates within the movement, which could then make it much easier to attract and retain other political advocates as needed.

2. Write and publish formal grantmaking criteria

If you work at a large AI safety donor, then I urge you to agree upon, write up, and publish formal grantmaking criteria. (If it's important to your organization to allow for a diversity of grantmaking priorities, then have each individual grantmaker write up and publish their own priorities and explain how they will evaluate which grants meet those priorities.)

What this means is not just listing some of the things that you believe are good to have in grants, but also ***developing an organized process that will prompt you to apply your criteria*** to each grant and decide for yourself how well each grant application meets each criterion, ideally by assigning a numerical score to each factor and weighting the results.

At least some of these scores should be based on objective or quantitative factors where different people will readily agree on what the approximate value of the score should be. In other words, don't just check whether a grant is "cost-effective" – check whether it delivers a certain number of meetings or papers or videos or laws per dollar. If you're measuring "expertise," don't just ask yourself whether you think an applicant's staff has good experience – write up a description of what kinds of experiences you would like their staff to have, and then award points based on how many of those experiences are present on the team.

Even if you ultimately decide not to follow the 'advice' of a particular set of scores, it's still an incredibly useful constraint to be forced to check to see what that advice is and to explain why you will or will not follow it in any particular case.

People are drawn to the field of AI safety based on their sincere concern for the public good, but even high-integrity people suffer from cognitive biases and can benefit from using tools (like written rubrics) to help manage those biases. Even if your grantmaking team is so sharp that it can perfectly uphold its values without referring to any written materials, your prospective grantees need formal grantmaking criteria so that they can see what they're supposed to do and intelligently direct their efforts. Not having formal

published criteria sets your grantees up to fail and forces you to read a glut of poorly tailored grant applications that are not a good fit for your organization's funding priorities. It's helpful to regularly write about what kinds of grants you want to fund, but it's even more helpful to publish quantitative criteria; the more information grantees have about your values, the more improvements you'll see in the average quality of the grant applications you receive.

3. Put social pressure on AI grantmakers

If your friends or colleagues or co-authors work for AI governance donors, then I urge you to pressure them to repair the gaps in their staffing and to publish more of their grantmaking criteria. Ask them what they're doing to respond to the concerns raised in this sequence. Follow up at parties and at co-working sessions. Making these changes will take several weeks, but there's no good reason for them to take more than a year – ask your friends what they plan to do, and when they plan to do it. The resulting conversations could be awkward, but failure to make these changes will increase the odds of an existential catastrophe.

4. Consider conducting your own grant investigations

If you are a medium-size AI safety donor, and you see that the larger donors are not responding to this advice, then I urge you to consider making some of your own grant evaluations. Evaluating the efficacy of political advocacy is challenging, and it is probably not something that you can do well in an afternoon, but depending on the size of your contributions, it might be worth your while to hire an experienced analyst on a short-term contract. A professional analyst might be able to write a set of criteria and manage one grant round, in, e.g., three months, suggesting a ballpark cost of $50,000. You might be able to share this cost with one or two friends, or you might be able to defray that cost by investing a significant portion of your own time in making your own independent evaluations.

Please consider the odds that you would actually agree with the judgments being made by institutional donors if you conducted your own investigation, and please consider how much more good you might be able to do by funding the organization that you judge to have made the best-available application in the field, rather than funding more academic research that lacks a clear theory of change. If you can do twice as much good with a well-

chosen advocacy grant compared to a well-chosen research grant, then hiring your own analyst could start to make sense on a total grantmaking budget that's as small as $100,000.

GOODBYE FOR NOW

Thank you for reading this extremely long sequence and for the financial and moral support that many of you have offered along the way. Despite your support, I cannot afford to pay any of the staff at the Center for AI Policy, and so we have ceased operations for now. The non-profit corporation and the website will remain active for the foreseeable future, and I myself remain willing and able to rejoin the movement on a full-time basis if and when sufficient funding becomes available. In the meantime, if any of you have questions about how to make your AI governance work more relevant, I am always happy to help at jason@aipolicy.us.

Previous:

Political Funding Expertise (Post 6 of 7 on AI Governance)

1 comments 37 karma

Comments 2

Sorted by New & upvoted

Kestrel 10mo 15 5 0

Thanks for this series. It helps illuminate a murky area.

Seems to me that it supports the same basic fact I've seen for a while now: the highest-impact stuff is increasingly in the political sphere, a community mismatch where engaged EAs and EA funders are drawn more towards soft political neutrality.

I have no good answers to this problem.

SummaryBot 10mo 1 0 0

Executive summary: The post argues that AI safety grantmaking is dominated by research-focused insiders who lack mainstream political and philanthropic experience, resulting in informal and often opaque funding decisions that neglect high-impact advocacy work; it calls for more rigorous, transparent procedures and greater investment in politically experienced staff to ensure funding decisions align with effective altruist goals.

Key points:

1. Mismatch Between Research and Advocacy Needs: The field disproportionately funds academic-style research over political advocacy, even though the latter is more likely to achieve real-world impact in reducing AI risk.
2. Insularity and Lack of External Expertise: Most AI safety grantmakers have no experience in mainstream philanthropy or professional advocacy, leading to poor adoption of established best practices.
3. Absence of Formal Evaluation Criteria: Funders often make grant decisions based on subjective judgments without using transparent rubrics, benchmarks, or quantifiable goals, leaving grantees without guidance.
4. Challenges of Feedback and Accountability: The lack of measurable feedback loops in AI governance is compounded by an unwillingness to establish even proxy metrics or structured evaluations.
5. Recommendations for Reform: The post urges funders to hire experts in politics and philanthropy, adopt formal grantmaking criteria, and support medium donors in conducting independent grant evaluations.
6. Call to Action and Closure: With CAIP suspended due to lack of funding, the author remains available for consultation and calls for systemic reform to prevent waste and better serve the mission of AI safety.

This comment was auto-generated by the EA Forum Team. Feel free to point out issues with this summary by replying to the comment, and contact us if you have feedback.

More from Jason Green-Lowe

141 **Please Donate to CAIP (Post 1 of 7 on AI Governance)**
Jason Green-Lowe · 1y ago · 40m read
25

63 **Shift Resources to Advocacy Now (Post 4 of 7 on AI Governance)**
Jason Green-Lowe · 11mo ago · 38m read
5

61 **The Need for Political Advertising (Post 2 of 7 on AI Governance)**
Jason Green-Lowe · 1y ago · 16m read
0

View more

Curated and popular this week

Impact Requires Implementation: What Dispensers for Safe Water Scale-Back Can Teach Us About Monitoring

Tony Senanayake · 13d ago · 17m read · 3

TL;DR Between 2022 and 2023, GiveWell and Open Philanthropy committed over $70 million to Evidence Action's Dispensers for Safe Water programme. An independent survey in 2025 found that in Uganda and Malawi the programme was reaching approximately 70% fewer people with detectable...

The AI people have been right a lot

Dylan Matthews · 4d ago · 7m read · 11

This post was crossposted from Dylan Matthew's blog by the EA Forum team. The author may not see or reply to comments. ---------- Subtitle: Try to keep an open mind as the world gets...

In Development Magazine Launch

Lauren Gilbert · 4d ago · 1m read · 4

We're announcing a new development magazine: In Development. Our goal is to promote discourse on development without dogma. In other words, we public stories doing deep dives into pressing development issues, backed up by evidence, argument, and experience. The first piece is by Paul...

Recent opportunities in AI safety

148 **The Anthropic IPO Is Coming. We Aren't Ready for It.**
Sophie Kim · 8d ago · 9m read
16

11 **Pivotal Research Fellowship applications are open (deadline May 3)**
Tobias Häberli, Tilman, James Nicholas Bryant, Ariel · 13h ago · 3m read
0

2 **Applications open for the Online wing of the AFFINE Superintelligence A...**
Mateusz Bagiński, Grace Roberts · 5d ago · 2m read
0

View more