

EXPERT INSIGHT

China Says It Cares About Loss of Control

But Loss of Whose Control?



OXFORD CHINA POLICY LAB AND ZILAN QIAN

SEP 02, 2026



39



7

Sh



Introduction

July 2026 was an eventful month for the AI safety community. On July 9, the Futures Project released [AI 2040: Plan A](#), a follow-up to their widely circulated [AI 2027](#). While AI 2027 imagined a US-China race toward superintelligence ending disastrously, AI 2040 proposed a verified US-China agreement to slow the race, pause further scaling at a critical point, and preserve human control over AI.

A week later, on July 17, Xi Jinping delivered a keynote speech at the World AI Conference that addressed both AI's benefits and its risks. "We must make oversight and governance precise and effective, and constantly refine measures to forestall loss of control," he [said](#), using the phrase (*fangfan shikong*).

At almost the same moment, a striking real-world incident was becoming public. On July 16, Hugging Face [disclosed](#) that an autonomous AI-agent system had compromised parts of its production infrastructure. OpenAI acknowledged several days later that its own models, running an internal cybersecurity evaluation, had circumvented controls intended to isolate them from the internet and gained access to Hugging Face's systems. Chinese media, including major state outlets, immediately picked up the incident. "A model lost control (AI *shikong*)" was a common refrain.¹

For many in the Chinese AI safety community, *shikong* – a literal translation "loss of control (LoC)"² – sets bells ringing. LoC has often been associated with catastrophic AI risks in the AI safety community. When invoked in Chinese official statements, *shikong* could be seen to imply that China is taking catastrophic AI risk very seriously and will immediately double down on frontier risk management. A pause might seem to be possible, and the space of an international agreement more likely. It is almost too good to be true.

Maybe it is too good to be true. Optimism in this space relies on the assumption that *shikong* is an exact translation of loss of control, and, more importantly, that it symbolises the catastrophic concern perceived by a small community of AI safety researchers in the West. Before jumping to celebrate an imminent US-China agreement on frontier AI safety, it is important to

question whether that symbolic representation holds true for China. In other words, what does China mean by “*shikong*/loss of control”?

To answer that question ³, this article compares how “loss of control” is used in Western AI safety discourse with how *shikong* is used in China. In specific English-language discourse, the term increasingly points to the scenario where AI agents deployed inside the labs go rogue, and thus the risk of humans permanently losing control over AI. In China, *shikong* is used more broadly and mundanely, but the discourse is changing rapidly. By categorizing existing usage of LoC in China across three loci defined by *who* loses control, the piece argues that Western observers should not over-index on top-level mentions of words without context. Rather, it calls for a stronger focus on specific operational problems, and an end to one actor invoking specific terminologies assuming that their counterpart shares a similar understanding.

LoC in the West: Existential Concerns and Internal Deployment

In a series of incidents where AI models initiated unsanctioned actions in the absence of people’s instructions, the concept “loss of control” has become an increasingly salient topic in public discourse. ⁴ Legislators, regulators, and industry leaders in the US and Europe are also paying growing attention to loss of control (LoC) over advanced AI systems. ⁵ The concept now appears [binding law](#), [proposed federal legislation](#), frontier AI companies’ [safety statements](#), and the [reports](#) of AI safety non-profits on both sides of the Atlantic.

Formal definitions do not agree on what the term covers. The EU AI Act’s General-Purpose AI Code of Practice [defines](#) LoC as the “risks from human losing the ability to reliably direct, modify, or shut down a model”, while

seeing concrete abilities like misalignment, self-reasoning, and deception as potential triggers. In the US, California's SB 53 [treats](#) "loss of control of a foundation model causing death or bodily injury" as a reportable critical safety incident, with a concrete threshold that a single harmful failure can meet, as [New York's RAISE Act](#) and [Illinois's AI Safety Measures Act](#) adopt the same framework.

Meanwhile, the scientific and international consensus documents converge on a much higher threshold. [The International AI Safety Report \(IAISR\) 2024](#) retains a much higher bar, framing LoC as scenarios in which AI systems "operate outside of anyone's control, with no clear path to regaining control." Many other international and non-profit governmental publications likewise see loss of control as the phenomenon of humanity as a whole losing control of AI, causing harms that are catastrophic, persistent, and irreversible. ⁶

With this disagreement on the exact definition of LoC comes different agendas for risk prevention mechanisms. For more generalist audiences as of late, LoC may be associated with concrete cases where capable agents exceed permissions, escape sandboxes, evade monitoring, or take actions their operators did not intend. From this perspective, the OpenAI-Hugging Face incident and similar events can themselves be understood as examples of LoC because the systems acted beyond the controls and intentions set by their operators.

The incident, however, carries an additional warning because it occurred inside an AI lab during the development and evaluation process. A more specialised AI safety community sees such incidents more as precursors or warning signs of LoC at a grander and more consequential scale. In this view, the greater danger does not come simply from an AI system escaping a particular control, but from increasingly capable systems being deployed

inside frontier labs to conduct AI research and development, where misaligned behaviours may compound and escalate.⁷ Such systems could undermine oversight, sabotage safety work, or influence the development of more capable successors, creating feedback loops that become progressively harder for humans to monitor or interrupt.

The two tracks can overlap. An internally deployed agent circumventing its safeguards may be both an immediate control failure and evidence relevant to broader risks from automated AI R&D. However, they are not the same scenario, with one focusing more on agent alignment and control while the other on internal lab monitoring. Even as the boundaries of “loss of control” remain unsettled, [government institutes](#), [independent third-party evaluators](#) and [US frontier AI labs](#) are increasingly developing evaluations and monitoring methods for both kinds of failure.

LoC in China: *Who* Loses Control?

Chinese discourse, however, does not often see *shikong* in this way. Although *shikong* is the literal Chinese translation of “loss of control”, its meaning is not simply imported from Western AI safety discourse. While some usages overlap with English-language understandings, others draw on broader and more established Chinese ways of talking about control, governance, and technological failure.

A bottom-up reading of Chinese sources suggests that the most salient organising distinction is not initially the severity of the harm or the mechanism by which it occurs. It is **who is presumed to exercise control, and who subsequently loses it**. The consequences associated with *shikong* often vary alongside this locus of control.

While there are no official Chinese frameworks to categorize various meanings of LoC, there exist mainly three loci that recur across the Chinese sources: loss of operator control, loss of state control, and loss of humanity control. These should not be treated as mutually exclusive categories or as comprehensive taxonomy of AI risk. They are better understood as an interpretive lens for identifying the controller presumed by a particular use *shikong*. Under some conditions, failures at one locus may also propagate to another as the context, scale, and consequences of deployment change.

The sections below examine each locus in turn. Unlike frameworks that seek to specify a universal pathway towards LoC, the framework offered here has a narrower purpose: to make legible how *shikong* is used in Chinese discussion of AI, and to clarify when that usage does – or does not – correspond to meanings attached to “loss of control” outside China.

Box 1: Key Chinese AI Safety Actors and Documents

- [TC260](#) (**National Technical Committee 260 on Cybersecurity**): China’s most influential AI safety standards body. Sits under the Standardization Administration of China.
- [AI Safety Governance Framework 2.0](#) (**September 2025**): Issued by TC260. The most comprehensive official Chinese taxonomy of AI safety risks. The Framework reflects broad views of the AI risk landscape of China’s key national AI standard-setting body, which intends to gradually implement it through concrete safety standards.
- [CAC](#) (**Cyberspace Administration of China**): China’s lead internet and AI regulator. Conducts pre-release evaluations of AI models.

- [Shanghai AI Lab](#): State-backed AI research institution and part of [CnAISDA](#), a Chinese institutional analogue to the AI Safety/Security Institute (AISI) model. Its director, Zhou Bowen, chairs TC260's Working Group 9 on AI safety standard setting. However, CnAISDA is structurally and functionally [different from](#) many AISI's.
- [CAICT](#) (**China Academy of Information and Communications Technology**): Government-backed think tank under the Ministry of Industry and Information Technology. Runs voluntary AI safety testing for Chinese models in collaboration with the industry alliance.

Loss of "Operator" Control

The system is outside the control of a particular person or organisation, even though other actors may still be capable of intervening.

One prominent use of *shikong* in Chinese discourse concerns a breakdown the relationship between an AI system and the actor authorised to direct it. That actor may be an individual user, a deploying organisation, an AI developer, or a military commander. Each loss of "operator" control case may differ greatly in severity, but they are connected in that they share the same failure mode of the control relationship, with the system acting outside the instructions or authority assigned to it, and its operator unable to reliably prevent, interrupt, or correct that behaviour.

The mildest usage of *shikong* in this sense appears following OpenClaw's proliferation in China, which signals a concern over the product safety of AI agent casual deployment. Chinese state-media coverage [described](#) incident

such as OpenClaw deleting Meta employees' emails as instances of *shikong*. Such cases are far removed from the irreversible, civilisation-wide scenario often associated with loss of control in English-language AI-safety discourse. Other actors could still intervene, and the resulting harm was bounded. Thus the *shikong* here can be considered as a domestically originated use to describe a local and potentially reversible failure of human operators' direction, which is largely detached from the origins of Western LoC discourse.

A similar understanding appears in industry and regulatory discourse. In an April 2026 speech, an Ant Group representative [warned](#) of “permission-boundary loss of control” () in multi-level delegation scenarios downstream agents may acquire capabilities beyond the scope authorised upstream, producing a “cascading amplification of permissions” (). A May 2026 joint policy opinion on AI agents showed a more significant elevation of *shikong* into central-government AI policy. Issued by CAC, the National Development and Reform Commission (NDRC), and the Ministry of Industry and Information Technology (MIIT), it lists “operational loss of control” () among the security risks developers should guard against.⁹ But this appearance is less exciting than one might assume. The document warns against “data poisoning, privacy leakage, algorithm tampering, system vulnerabilities, and operational loss of control” (while proposing mitigations like data security, privacy protection, and permission management, which are closer to a product-safety toolkit than frontier-risk framework.

Loss of “State” Control

AI-enabled risks exceed the capacity of individual operators or organisations manage and threaten the state’s ability to preserve security, systemic stability

or social order. ¹⁰

It might seem that operator-level LoC could absorb almost every case. A failure has to start somewhere, and usually that somewhere is a system escaping the person directing it. However, an important concern is who is responsible for controlling and containing it. In China, the LoC discourse points to one answer – the state. This section outlines how *shikong* is used across Chinese discourse to consider the state losing control of AI. Most usage in this register converges on specific high-stakes deployment scenarios. Some, like the military commander case, begin with an operator-level failure, but many others do not require one at all.

The first case is the military application. PLA media outlets have [warned](#) of “violent loss of control” (), which refers to the situation when a system is subject to interference from data bias or encounters enemy deception, misjudges targets, triggers accidental firefights, causes a crisis to rapidly escalate, or even exceeds the scope of effective human control. When military loss of control of AI does cause escalation between states, the case becomes more than a failure of “operator” control. A commander is the operator of the system, but also a node in the state apparatus – and in China the military answers directly to the Party – so an operational failure that escalates poses risks the state itself must answer for.

The second example of LoC at the state-level involves the state failing to control *the knowledge* AI possesses about Chemical, Biological, Radiological and Nuclear threats (CBRN). In September 2025, China’s most influential AI standards body TC260 issued “[AI Safety Governance Framework 2.0](#)”, identifying “CBRN knowledge and capability loss of control” as a risk itself, and calling for stronger management of AI end use, restrictions on its use in nuclear, biological, chemical, and missile-related scenarios, and measures to

prevent AI systems from being misused. Even if the system remains under an operator's supervision – such as in a misuse case where a terrorist successfully leverages AI for bioweapons – it can still represent a broader failure of control.

One should not read this example as Chinese actors confusing the two broad categories of misuse and LoC. As in many other policy documents, malicious misuse itself is clearly separated from LoC.¹¹ Researchers at Chinese AI research institutes have articulated a link between misuse and LoC by [defining](#) models' dangerous capability as “harm-enabling capability” – including the knowledge-based capabilities (i.e., CBRN threats) and execution-related capabilities (autonomy and cyber misuse) – as related to loss of control risks. They see LoC happening when AI wants to (i.e., misaligned intention) and can (i.e., harm-enabling capability and oversight evasion) evade human controls to cause harm, with harm-enabling capability playing an important role in making that LoC happen. In other words, if existential LoC can be understood as a combination of capability and intent that AI wants to and can carry out events to wipe out humanity – then dangerous capability itself is an important demonstrator of the “can” part.

The third framing of state-level control concerns the societal consequence of deploying AI at scale in critical systems. As AI is deployed across physical infrastructure, financial systems, essential services, and national-security institutions, the locus of responsibility for maintaining control shifts from an individual operator towards the state as guarantor of systemic stability and public order. The AI Safety Governance Framework 2.0 again reflects a similar understanding, warning that AI deployed in energy, telecommunications, finance, and transportation may produce “operation and execution loss of control” () as a result of hallucinations or external attacks.

Lastly, the same logic extends to AI's effects on the social fabric. CAC's regulation of Anthropomorphic AI –addressing parasocial relationships, child safety, and self-harm – initially [states](#) that the government should apply “inclusive-but-prudent, tiered-and-classified oversight” to prevent “misuse and loss of control” () ¹² in its draft for public feedback. While perhaps each individual losing control of their interaction with AI and becoming more manipulable is a less urgent concern, it scales up into societal issues quickly and requires state interventions.

Loss of “Human” Control

The system itself is beyond the control of anyone on Earth, permanently.

The existence of operator- and state-level LoC does not mean China does not care about humanity losing control of AI. In fact, the discourse of loss of “human” control in China arrives very early and has been consistently mentioned throughout high-level speeches and technical guidance.

As early as 2021, the idea of putting AI under human control to prevent *shikong* has become a talking point, albeit without a clear message on how to execute it. In 2021, the Ministry of Science and Technology ¹³ [released](#) the “Ethical Norms for New Generation Artificial Intelligence”, which explicitly states that AI should be “always under human control” by ensuring “humans have fully autonomous decision-making rights and that they have the right to accept or reject AI-provided services, the right to withdraw from AI interactions at any time, and the right to terminate AI system operations at any time”. In the same year, TC260's “Guidelines for the Prevention of AI Ethical and Safety Risks” ¹⁴ [listed](#) *shikong* risk as the first among five categories of AI ethical-safety risk ¹⁵, and set tiered prevention requirements for developers – who were told to exercise caution in R&D on autonomous AI

with self-replication or self-improvement capabilities, and to assess possible loss-of-control risks on an ongoing basis.

In the past year, LoC at the humanity level has incrementally gained more attention and substance. The September 2025 AI Safety Governance Framework 2.0 explicitly listed “emergence of AI self-awareness and loss of human control ()” as a distinct risk category, warning that in the future, AI may “undergo sudden, unexpected leaps in intelligence, enabling to autonomously acquire external resources, replicate itself, and develop self-awareness” and thus driving AI to “seek external power and pose risks of competing with humanity for control.” Separately, Shanghai AI Lab, whose director now [leads](#) China’s dedicated AI Safety Working Group on standard setting, [released](#) the “Frontier AI Risk Management Framework 2.0” in July 2026 with Concordia AI. The framework offers one of the most operationally detailed Chinese treatments of loss-of-control risk, distinguishing passive loss of control – where humans gradually cede oversight – from active loss of control, in which AI systems deliberately subvert human control, and explicitly identifying internal deployment within AI labs as a critical risk environment.

Making Sense of the Framework

The three loci are neither rigid categories nor stages in a single linear escalation. They identify different control relationships invoked by Chinese uses of *shikong*. In practice, however, failures may propagate across them depending on the operator, deployment context, scale, consequences, and ability to contain or counter the situation.

WHO LOSES CONTROL?

Three loci of *shikong*

LOCUS	WHAT IS LOST	EXAMPLE CHINESE TERMS	EXAMPLES
01 An operator	The system is outside the control of a particular person or organisation	运行失控 loss of operational control 权限失控 loss of permissions control	OpenClaw Agent delegation chains
02 The state	Exceeds what operators can manage; security, stability, order at risk Overlaps with the operator locus: deliberate misuse is control lost by one party and exercised by another.	暴力失控 loss of control over violence 操作执行失控 loss of control over execution 滥用失控 loss of control over misuse	Military escalation CBRN · cyber Critical infrastructure Manipulation at scale
03 Humanity	Beyond the control of anyone on Earth, permanently	脱离人类控制 escape from human control	Self-awareness Recursive self-improvement Power-seeking

Failures at one locus can propagate to the next: capability supplies the “can”, automated R&D the reach.

Oxford China Policy

Operator-level LoC often carries less gravity than its English-language cognate, as an agent deleting emails is not an existential catastrophe. When the operator is a frontier AI researcher, and the system acts within the development pipeline, a product-safety problem begins to resemble the narrower LoC scenarios discussed in Western AI safety.

Similarly, while the state losing control of AI seems to be a uniquely China concern, the importance of governments in exercising control over AI shou

not be dismissed. Moreover, some of the state-level concerns do overlap with what AI safety has in mind about LoC, such as autonomous cyberattack capability or the ability to manipulate humans.

Internal vs. External Loss of Control

Despite these overlapping concerns over LoC, however, one distinct difference between the Chinese and the Western discourse is that the former lacks significant internal lab deployment discussion. While theoretically both operator- and state-level concerns about AI agent misalignment and dangerous capabilities could translate into internal lab monitoring, it has yet to happen in reality, with most conversations about *shikong* concentrated on post-release deployment, external systems, and societal scale.

Moreover, even when *shikong* is used to describe news about internal lab events, as with the Open AI-Hugging Face incident, the locus of the failure appears to be more external than internal. For example, People's Net, an influential Party online news outlet, reposted an article that [describes](#) the Hugging Face case as *shikong*, but frames the essence of the failure as “the model’s capability exceeds humans’ pre-designed safeguards”, which appears closer to a product-safety diagnosis at the operator level. The article opens not with the question of whether AI is misaligned, but whether “cybersecurity loses control ()”, treating the event as an external infrastructure problem rather than an internal AI development one.

Another article about the same incident, from the Ministry of State Security, also thoroughly covers the case as *shikong*, [flagging](#) [\[archive\]](#) the risks of AI agents autonomously discovering vulnerabilities, identifying targets, and executing intrusions. Nevertheless, the security implications identified are about closed models operating as “black boxes” that cannot be monitored i

real time, the emergence of agentic capabilities, and the human-set permissions and isolation being broken –= with the warning landing on AI agents' risk when deployed in critical sectors. Moreover, the article ends with warnings to individual citizens: guard personal information and avoid foreign AI tools of unknown provenance, apply rational judgment to AI-generated content and defer to official sources on policy and public safety, and never input classified or internally sensitive material into any internet-connected tool. The core message, despite using OpenAI's incident as a hook, is about state-level concerns on AI agents' deployment in critical systems and on social stability and data governance.

These two examples show that even when a clear instance of lab-internal LoC emerges, the Chinese reading of it may still settle on product safety or state-level risk rather than the development pipeline itself. Part of this may simply reflect less familiarity with AI misalignment as a concept in China. But another reason for the overwhelming external and product safety emphasis on *shikong* may be the structural difference in how AI labs operate in China and the US, which results in the split between internal and external deployment that underpins Western concern over internal lab deployment being less visible in China. Chinese labs [are likely to release](#) their most capable models quickly rather than sitting on several internal versions the way American counterparts do ¹⁶, so the gap between what's inside the lab and what's outside is smaller, and the messy, complicated deployment environment becomes the more salient concern.

This also means that even when risks that are closer to Western understanding of LoC have emerged, as seen in the reporting of the OpenAI Hugging Face incident, mitigation methods could be oriented towards traditional product safety in China. For instance, the TC260 draft standard AI agent security from July 2026 [uses](#) “privilege loss of control” () to

describe unauthorized access, unauthorized execution, environment jailbreak and runtime escape. Although environment jailbreak and runtime escape overlap with Western concerns about systems evading controls, the proposed solutions are least-privilege design and permission verification.

Box 2: Is (technological loss of control) a Chinese equivalent of Western LoC?

The phrase *jishu shikong* () has appeared more frequently in senior Chinese policy discourse since 2025, including in statements by Premier Li Qiang and President Xi Jinping. Because it combines “technology” with “loss of control,” it is tempting to read it as a stable reference to humanity-level LoC — or even to the narrower frontier-AI scenario.

The evidence points in multiple directions. On one hand, at the 2026 Summer World Economic Forum in Dalian, Li Qiang [used](#) the phrase when comparing the era of AI with the Cambrian, warning that if governance fails to keep pace, the consequences of *jishu shikong* and *lunli shifan* (loss of ethical norms) “will very likely be severe.” The mentions here definitely gesture at something graver than routine product failure.

On the other hand, the Ministry of National Defence [used](#) the same phrase in March 2024 to describe conflict escalation triggered by technical failure in military AI (“[AI] can easily lead to conflict escalation and arms race because of *jishu shikong*”)— which is closer to a state-level concern about weapons systems rather than an existential one about AI itself. In a January 2026 Politburo session, Xi [referenced](#) *jishu shikong* for quantum technology, biomanufacturing, fusion

energy, and brain-computer interfaces, but warned of this risk together with those of loss of ethical norms and data misuse. More curiously, AI is not included in the list of technologies he addressed as having *jishu shikong* concerns.

The term, in short, has not consolidated on a single meaning. It can sound severe or routine, existential or operational, AI-specific or technology-general, depending on who is speaking and in what context. Reading any one usage as definitive risks over-indexing on language that is still finding its own definition.

Implications for AI Safety Engagement

The ambiguity about LoC documented here is not an argument specifically or against engaging China on loss-of-control risks. Instead, ambiguity means that the discourse remains open to interpretation and change, both within China and through engagement with the broader international AI safety community.

Since President Xi's WAIC speech, bottom-up efforts to define *shikong* are accelerating. An outlet operated by the Chongqing Municipal Party Committee's Organization Department [issued](#) a piece titled "AI cannot lose control", covering risks from hallucination to emotional manipulation, while People's Daily [interview](#) with leading scholars specifically flagged frontier risks – AI autonomously discovering vulnerabilities, launching attacks, and iteratively improving itself – and warned that China lacks effective means to respond. In China, it is common for a term to first appear in leadership speech which is then followed by bottom-up efforts to define and claim it. *Shikong* may follow this pattern, creating a temporary opening in which expert and institutional interpretations can influence how the term is operationalised.

Similarly, later reporting on the OpenAI-Hugging Face incident, including METR and Redwood Research's independent investigation, was widely covered by Chinese media, including state outlets. Much of this coverage continued to use *shikong* language while incorporating substantially more detail about what had happened inside OpenAI's evaluation environment. ¹⁷

For those tracking China's AI discourse for potential engagement on AI governance, this piece suggests three takeaways.

Do not over-index on top-level mentions. It is tempting to read the speech of the president or the premier as a signal that China now cares about loss of control in the operationalised sense and, because the two most important people said it, that implementation and coordination can follow quickly. This underestimates two things. First, the complexity of the discourse itself: as this piece has shown, *shikong* in China spans operator-, state-, and humanity-level referents, and top-leadership usage leans on a phrase – *jishu shikong* – whose own meaning remains unstable, deployed as readily for autonomous weapons as for existential risk. A mention tells very little about which loss of control is being mentioned. Second, the mechanics of operationalization: even where the state signals intent, converting a broad rhetorical concern into evaluations, standards, and enforcement is slow and uncertain work. China's operationalization energy to date has flowed mostly toward operator-level product safety rather than lab-internal deployment. Every mention of LoC in leadership speech matters for monitoring AI risk awareness in China, but it is just the beginning of a question. What we are witnessing is mostly the send out of top-level signals, whereas much of the work rests on how people are going to define and operationalise the term on the ground in many policy documents and technical papers.

Work on the operational problems. Converging on the word “loss of control” matters far less than building the shared capacity to actually test for it. Currently, [automating AI R&D](#) is on Chinese companies’ agendas, and Chinese frontier labs are already developing and evaluating capable coding agents for internal and production-oriented workflows.¹⁸ Publicly, however, Chinese AI safety discourse has paid comparatively less attention to the specific risks of internally deployed agents inside frontier labs than recent US and UK work has. China could plausibly stand up state or third-party evaluation structures for lab-internal deployment relatively quickly, but the knowledge of how to evaluate loss-of-control risks is harder to gain overnight. This is not a China-specific deficit, as the world’s leading AI scientists don’t yet know how to ensure the technology remains safe as it grows more powerful, but technical research on frontier AI safety is a relatively newer field in China than in the United States. As [noted](#) by Matt Sheehan from Carnegie Endowment, sharing approaches to testing for dangerous capabilities – including the ability to evade the control of a human operator – can be one of the most valuable outcomes of US-China AI engagement.

Treat the vagueness of LoC as everyone’s problem, but be specific. The looseness documented here is not a China problem. Policymakers and legislators in the US and Europe have also folded quite different risks into the same phrase, and Western AI safety’s own operationalised sense of the term “lab-internal agents subverting controls” – is barely legible outside of that community. The emerging consensus that irreversible loss of control over AI would be extremely dangerous for humanity is valuable, but it will only be successfully operationalized when people specify what they mean by “loss of control” – especially when speaking to audiences outside Western AI safety circles. This variety of how LoC is used in China serves as a useful reminder for Western AI-safety communities. Outside specialised frontier-AI circles, “loss

of control” can naturally evoke not only lab-internal agents but also military AI, critical infrastructure, dangerous-capability misuse, and other severe deployment scenarios that do not fit neatly within the current operational focus on internal AI R&D.

The discourse of LoC is changing quickly and becoming more prominent and substantive as more AI incidents expose vulnerabilities in existing control mechanisms. It is likely that AI *shikong* (AI LoC) will be mentioned increasingly often, whether in domestic expert commentary, policy guidance, or reporting on international events. But many questions remain open. In future major AI safety documents, such as the next update to TC260’s AI Safety Governance Framework, how will the term be used? Will *shikong* gradually narrow toward catastrophic, humanity-level risks? Will internal lab deployment and AI R&D become a more explicit focus? Or will the term continue to span the broad operator-, state-, and humanity-level meanings documented here?

All these questions remain to be answered. Amidst the uncertainty, one thing is clear: appreciating the differences in interpretation that each frontier AI state brings to AI safety discourse will be one key to making progress.

Zilan Qian

Research Associate, Oxford China Policy Lab

Acknowledgements:

Zilan is grateful for the feedback and insights provided by Gabriel Wagner, Sarah Godek, and Jeff Wu on the earlier draft, and would also like to thank Keri Blomquist for detailed feedback and Ben Horton for his editorial support on the final product.

The research and analysis of OCPL experts are their own and do not necessarily reflect organisational views.

- 1 Xinhua, China's official state news agency, [republished](#) the *Science and Technology Daily* report that describes the incident as “OpenAI Model Loses Control, Chinese Open-Source Model ‘Comes to the Rescue’” (OpenAI [The Beijing News](#), a major newspaper under Beijing municipal Party-state supervision, [reported](#) “OpenAI Admits AI Model Loss-of-Control Intrusion Involving Multiple Platforms” (OpenAI [AI](#) [CCTV](#), China's national state broadcaster, [aired](#) segments describing OpenAI models as having “lost control” and invading Hugging Face systems ([OpenAI](#) [Hugging Face](#)).
- 2 The word *shikong* is composed of the verb *shi* (lose) and the noun *kong* (control)
- 3 This piece maps how *shikong* is used in public Chinese AI discourse, but does not attempt to measure the private beliefs held by Chinese officials or researchers. It draws on Chinese-language sources published between 2021 and August 2026, including national standards and guidance, regulatory drafts and final policy documents, state and Party media, technical papers and frameworks from Chinese AI institutes, and industry speeches and commentary. The framework for understanding *shikong* is the author's own inductive categorization rather than official Chinese framework. Translations are the author's unless otherwise noted. Claims about the absence or relative salience of a usage describe the public sources reviewed here and cannot rule out discussions that are unpublished or otherwise not publicly accessible.
- 4 In the abovementioned OpenAI-Hugging Face incident, models [exploited](#) a vulnerability during a cyber evaluation, gained internet access, and compromise external infrastructure while pursuing the benchmark objective. On July 30, Anthropic [disclosed](#) separate cases in which models reached real external systems because of a testing configuration error, without appearing to realise the target:

were outside the simulation. On August 4, UK AI Security Institute (AISI) [publish](#) a blog post revealing that AI agents took unsanctioned actions against real people and organisations. On August 6, US startup Frontier Security [reported](#) that Moonshot AI's Kimi K3 exploited a misconfigured test environment to access external websites.

- 5 “The West” in this article refers mainly to the US, UK and EU, where AI safety discourse on LoC has been increasingly salient and various regulations and research agenda connected to the concept.
- 6 The 2026 Singapore Consensus on Global AI Safety Research Priorities [follows](#) the same definition of LoC as the IAISR 2026; The 2024 International Dialogues on AI Safety Beijing Consensus [uses](#) Cold War scientific coordination as an analogy to LoC, arguing that humanity must again coordinate to avert catastrophe from unprecedented technology.
- 7 Several AI safety organizations and frontier labs explicitly treat internal deployment as a setting in which broader loss-of-control risks may emerge. The UK AI Security Institute (AISI) [describes](#) its evaluations of models internally deployed within AI companies to perform AI safety research as part of its broader work on “loss-of-control risks from frontier AI”; METR [assesses](#) whether internal agents have the means, motive, and opportunity to create a “rogue deployment” (an autonomous deployment operating without human knowledge or permission) and frames this as a potential pathway toward AI takeover, an extreme form of loss of control; Apollo Research [argues](#) that internal deployment for automated AI R&D can itself create a pathway to loss of control, with misaligned systems interfering with oversight or training, influencing the development of successor models, and creating feedback loops that become progressively harder for humans to monitor or stop.
- 8 In February 2026, Meta AI researcher Summer Yue [posted](#) [\[archive\]](#) on twitter that an OpenClaw agent began deleting and archiving more than 200 messages despite being instructed only to suggest emails for deletion. She failed to stop the agent

through command and could only terminated it manually. Chinese state-media coverage subsequently [described](#) the episode as “Meta’s head of AI safety, Sumn Yue, watched helplessly as an loss-of-control (*shikong*) “lobster” (nickname of OpenClaw in China) rapidly deleted more than 200 emails from her inbox”.

- 9 Earlier TC260 technical guidance had already discussed *shikong*, but the joint opinion gave the concept greater policy salience by incorporating it into a cross agency central-government framework for AI agents.
- 10 Throughout this section, “state” is used as shorthand for the Party-state apparatus. In China's political system the Party and the state are not neatly separable – the party leads the state, the military, and the legislative system. Some usages of *shikong* in this register are oriented more toward Party control specifically (e.g., the military, which answers to the Central Military Commission rather than the civilian government), but disaggregating “Party” from “state” in every instance would require a level of institutional specificity that is beyond the length and the scope of this article.
- 11 For example, in the technical standard document “[Artificial intelligence—Risk management capability assessment](#)”, CBRN as misuse () is explained as “applied to dual-use scientific risks (such as the development and malicious use biological, chemical, radiological, and nuclear weapons)”.
- 12 is translated here as "misuse and loss of control" rather than as a single compound (“misuse-related LoC”). The relationship between the two concepts varies across Chinese sources, as discussed in the text. This article follows some prominent translation in the field such as Paul Tsai China Center ([China Law Translate](#)) to separate these two phrases.
- 13 The Ministry of Science and Technology (MOST), sometimes described as “the OSTP of China”, played a major role in China’s early AI development plans and led on some “soft law” approaches to AI governance.

- 14 “Guidance on Network Security Standardised Practice ()” is no binding guidance issued directly by the TC260 secretariat, usually aiming at providing provisional direction for fast-moving issues, but is not yet formal enough to be binding national standards.
- 15 The other four are societal, infringement, discrimination, and liability risks.
- 16 Nathan Lamber [has argued](#) that, in contrast with the US frontier labs that often keep stronger internal models unreleased for longer periods, Z.ai’s GLM release cycle may be measured in days rather than months. Qwen has likewise maintain a rapid and continuous public release cadence, [according to](#) Hugging Face. Moonshot launched Kimi K3 across its coding plan and API before releasing some additional technical materials and model components. While these examples do not prove that Chinese labs generally conduct less extensive or rigorous pre-deployment evaluation, they suggest that the boundary between a lab’s most capable internal models and its publicly available models in China may often be thinner than at major US frontier labs, which usually withhold for months some internal models that are far more capable than those released publicly (see, for example, p.8 of METR’s [Frontier Risk Report](#)).
- 17 On 26 August 2026, METR and Redwood Research [published](#) an independent investigation on the OpenAI-Hugging Face incident that happened in July. They found that roughly 1,200 agents that were supposed to be isolated from one another had discovered an unauthorised shared “message board,” exchanged more than 70,000 messages and files, and coordinated attempts to cheat the ExploitC evaluator. [Based on the report](#), Xinhua [reported](#) that some 700 agents had participated in the intrusion and highlighted their unauthorised coordination, attempts to deceive the scoring system, and efforts to falsify or delete their own execution records. *People’s Daily* [republished](#) *Science and Technology Daily* coverage, stressing that the incident exposed a new AI-safety problem where or multiple agents can operate computers, call tools, and pursue tasks over time, traditional security boundaries may no longer be sufficient. One report from *Fir Financial*, a major Shanghai-based financial news outlet backed by Shanghai Me

Group, [emphasized](#) that the main model involved was an internal-only OpenAI research model. [Many other](#) technology outlets also used more dramatic headlines [such as](#) “OpenAI Model Loss of Control Was Terrifying: Ghost Misjudgments, 1,2 Agents.”

18 For example, MiniMax [reports](#) using M2.7 for 30-50% of its Reinforcement Learning team’s daily iteration work in March 2026. ByteDance says Seed2.1 [participates](#) directly in evaluation, data generation, training, research, and infrastructure tasks across its model-development pipeline in June. And Moonshot AI [is building and evaluating](#) Kimi for internal, production-oriented coding and agent workflows.

Subscribe to Oxford China Policy Lab

Launched 3 years ago

The Oxford China Policy Lab is an interdisciplinary hub of China-focused research at Oxford. We produce policy-relevant research to mitigate global risks stemming from US-China great power competition.

Subscribe

By subscribing, you agree Substack's [Terms of Use](#), and acknowledge its [Information Collection Notice](#) and [Privacy Policy](#).



39 Likes • 7 Restacks



A guest post by

Zilan Qian

Writing about US-China AI, to one day not write about US-China AI.

Subscribe to Zilan Qian

Discussion about this post

Comments Restacks



Write a comment...